



KubeCon



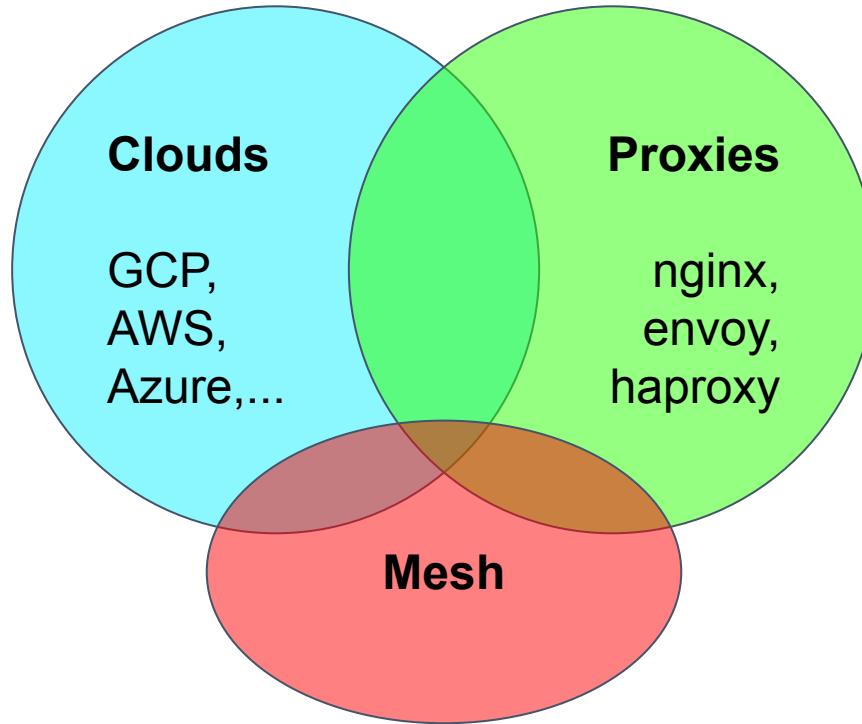
CloudNativeCon

Europe 2019

Ingress V2 and Multi-Cluster Services

Rohit Ramkumar (rramkumar1@)
BOWEI DU (bowei@)

Next “Ingress” API



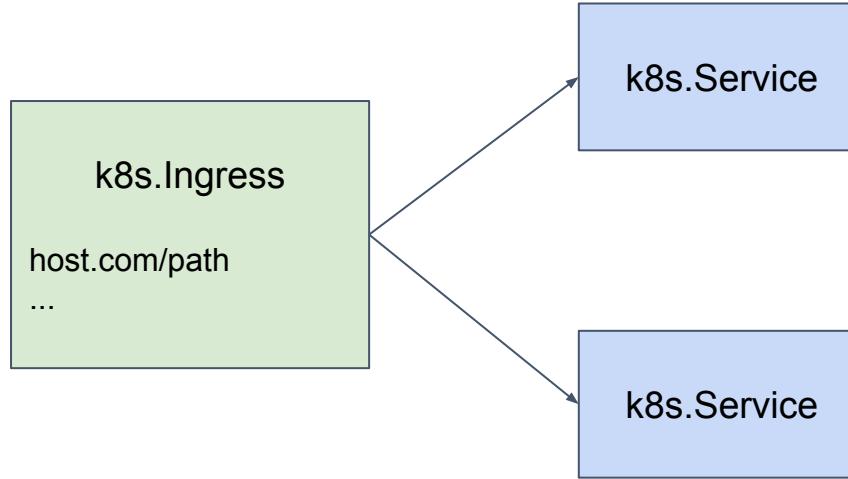
Next “Ingress” API

Challenges: portability, expressiveness, features

Lots of APIs now in this space.

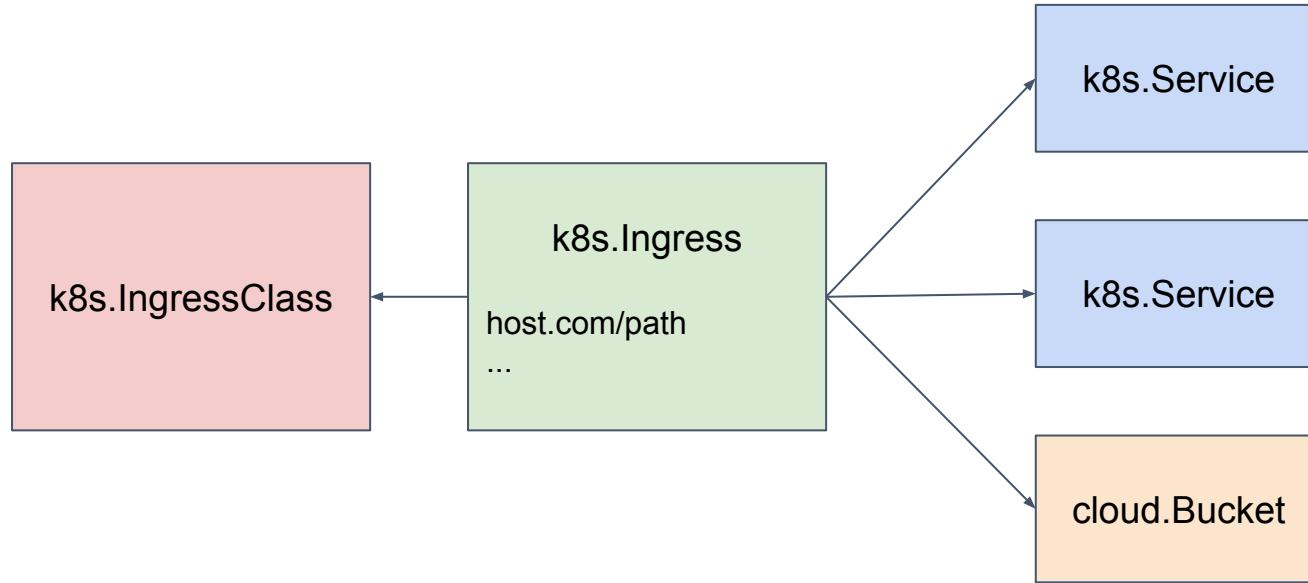
What can we learn from them?

API models: Ingress



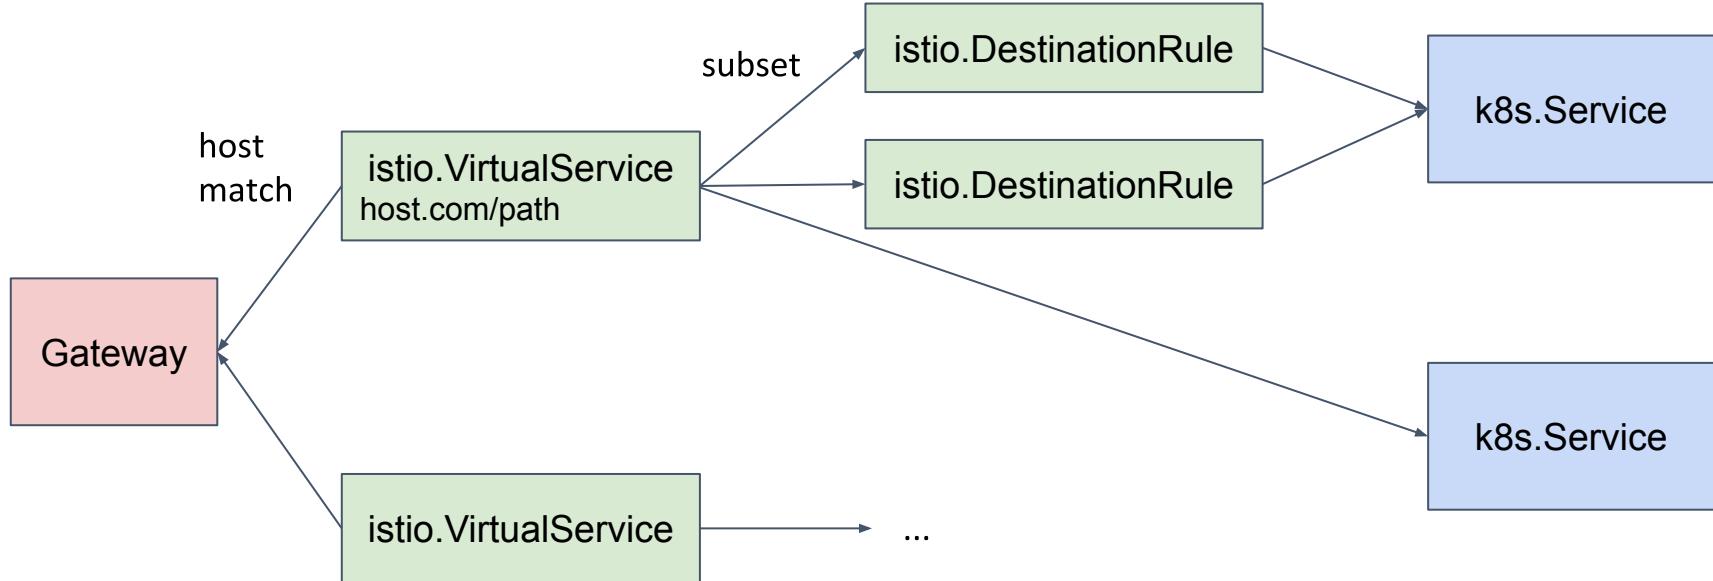
Key attributes: self-service, single-role, simple

API models: Ingress GA (proposed)



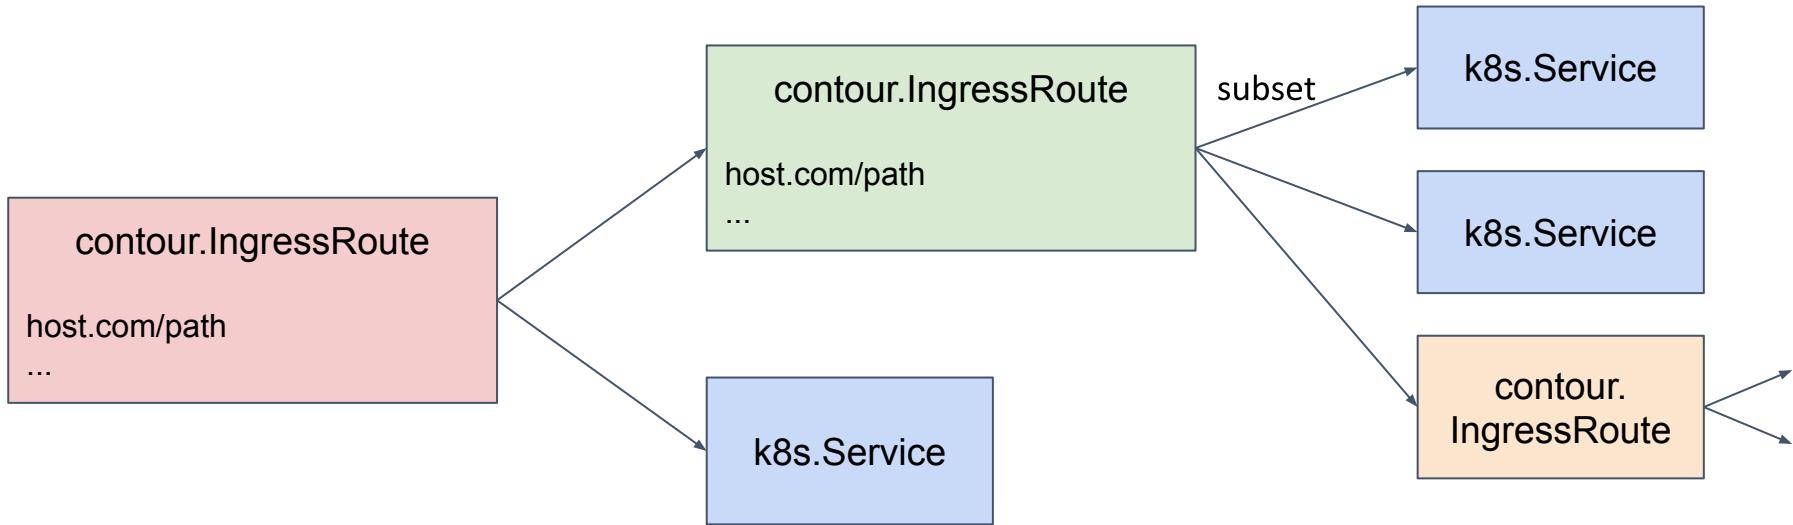
Key attributes: “flavor” of Ingress via class, heterogenous backends.

API models: Istio



Key attributes: separate roles for proxy infrastructure, application definition; rich feature set based on Envoy

API models: IngressRoute



Key attributes: Separate roles and recursive delegation, composable, additional features such as traffic splitting.

API models: Others

Lots of annotations for features (Everybody)

Decorators on k8s.Service instead of k8s.Ingress (Ambassador)

Use Custom Resources (Gloo)

What have we learned?



KubeCon



CloudNativeCon

Europe 2019

Multi-role environments

- Infrastructure vs App dev

How much portability?

- Should be a user choice.

Support future API growth

- Claim: providers/features will converge over time (but quite slowly).

A modest proposal

Warning: this is very early...

Taking from existing work...

Shape of the resources (model)

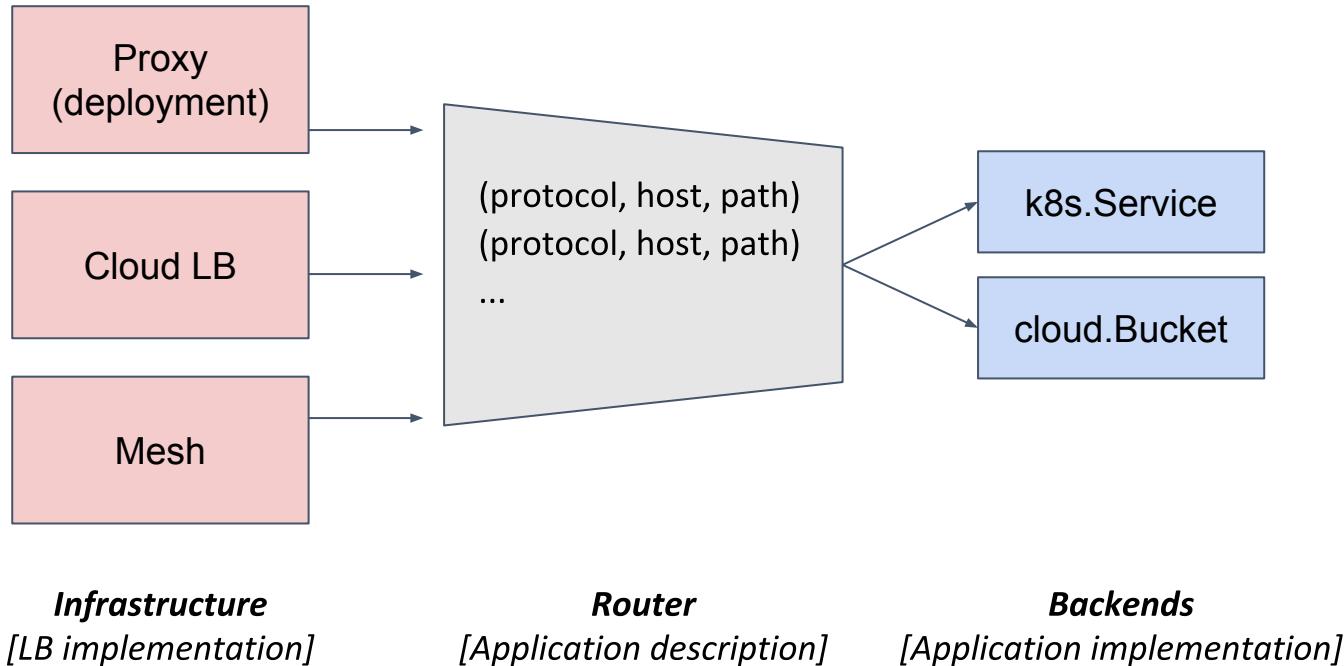
Addressing portability and extensibility

Future directions

A modest proposal: the model



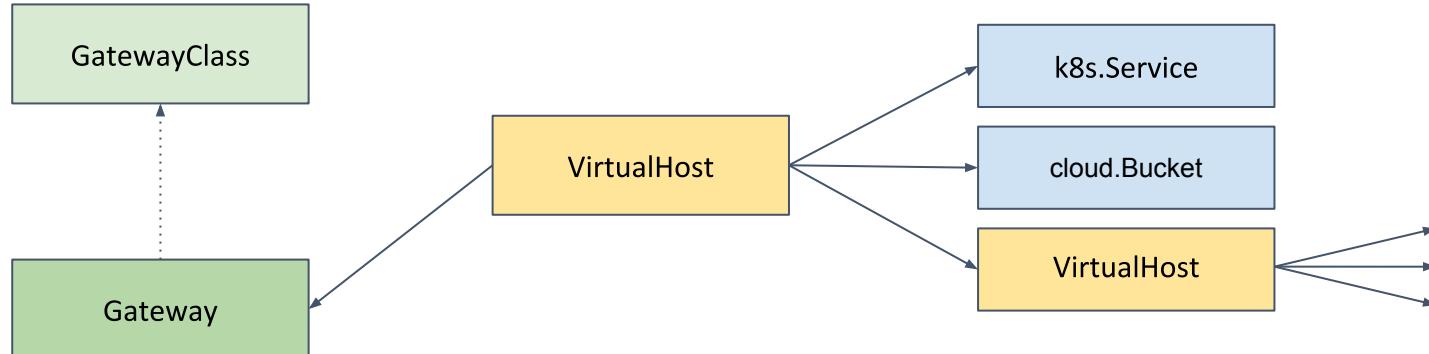
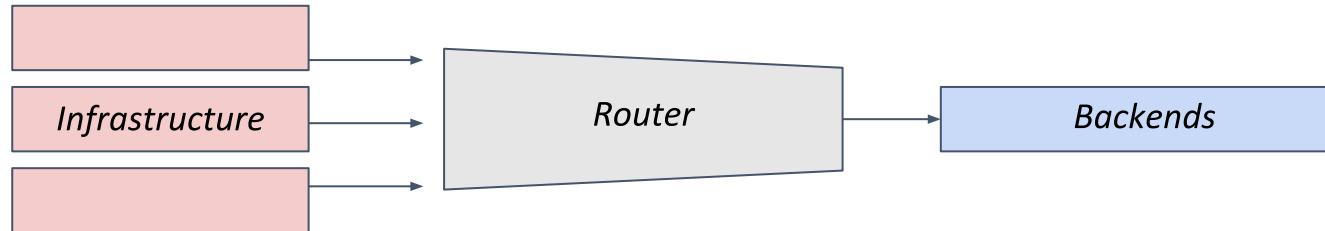
CloudNativeCon
Europe 2019



A modest proposal: the model

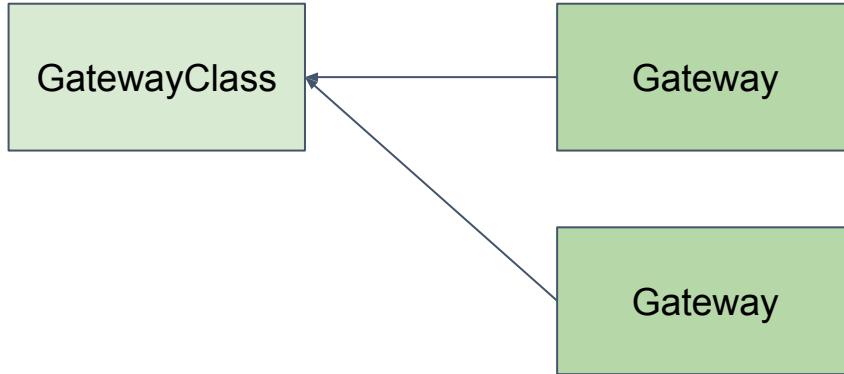


CloudNativeCon
Europe 2019



*Names are temporary

Gateway



Deployment specific options - e.g. mergeable

Abstracts available implementations

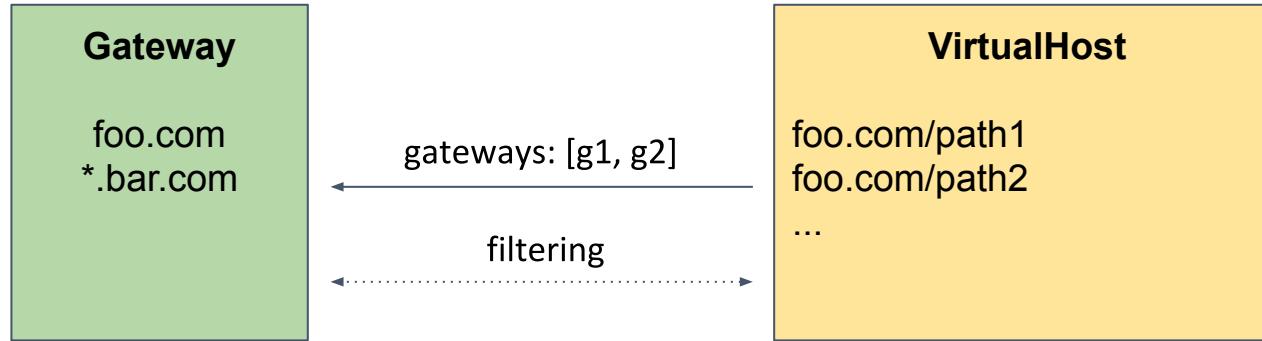
Represents actual instance of LB/proxy infrastructure, capacity

Protocol termination (<IP:port>, TLS)

Gateway ↔ VirtualHost



CloudNativeCon
Europe 2019



How to match Gateway and VirtualHost?

- VirtualHost attaches to a Gateway
- Gateway filters VirtualHosts -- wildcard allows for self-service.

Portability



KubeCon | CloudNativeCon
Europe 2019

Core

MUST be supported.

Extended

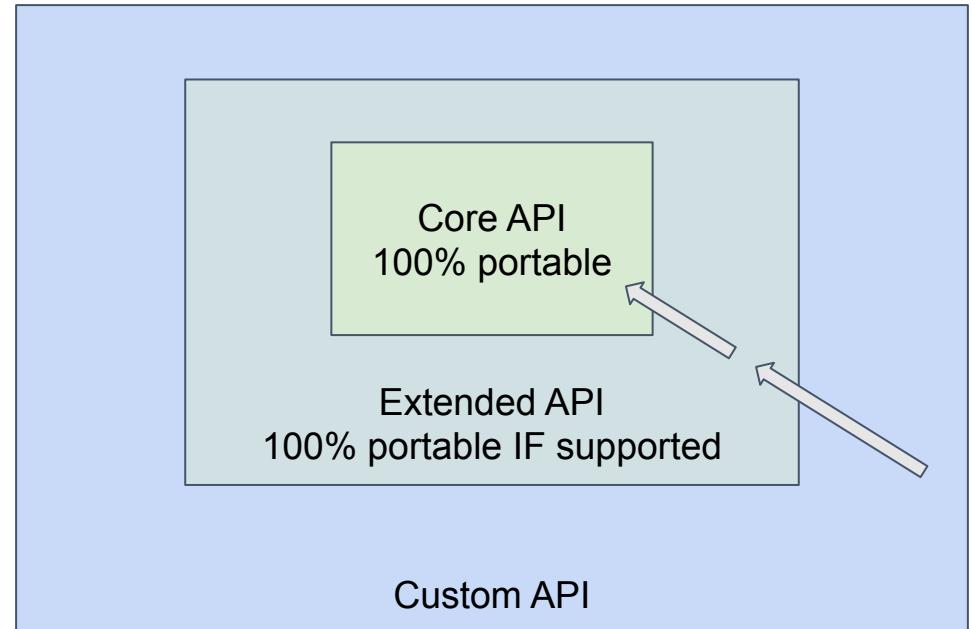
Feature by feature.

MAYBE supported, but
MUST be portable.

Part of k8s API schema.

Custom

No guarantee for portability,
No k8s API schema.



Portability

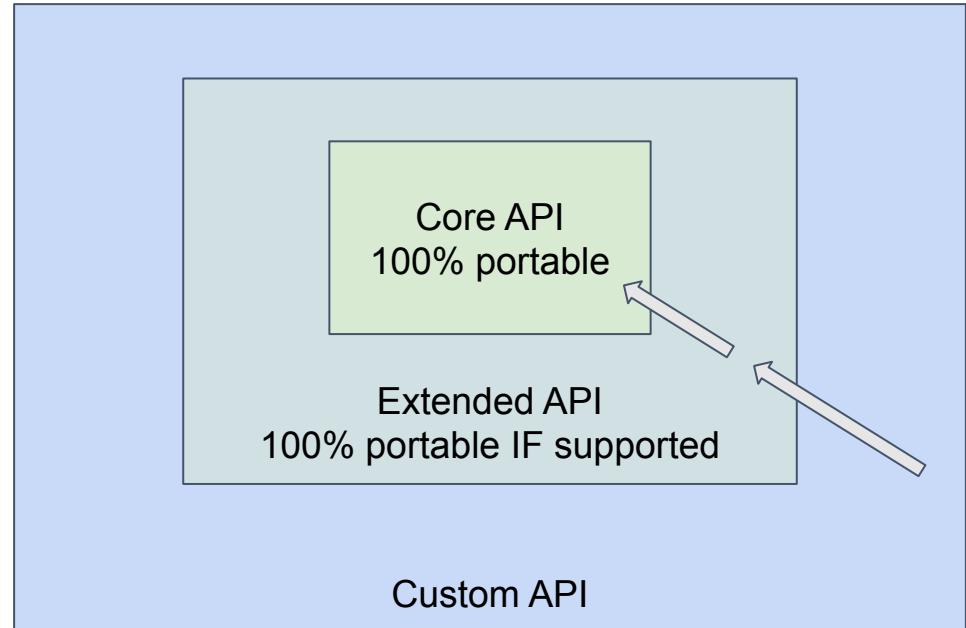


KubeCon | CloudNativeCon
Europe 2019

Enforcement by conformance tests

Extended feature definition requires self-contained conformance.

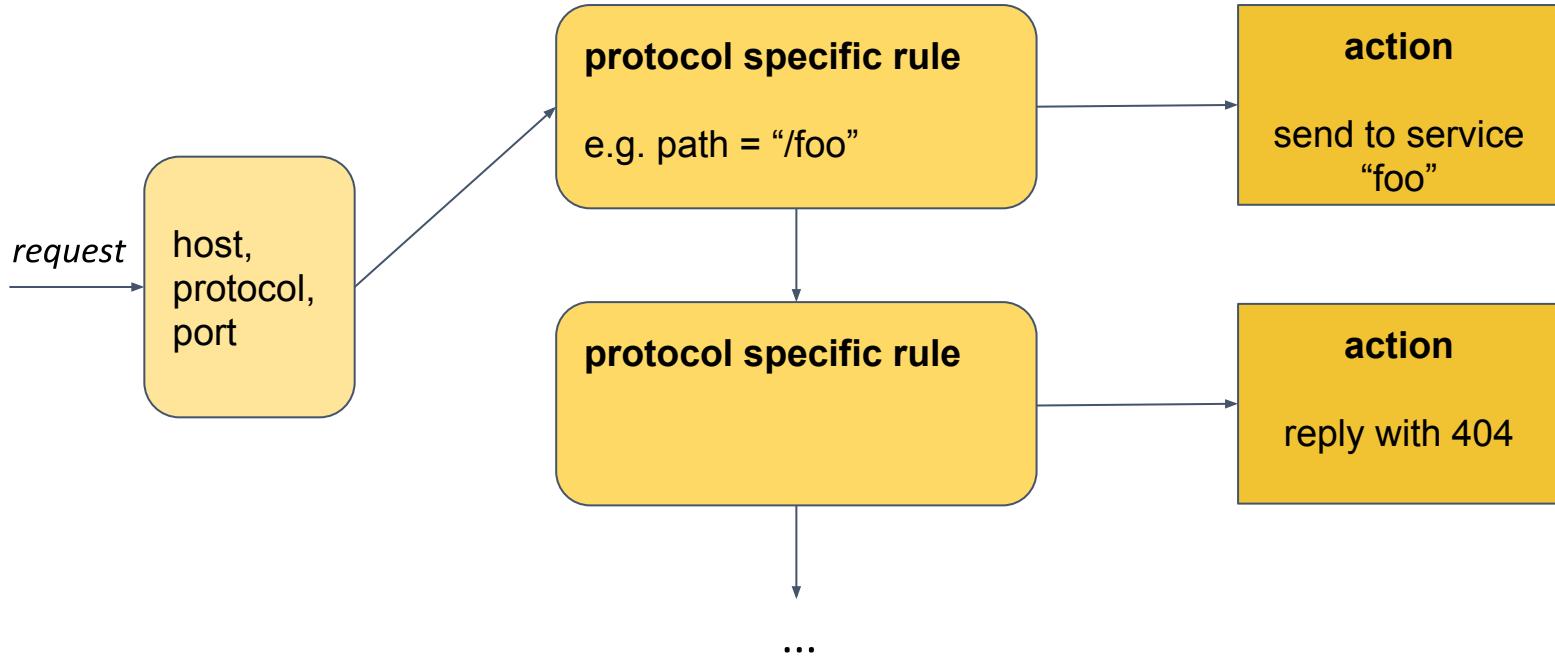
Require all extended features be checked statically?



VirtualHost request matching



CloudNativeCon
Europe 2019



VirtualHost (warning: sketch)



CloudNativeCon
Europe 2019

```
type VirtualHost struct {
    Routes []Route
}

type Route struct {
    HTTP *HTTPRoute
    TCP  *TCPRoute
}

type HTTPRoute struct {
    Host string
    Rules []HTTPRouteRule
}

type HTTPRouteRule struct {
    // Matchers, with room for extension
    Path      string
    // Actions, with room for extensions
    Backend   *HTTPRouteActionBackend
    StatusCode *HTTPRouteActionStatusCode
}
```

```
# VirtualHost
spec:
  routes:
    - http:
        host: "foo.com"
        rules:
          - path: "/"
            backend:
              service: foo-app
              servicePort: www
    - tcp:
        port: 9000
        rules:
          - backend:
              service: tcp-app
              servicePort: my-protocol
```

VirtualHost extensibility

Complex syntax tree -- needs a decorator pattern

- Better API machinery?
- String/String vs Raw Objects (inline CRDs) vs CRD link

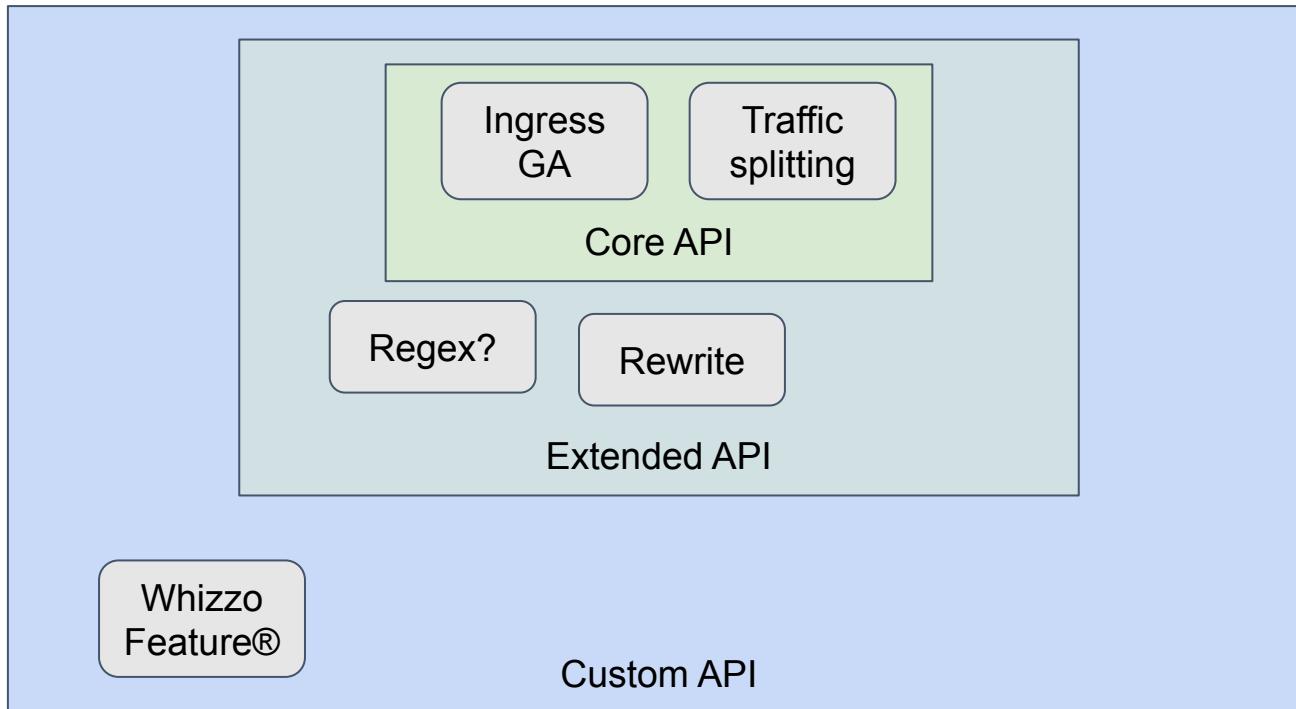
Portability



KubeCon

CloudNativeCon

Europe 2019



Future directions

Interesting alternative backends:

- Storage bucket
- Multi-cluster Services



KubeCon



CloudNativeCon

Europe 2019

Multicluster Service

Multi-Cluster Services: Use Cases



KubeCon



CloudNativeCon

Europe 2019

- Canary - Deploy new version of application in one Kubernetes cluster in one locality before rolling it out worldwide.
- Low latency - Users get routed to the application backends that are closest to them geographically.
- High availability - Users requests are served even if one cluster holding application backends is completely down.

Multi-Cluster Services: User Journey



KubeCon



CloudNativeCon

Europe 2019

Persona: Network Admin

- I want **L7 load balancing across** a global service that is deployed on Kubernetes clusters.
- I want the ability to add / remove clusters from this load balancing based on business requirements.
- I want teams in my organization (service owners) to manage their own ingress traffic but maintain a clear boundary between the “admin” and the “Kubernetes end user”.
- I want my teams to use a native Kubernetes API to leverage default features (e.g namespacing, labels) & to reduce the learning curve.

Multi-Cluster Services: APIs



KubeCon

CloudNativeCon

Europe 2019

```
kind: MultiClusterIngress
metadata:
  name: my-mci
spec:
  template:
    spec:
      backend:
        serviceName: my-mc-service
        servicePort: 80
      rules:
        - host: foo.bar.com
          http:
            paths:
              - backend:
                  serviceName: my-mc-service
                  servicePort: 80
```

```
kind: MultiClusterService
metadata:
  name: my-mc-service
spec:
  template:
    selector:
      app: shopping
    ports:
      - name: web
        protocol: TCP
        port: 80
        targetPort: 80
    clusters:
      - selector:
          matchLabels:
            region: us
```

Multi-Cluster Services: APIs



CloudNativeCon
Europe 2019

```
kind: MultiClusterIngress
metadata:
  name: my-mci
spec:
  template:
    spec:
      backend:
        serviceName: my-mc-service
        servicePort: 80
      rules:
        - host: foo.bar.com
          http:
            paths:
              - backend:
                  serviceName: my-mc-service
                  servicePort: 80
```

```
kind: MultiClusterService
metadata:
  name: my-mc-service
spec:
  template:
    selector:
      app: shopping
    ports:
      - name: web
        protocol: TCP
        port: 80
        targetPort: 80
    clusters:
      - selector:
          matchLabels:
            region: us
```

Multi-Cluster Services: APIs

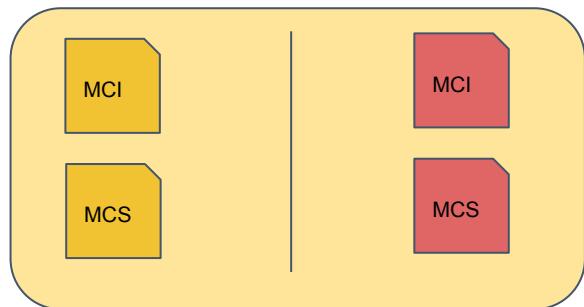


CloudNativeCon
Europe 2019

```
kind: MultiClusterIngress
metadata:
  name: my-mci
spec:
  template:
    spec:
      backend:
        serviceName: my-mc-service
        servicePort: 80
      rules:
        - host: foo.bar.com
          http:
            paths:
              - backend:
                  serviceName: my-mc-service
                  servicePort: 80
```

```
kind: MultiClusterService
metadata:
  name: my-mc-service
spec:
  template:
    selector:
      app: shopping
  ports:
    - name: web
      protocol: TCP
      port: 80
      targetPort: 80
  clusters:
    - selector:
        matchLabels:
          region: us
```

Multi-Cluster Services: Workflow



“Config Source”



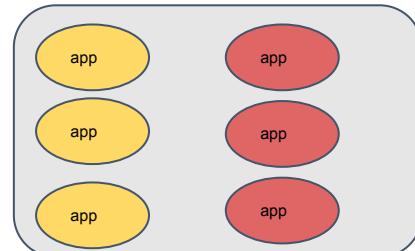
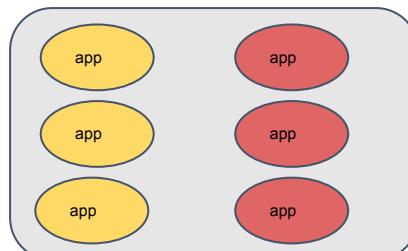
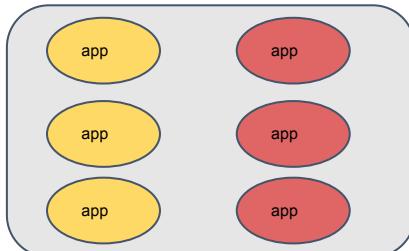
Admin



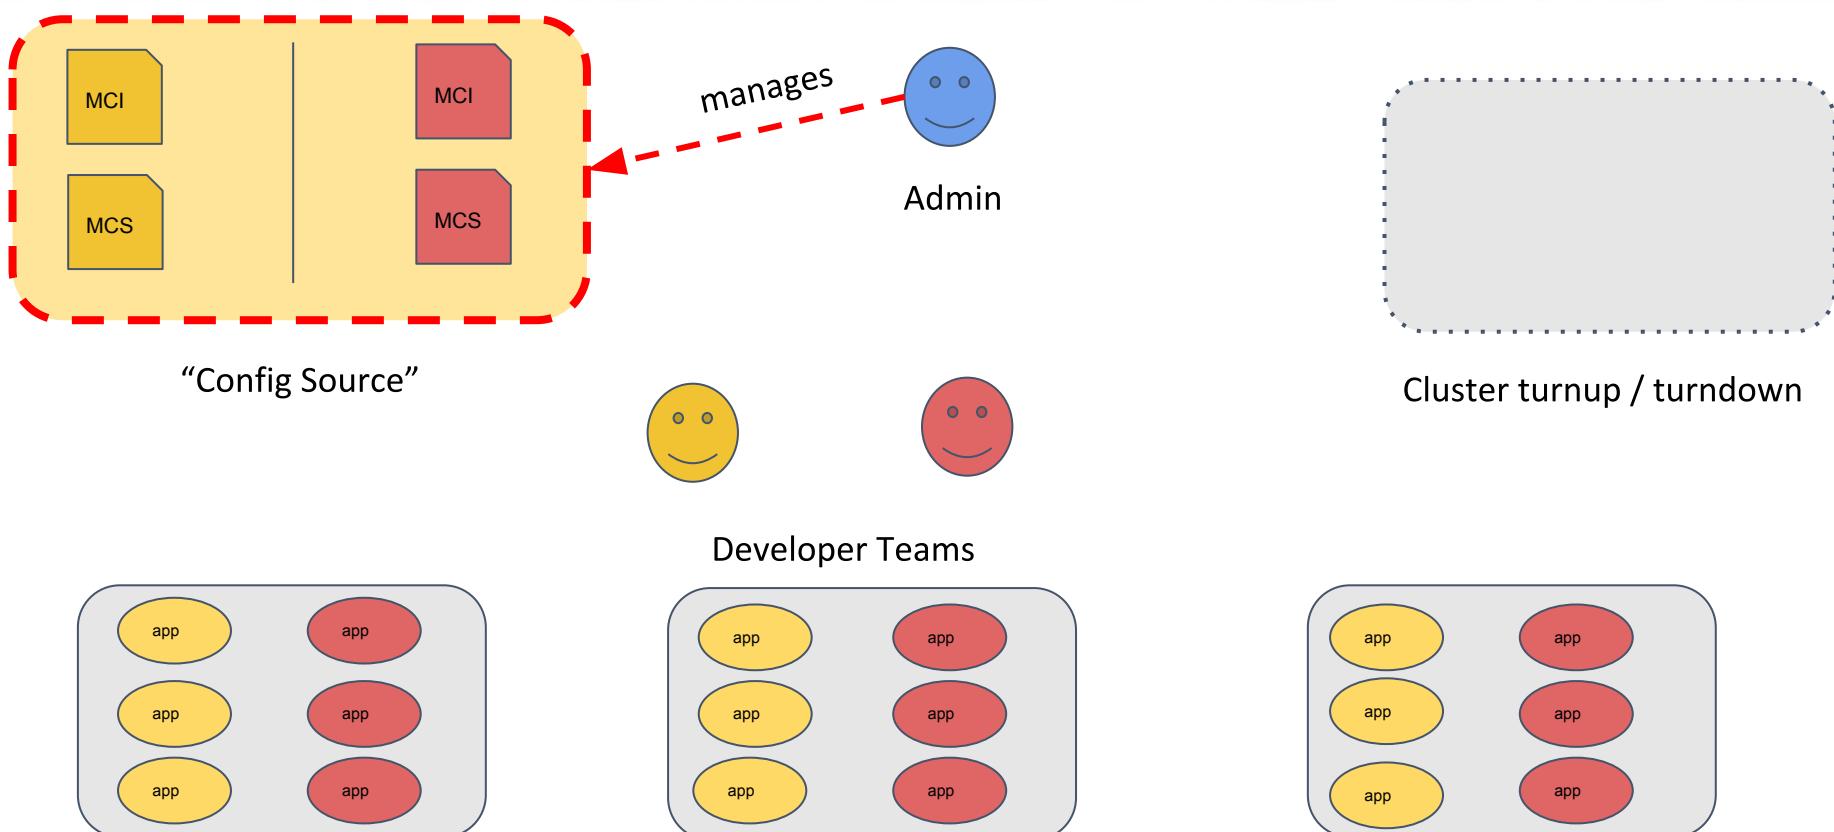
Cluster turnup / turndown



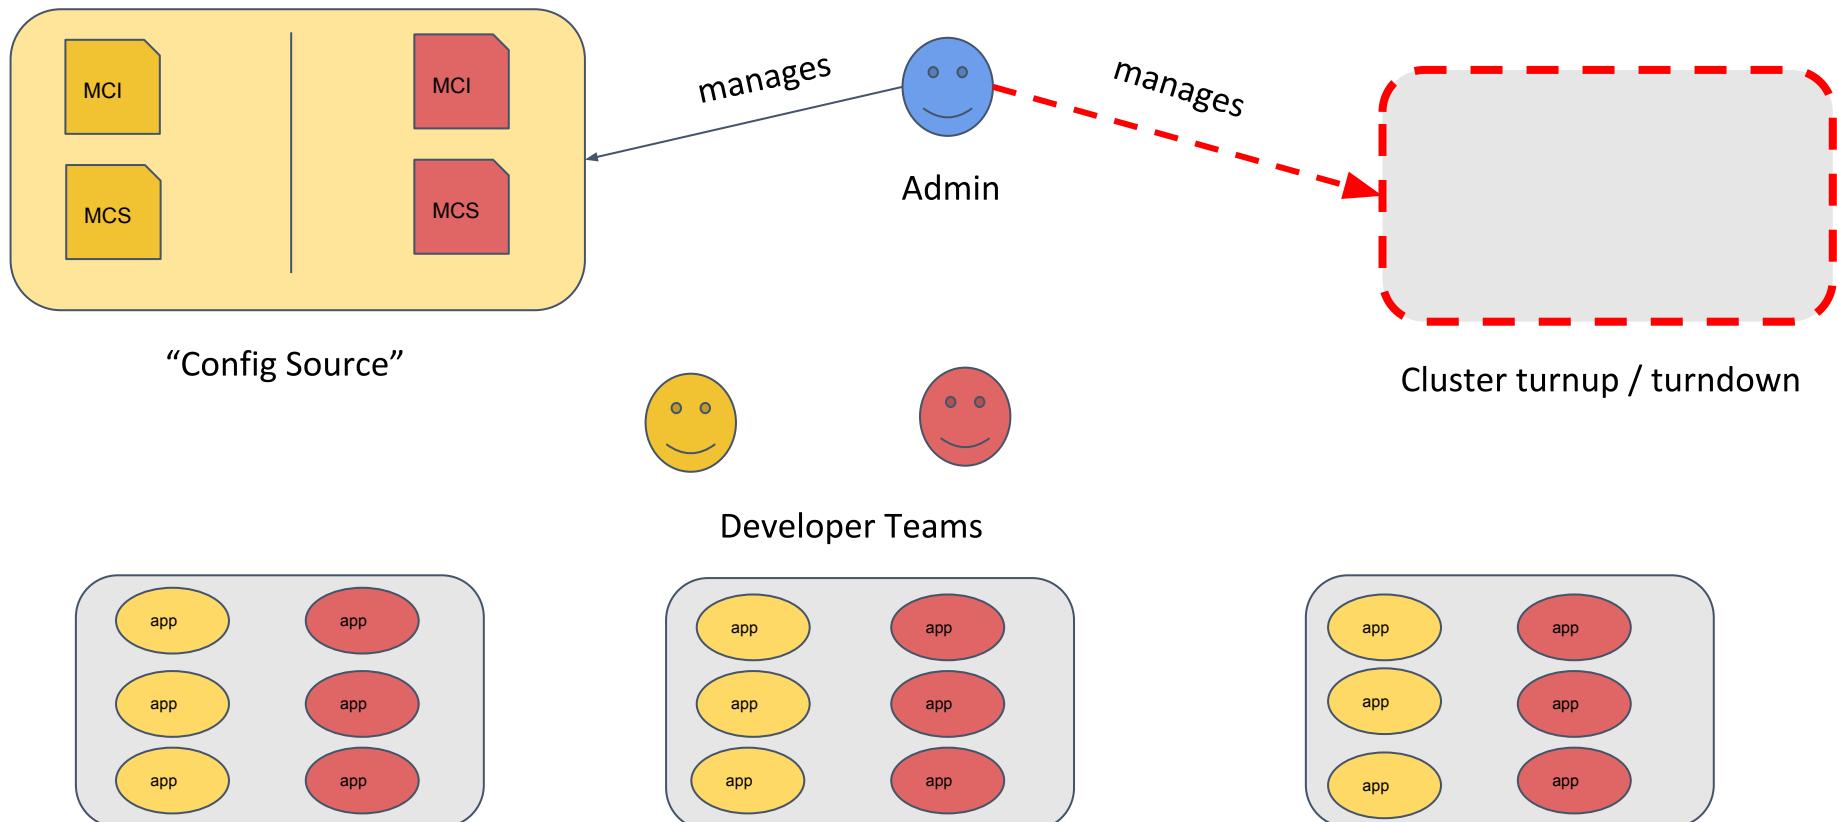
Developer Teams



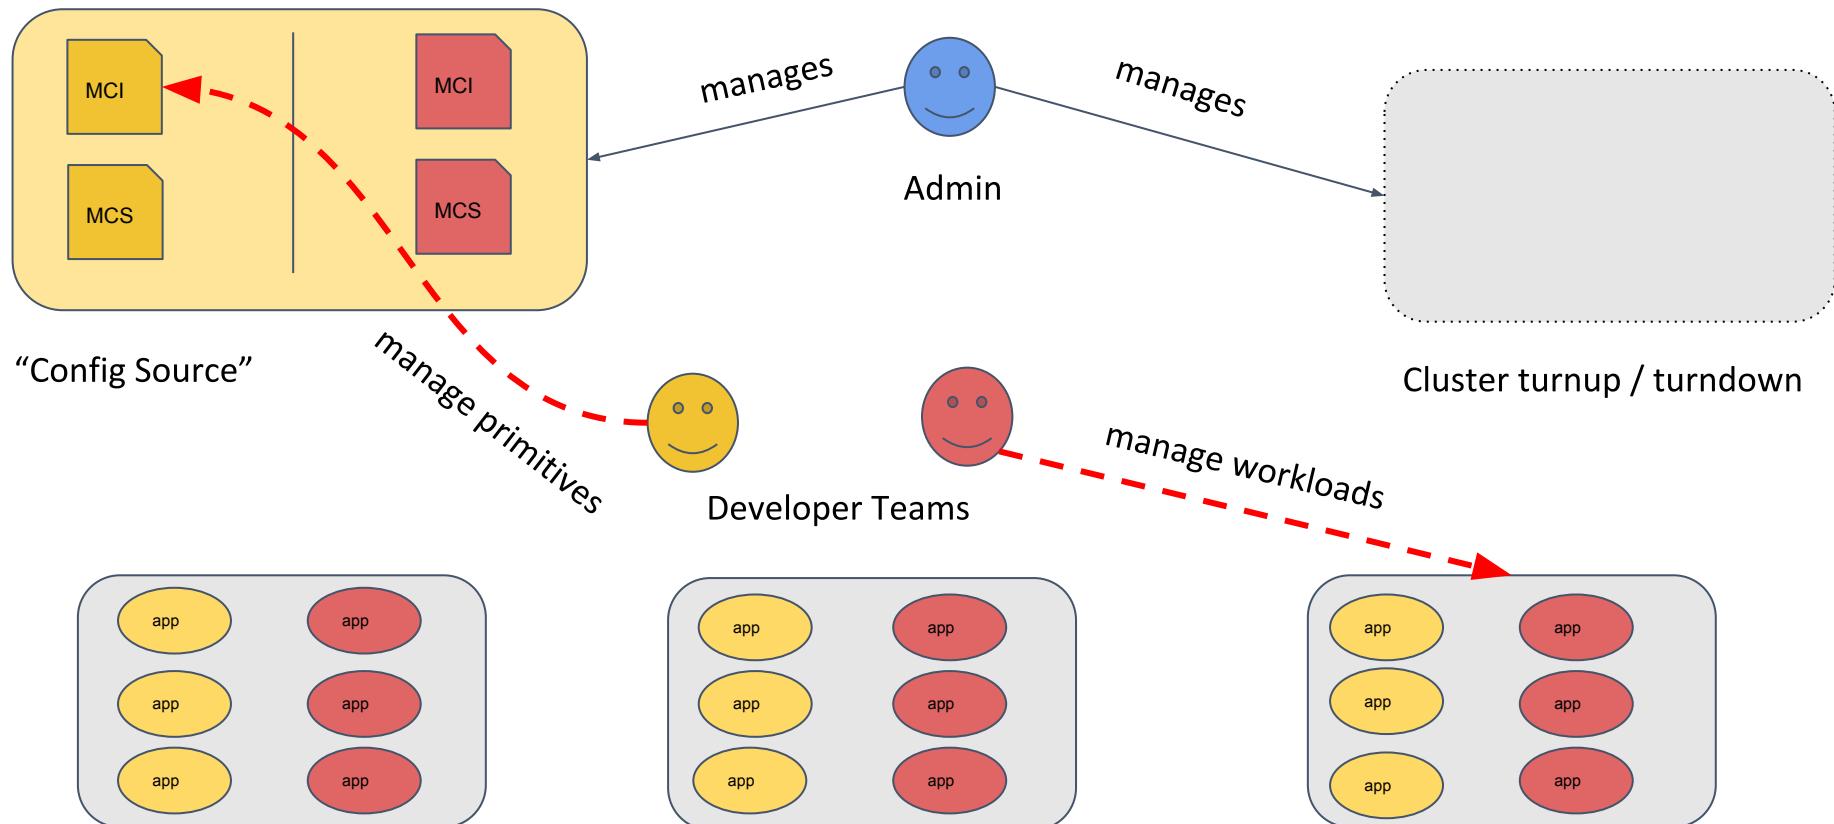
Multi-Cluster Services: Workflow



Multi-Cluster Services: Workflow



Multi-Cluster Services: Workflow



Multi-Cluster Services: FAQ



KubeCon



CloudNativeCon

Europe 2019

- *Why not have a MultiClusterDeployment? That alleviates the need for having to do the manual work of replicating their workload across clusters.*
 - Do not want to be opinionated on how users deploy their workloads. It's not really the spirit of what the API is trying to provide.
 - Best practices in this space are divided, no one-size fits all approach.

Multi-Cluster Services: FAQ



KubeCon



CloudNativeCon

Europe 2019

- *What about Federation v2?*
 - Multi-cluster models & assumptions are slightly different.
 - E.g FederatedIngress is replicated across clusters.
 - Federation v2 is opinionated on cross-cluster DNS & SD, deployment of workloads, etc.
 - In general, Federation v2 is trying to target a much wider set of use cases.

Multi-Cluster Services: FAQ



KubeCon



CloudNativeCon

Europe 2019

- Where does Istio's Multicloud story fit?
 - Potential for existing API's to support use case.

Conclusion

- Expect the KEP proposal to come shortly
- Feedback from community is key
- Reference implementation



KubeCon



CloudNativeCon

Europe 2019

Thanks!

Rohit Ramkumar (rramkumar1@)

Bowei Du (bowei@)

Backup



KubeCon



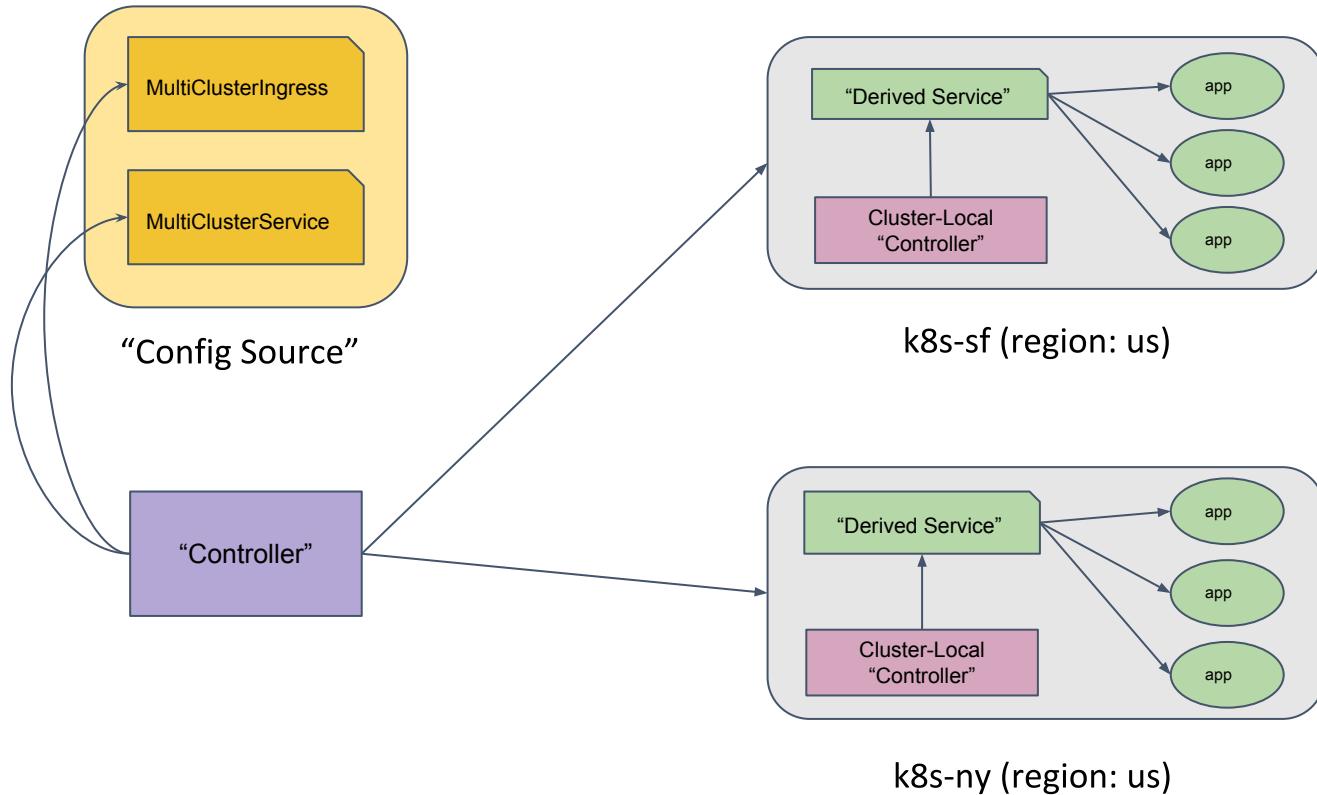
CloudNativeCon

Europe 2019

Multi-Cluster Services: Reference Implementation



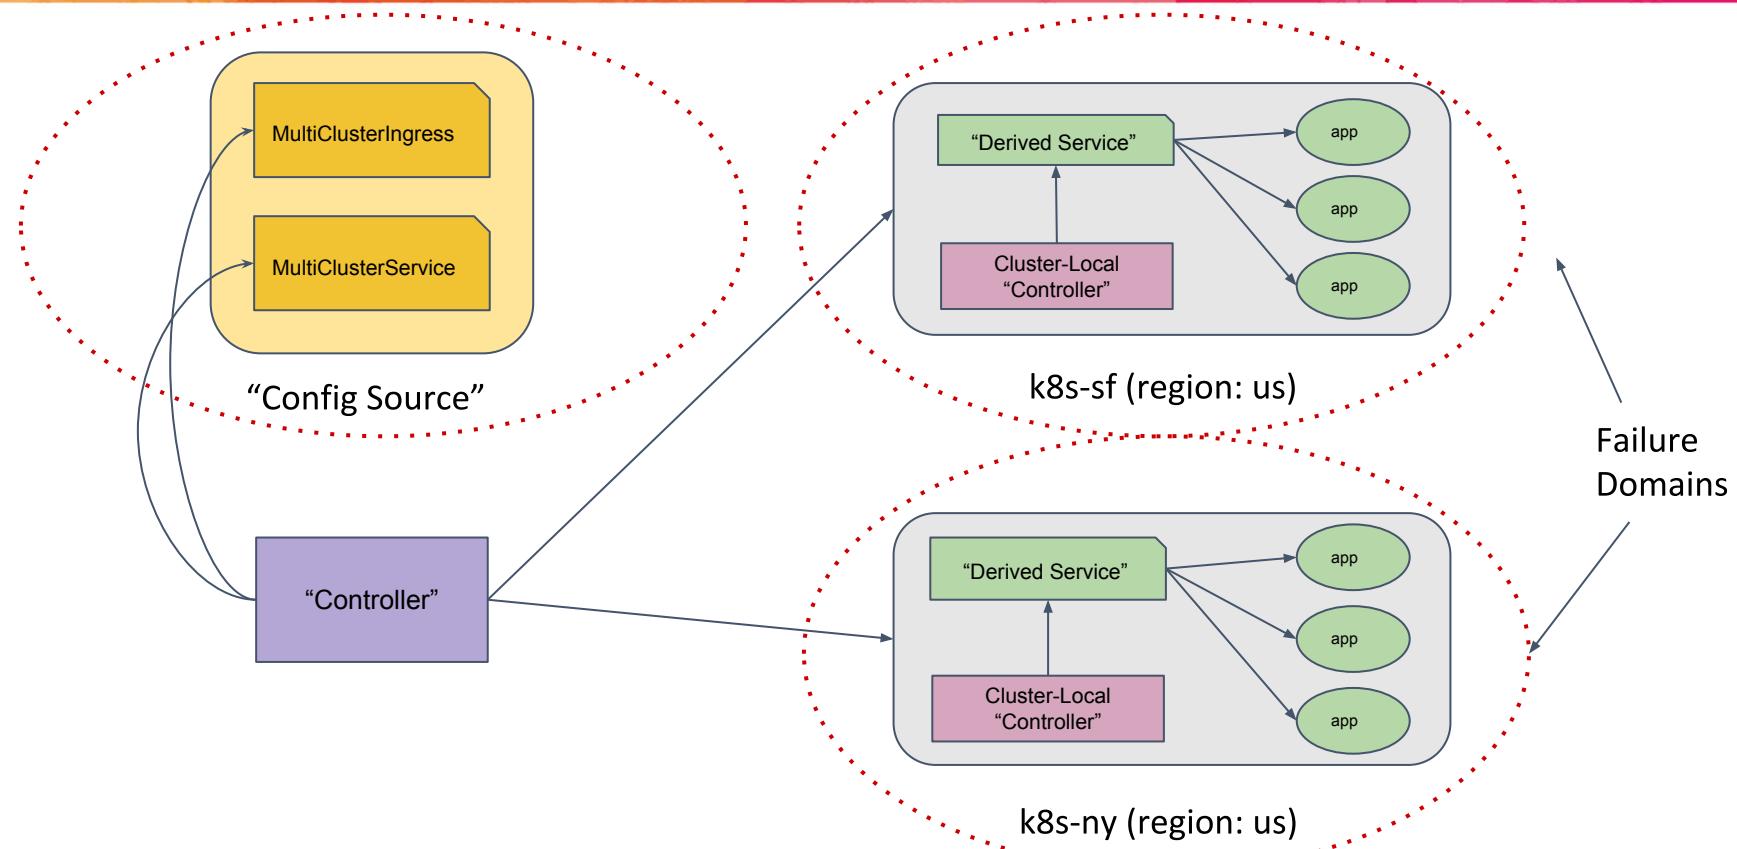
CloudNativeCon
Europe 2019



Multi-Cluster Services: Reference Implementation



CloudNativeCon
Europe 2019



Multi-Cluster Services: Reference Implementation



CloudNativeCon
Europe 2019

