

# SERVICE MESH ❤️ LEGACY

Benefits of a service mesh when integrating Kubernetes with legacy services



Dr. David Meder-Marouelli, Stephan Fudeus





## Stephan Fudeus

- Expert for Container Platforms
- 16 years of distributed systems
- Since 2017 with 1&1 Mail & Media



## Dr. David Meder-Marouelli

- Architect for Automation
- 15 years of IT experience
- Since 2015 with 1&1 Mail & Media

# United Internet / 1&1 Mail & Media



## United Internet

- Leading European internet specialist
- > 9000 employees
- > 90.000 servers in 10 data centers worldwide
- Access and Applications



## 1&1 Mail & Media

- Several free basic services and professional fee-based e-mail solutions
- One of the most powerful online marketing platforms
- > 33 million active users / month
- Multiple brands

# Product Overview

**Communication  
and Organisation**  
E-Mail, Calendar,  
Contacts, SMS, Fax

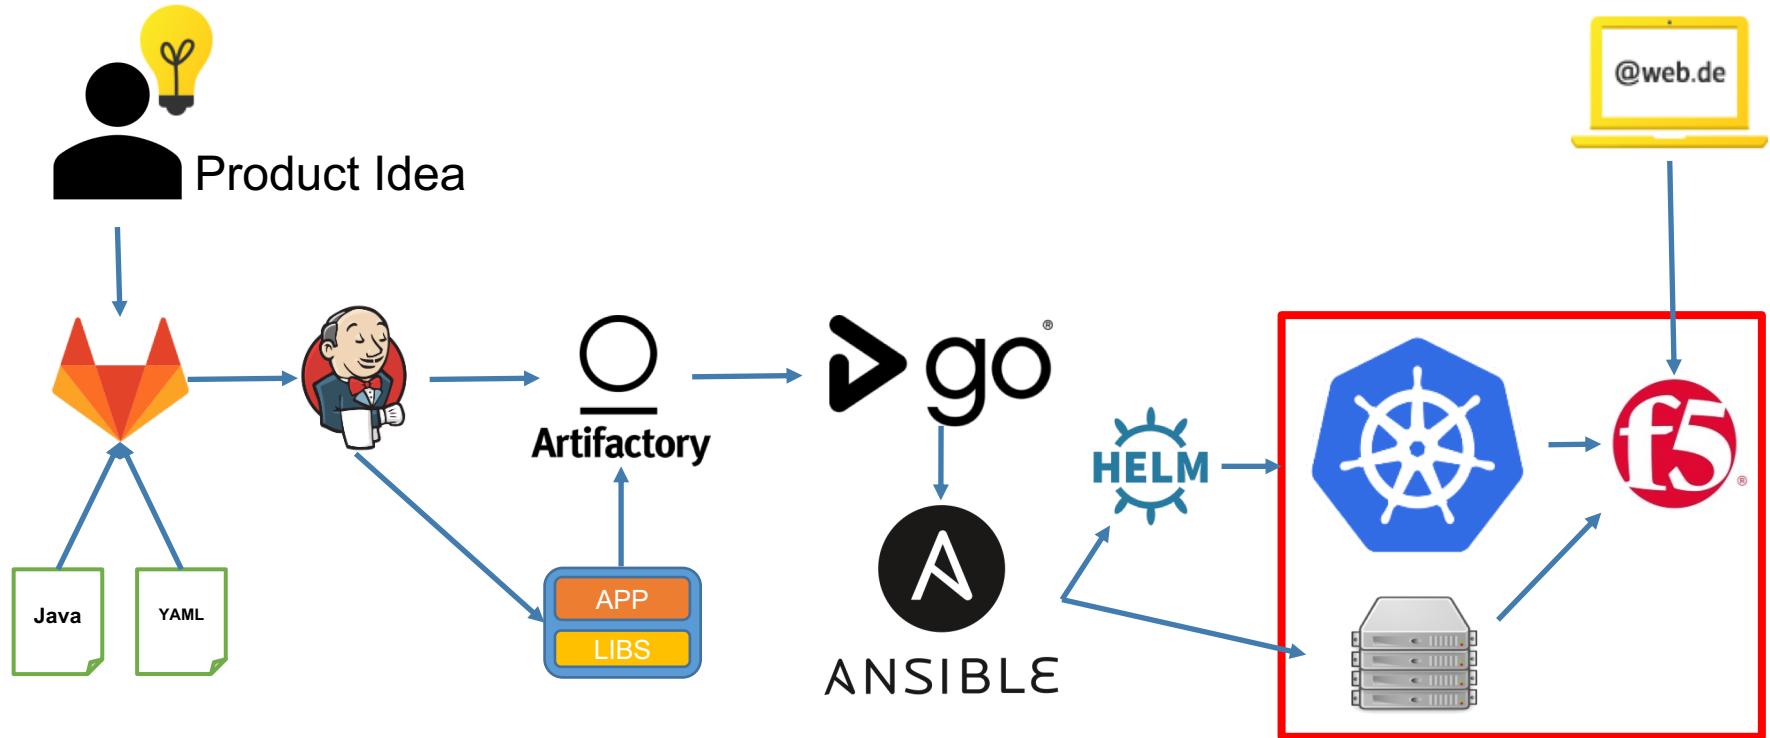


**De-Mail**  
Legally compliant  
communication and  
identity management

**Online Office**  
documents,  
spreadsheets,  
presentations

**Cloud-Storage**  
for photos, videos,  
music and documents

# Ideal Software Lifecycle



# Container Strategy

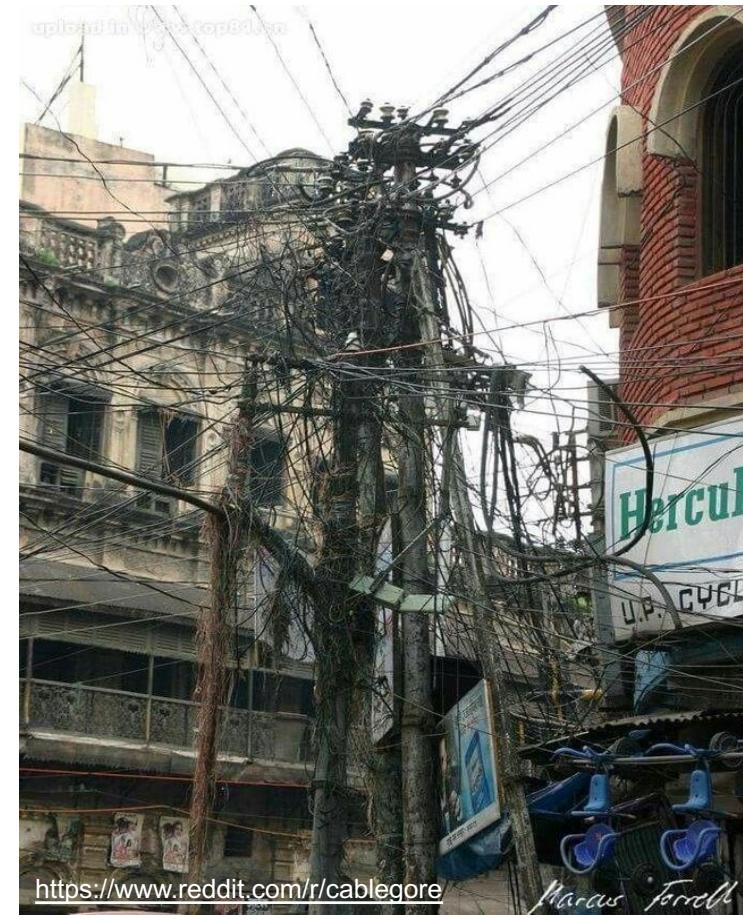
- Kubernetes as centrally provided orchestration platform
  - Fast deployment cycles
  - Focus on soft multi-tenancy
    - Friendly users, but with security in mind
  - Focus on microservices
- Multiple clusters decoupled on network dimensions
  - fe/be/infrastructure, data center, live/non-live
  - bare-metal on-premise
  - non-routable podCIDR and serviceCIDR (RFC 6598 / CGNAT / 100.64.0.0/10)



**kubernetes**

# Legacy Gap

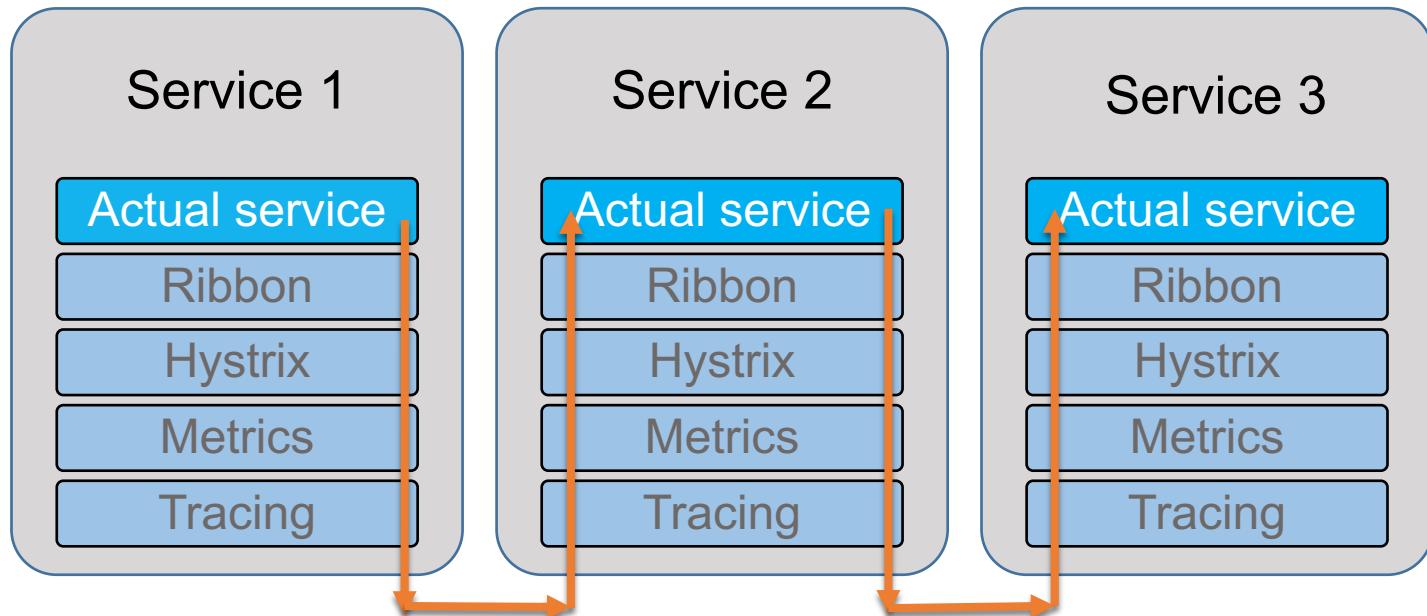
- Ca. 1000 existing services
  - 90% on VMs and bare metal
  - 10% on Kubernetes
- Complex dependencies
- Not cloud (native) ready
  - Not stateless
  - IP based ACLs
- ... some never can/will be
- Both worlds need to interact



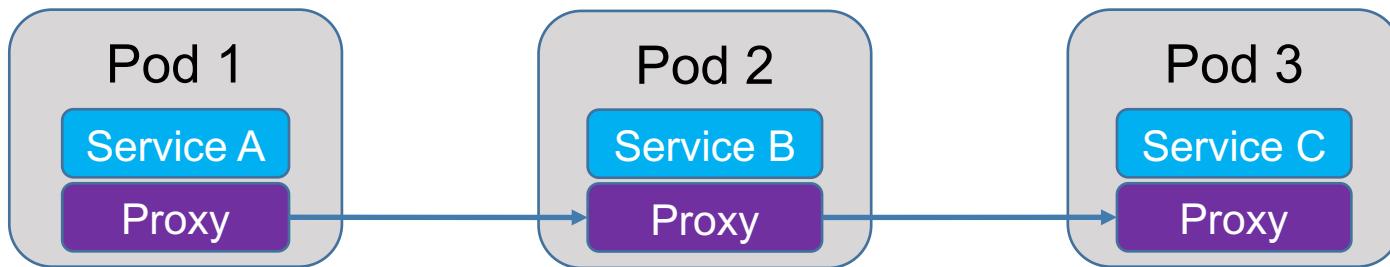
<https://www.reddit.com/r/cablegore>

# Typical Application Stack

- Mostly Java SpringBoot
- Several PHP
- Few Node.js SPA



# Cross cutting concerns @ infrastructure level

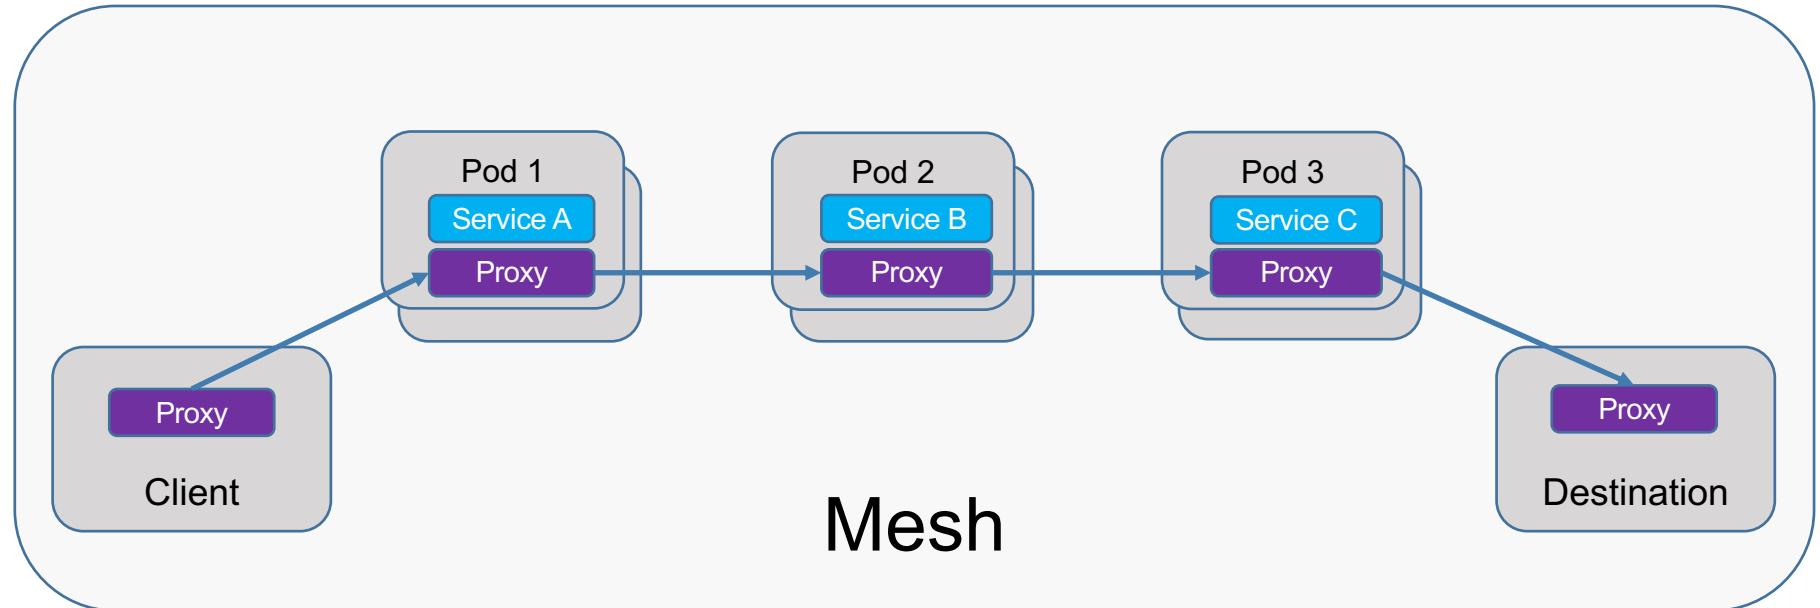


## Benefits of using a Service Mesh

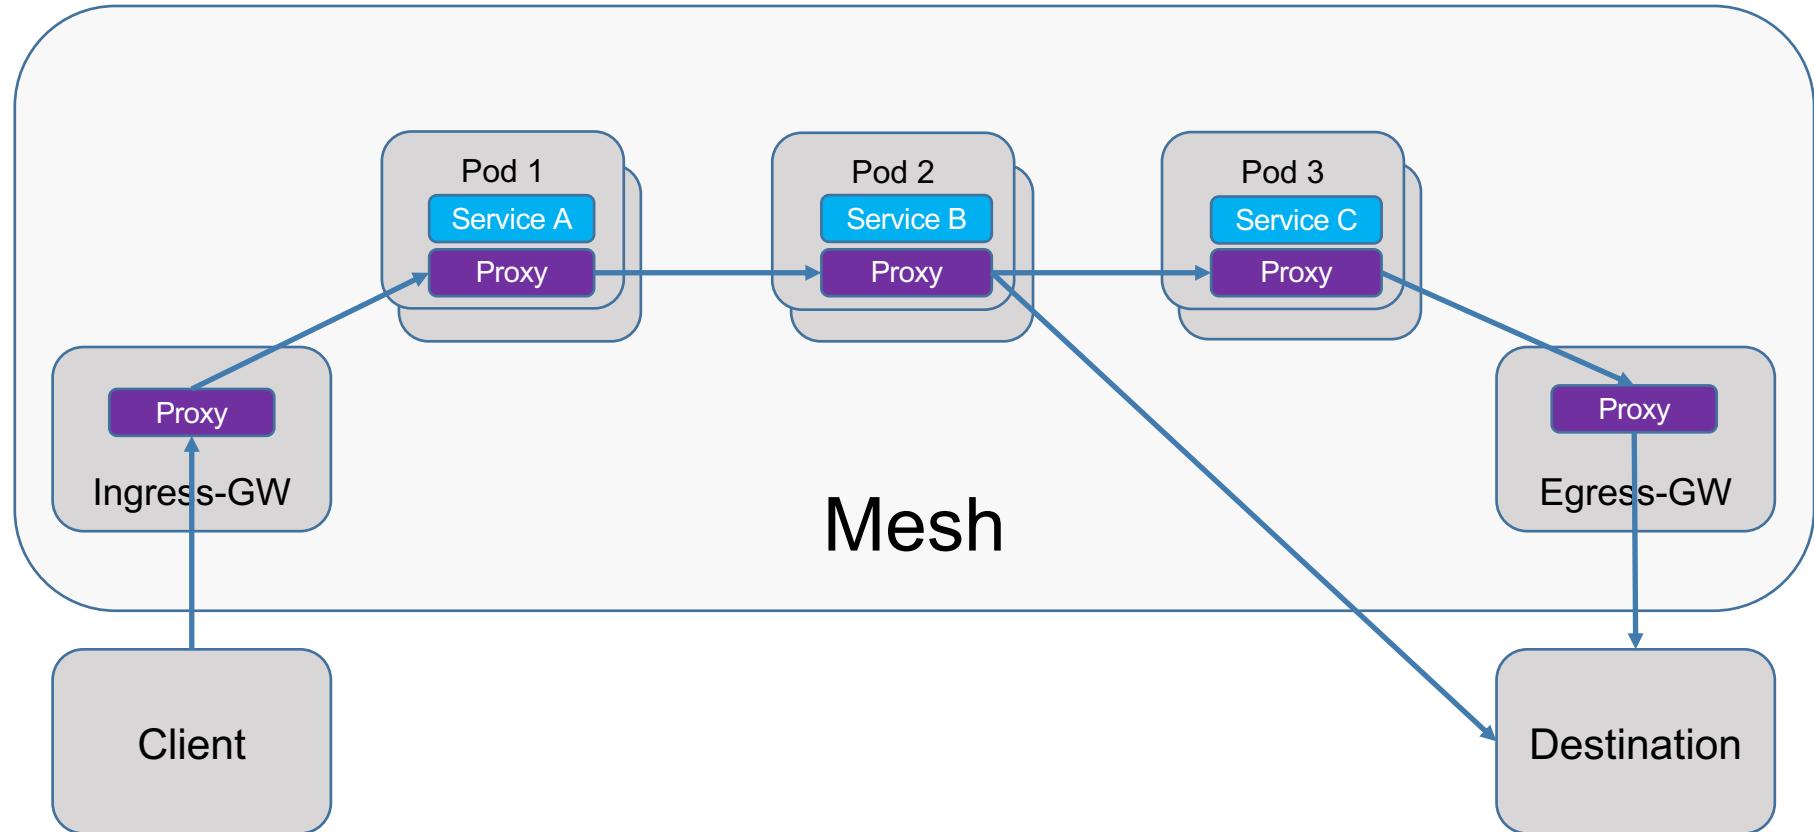
- Deal with the service mess of microservices
- Externalize required standard functionality
  - Telemetry
  - Request routing
  - Circuit breaking
  - Rate limiting
  - TLS/Authn/Authz
  - ...
- Declaratively configured
- Centrally maintained
- Language agnostic



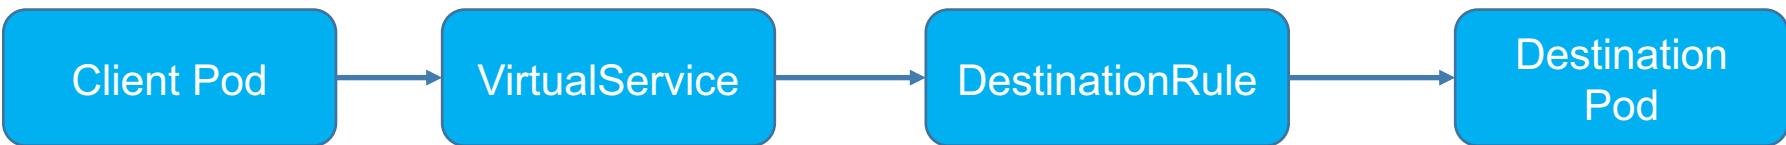
# General connectivity – in-mesh



# General connectivity – out-of-mesh



# Configuration Objects



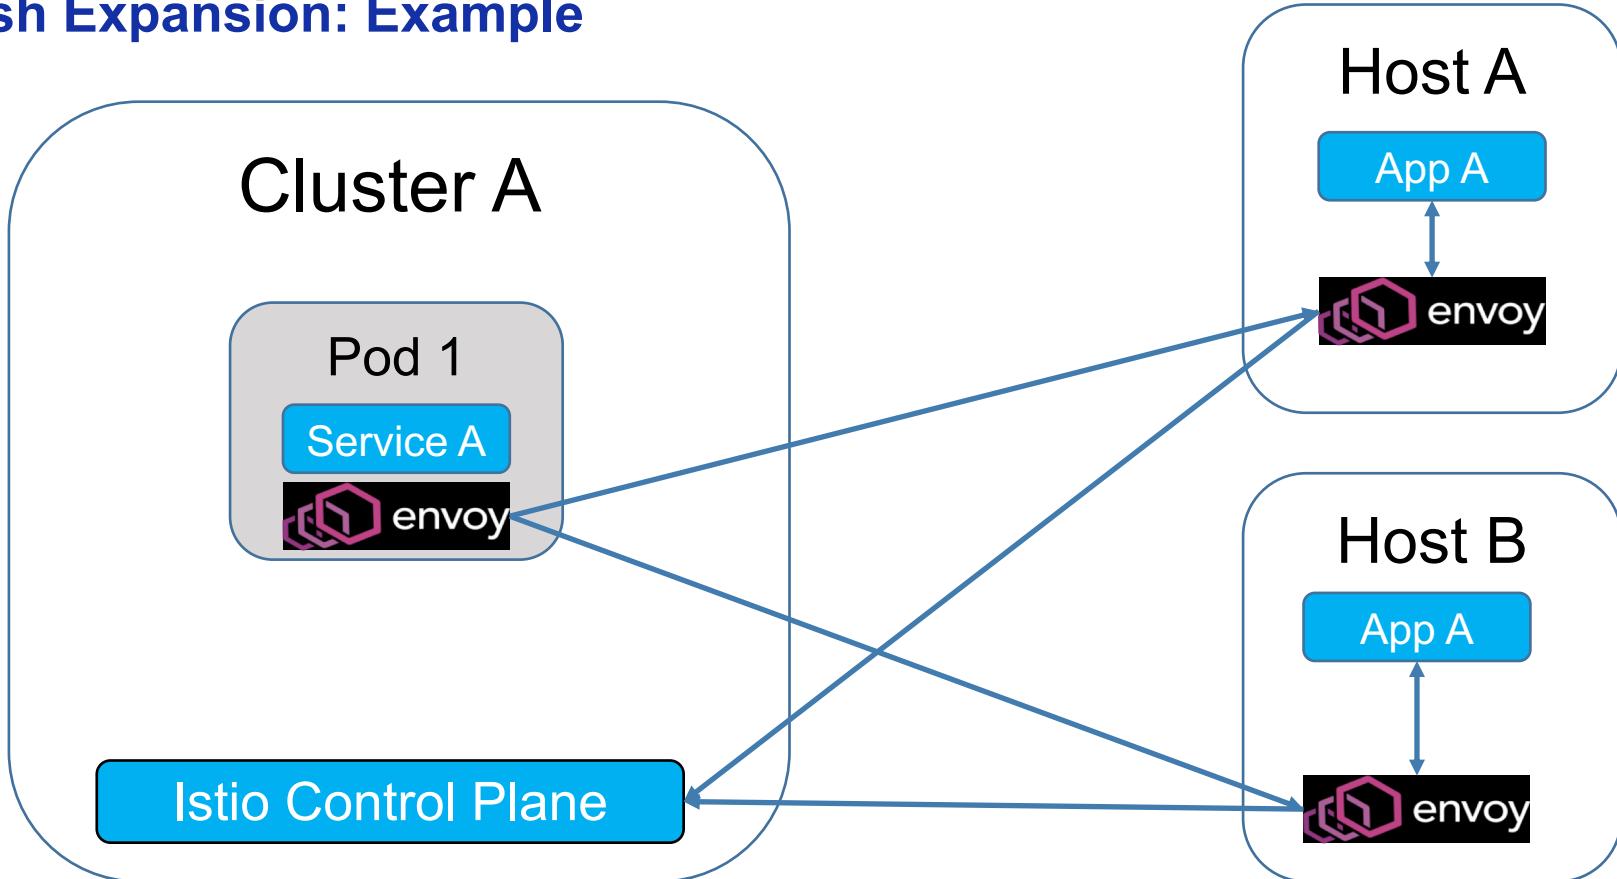
```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - route:
      - destination:
          host: reviews
          subset: v1
          weight: 75
      - destination:
          host: reviews
          subset: v2
          weight: 25
```

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  trafficPolicy:
    loadBalancer:
      simple: RANDOM
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
      trafficPolicy:
        loadBalancer:
          simple: ROUND_ROBIN
    - name: v3
      labels:
        version: v3
```

## Connecting to the „Old World“

- „Mesh expansion“ (based on Istio 1.0.x)
- Install Envoy on hosts and connect to Istio control plane
- mTLS: Universal transport encryption and authentication
  - Uniformly configured with Kubernetes resources
- Central observability by Istio telemetry

# Mesh Expansion: Example



# Problems with Operating Istio

- Security concerns
  - High privileges for control plane components
    - run as root
    - writable root filesystem
  - High privileges for admins and serviceaccount
    - net\_admin capabilities
    - run as root
  - Same problem with bookinfo sample application
- Sidecar injection
  - Problematic order of automatic sidecar injection vs. PSP evaluation



# Problems with Mesh Expansion

- Connectivity to control plane
  - Manual DNS tweaks via dnsmasq
  - Exposure of K8s-DNS required
- Telemetry / Policy
  - Mixer connections only to pod IP address
  - Additional tweaks to iptables required for PoC
- Inbound (in-mesh) calls
  - Would proxy to pod IP addresses, which are not routable
- mTLS setup

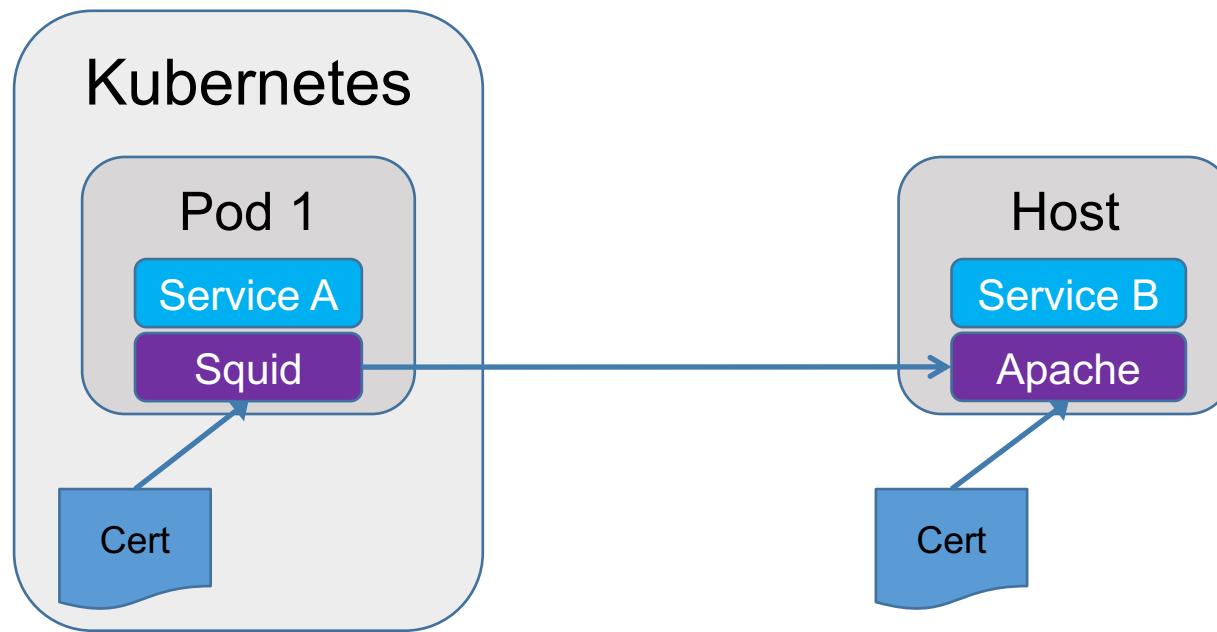
## Result for Istio 1.0

- Not suitable for production in our setup
  - Too many unstable tweaks necessary
  - Too much interference on expansion nodes
- Blocked migration of services to Kubernetes
  - Need for intermediate short term solution



## „Service Mesh Light“

- Mimic mTLS with custom proxy setup and static certificates
- Automatic sidecar injection as part of deployment pipeline



## Istio 1.1 to the Rescue?

- Control plane connection ✓
- Outbound expansion ✓
- mTLS setup ✓
- Istio-CNI / Security concerns ✓

But:

- Inbound expansion ✗
- Automatic sidecar injection ✗
- Documentation complex
- Documentation partially inconsistent
- Multi-tenancy unclear



## What's up next?

- Solving the remaining issues
  - Inbound mesh expansion
  - Large scale evaluation
  - Production readiness
    - Redundancy
    - Permissions
    - Scalability
  - Multi-tenancy questions
    - Responsibilities
  - Performance and overhead
    - Costs
    - Runtime Overhead

## What's up next?

- More benefits to be leveraged
  - Multi-cluster connectivity
  - SPIFFE interoperability
  - Locality based routing
  - Error scenarios
  - Authorization (request-based with JWT)
- Evaluate alternatives (e.g. Linkerd2)
- Provide upstream patches

## Summary

- Istio (>= 1.1) fits our needs
  - Feature set and K8s integration look fine
  - Not yet fully covering our requirements
- Hard to set up the right way
- Organisational questions need to be solved
- Direction of development is promising

