

KubeCon



CloudNativeCon

Europe 2019



KubeCon



CloudNativeCon

Europe 2019

What's the Performance Overhead?

Answering the Biggest Question in Tracing

Gabriela Soria

Former Outreachy intern for CNCF

About me - Outreachy



KubeCon



CloudNativeCon

Europe 2019

[Outreachy](#) provides three-month paid internships to work in Free and Open Source Software (FOSS) for under-represented people.



OPENTRACING



@gabrielasoriag

Agenda

- Introduction
- Benchmark tests
 - Scope and considerations
 - JMH Benchmark application example
 - Results
 - Conclusions
- Next steps

Introduction – Microservices



KubeCon



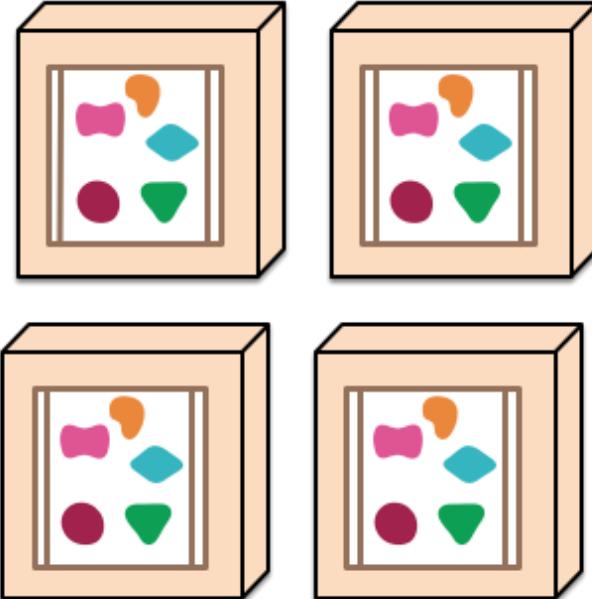
CloudNativeCon

Europe 2019

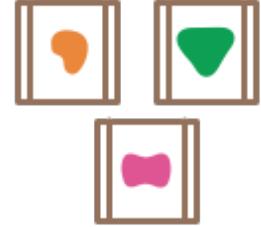
A monolithic application puts all its functionality into a single process...



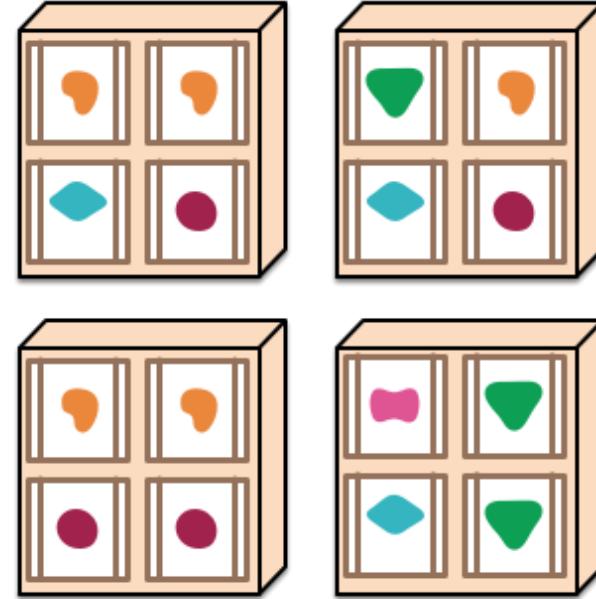
... and scales by replicating the monolith on multiple servers



A microservices architecture puts each element of functionality into a separate service...



... and scales by distributing these services across servers, replicating as needed.



Distributed tracing – OpenTracing



KubeCon



CloudNativeCon

Europe 2019

Distributed tracing profiles and monitor applications, especially those built using a microservices architecture.

OpenTracing is working towards creating more standardized APIs and instrumentation for distributed tracing.

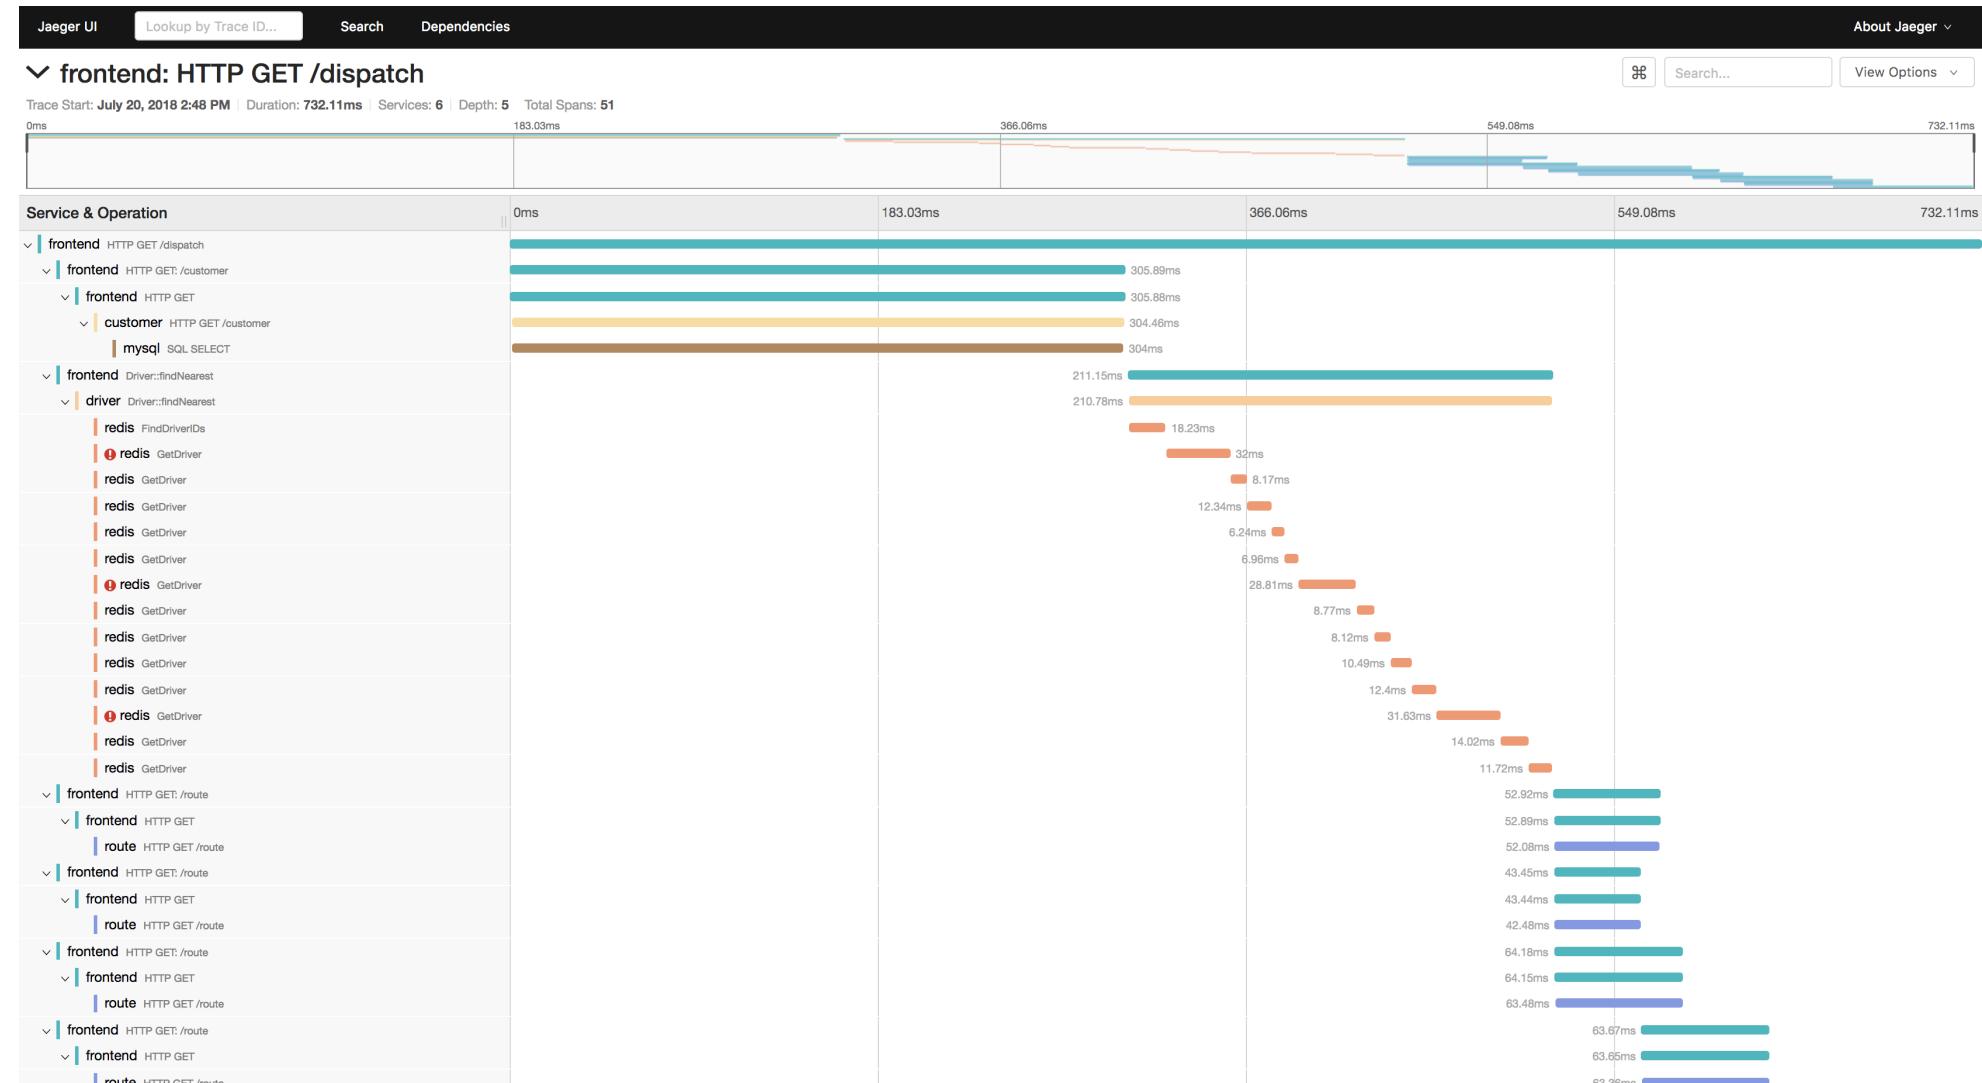
Jaeger



KubeCon

CloudNativeCon

Europe 2019



@gabrielasoriag

Source: <https://www.jaegertracing.io/docs/1.11/>

Scope and considerations



KubeCon



CloudNativeCon

Europe 2019

- We used Java Microbenchmark Harness (JMH)
- We used JMH Visualizer to present the results
- Tested Opentracing **Java** API only
- We tried to avoid network calls (in the scenarios that was possible)

JMH Benchmark Application Example



KubeCon



CloudNativeCon

Europe 2019

- Create a JMH application:

```
mvn archetype:generate \  
-DinteractiveMode=false \  
-DarchetypeGroupId=org.openjdk.jmh \  
-DarchetypeArtifactId=jmh-java-benchmark-archetype \  
-DgroupId=org.sample \  
-DartifactId=jmh-examples \  
-Dversion=1.0
```

JMH Benchmark Application Example



KubeCon



CloudNativeCon

Europe 2019

```
public class BenchmarkPetclinicSampleTime extends BenchmarkPetclinicBase {
```

```
    @Benchmark
    @BenchmarkMode(Mode.SampleTime)
    public Owner noInstrumentation(StateVariablesNoInstrumentation state) {
        return findPetOwnerById(state);
    }
```

```
    @Benchmark
    @BenchmarkMode(Mode.SampleTime)
    public Owner noopTracer(StateVariablesNoopTracer state) {
        return findPetOwnerById(state);
    }
```

```
    @Benchmark
    @BenchmarkMode(Mode.SampleTime)
    public Owner jaegerTracer(StateVariablesJaeger state) {
        return findPetOwnerById(state);
    }
```

....

@gabrielasoriag

JMH Benchmark Application Example



KubeCon



CloudNativeCon

Europe 2019

```
public class BenchmarkPetclinicBase {  
  
    public Owner findPetOwnerById(StateVariables state) {  
        return state.petcontroller.findOwner(1);  
    }  
    @State(Scope.Benchmark)  
    public static class StateVariables {  
        public PetController petcontroller;  
        public ConfigurableApplicationContext c;  
  
        @TearDown(Level.Iteration)  
        public void shutdownContext() {  
            c.close();  
        }  
  
        public void initApplication() {  
            c = SpringApplication.run(PetClinicApplication.class);  
            petcontroller = c.getBean(PetController.class);  
        }  
    }
```

@gabrielasoriag

JMH Benchmark Application Example



KubeCon



CloudNativeCon

Europe 2019

```
public static class StateVariablesNoInstrumentation extends StateVariables {
    @Setup(Level.Iteration)
    public void doSetup() {
        System.setProperty("tracerresolver.disabled", Boolean.TRUE.toString());
        initApplication();
    }
}

public static class StateVariablesJaeger extends StateVariables {
    @Setup(Level.Iteration)
    public void doSetup() {

        System.setProperty(AbstractEnvironment.ACTIVE_PROFILES_PROPERTY_NAME,
TracerImplementation.JAEGERTRACER);
        initApplication();
    }
}
```

JMH Benchmark Application Example



KubeCon



CloudNativeCon

Europe 2019

- Run the tests:

```
mvn clean install  
java -jar target/benchmarks.jar
```

▼ opentracing-benchmark-spring-cloud	Yesterday at 01:45
► .idea	Today at 00:15
► results-md	Yesterday at 02:55
↳ README.md	Yesterday at 01:45
► results-imgs	Yesterday at 01:39
▼ results	Yesterday at 01:31
↳ jmh-2019-05-03-00-08-35.json	May 3, 2019 at 09:25
↳ jmh-2019-05-03-01-05-59.json	May 3, 2019 at 09:25
↳ jmh-2019-05-03-00-37-27.json	May 3, 2019 at 09:25

JMH Benchmark Application Example



KubeCon



CloudNativeCon

Europe 2019

- Visualize the results

<http://jmh.morethan.io/>

JMH Visualizer

Dropzone

Drop your JMH JSON report file(s) here!

"JMH is a Java harness for building, running, and analysing nano/micro/milli/macro benchmarks written in Java and other languages targeting the JVM."

@gabrielasoriag

List of tests



KubeCon



CloudNativeCon

Europe 2019

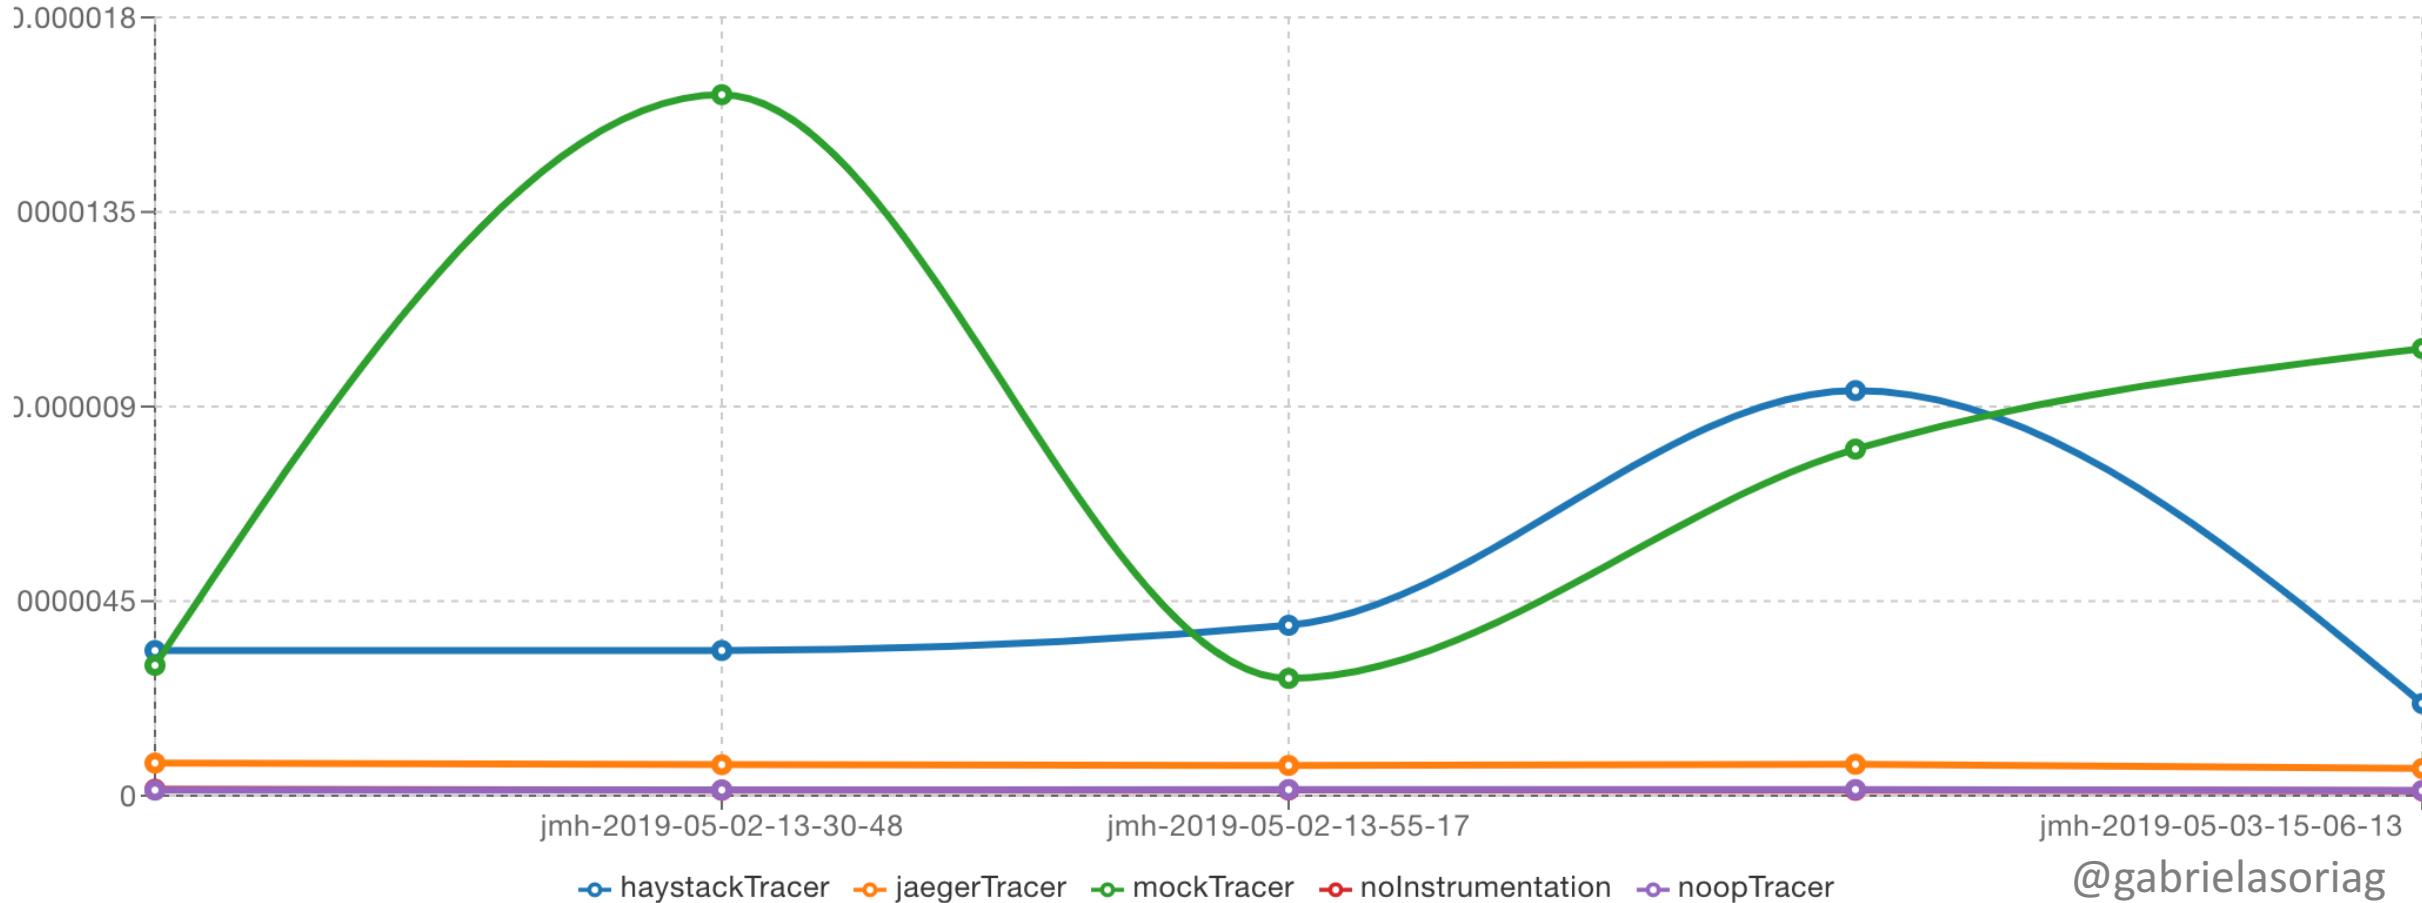
Scenario	Dependency	Description	HTTP calls	Framework
Simple Java	opentracing-java (manually)	String concatenation	No	No
Spring Boot	opentracing-java (manually)	Billing process	No	Yes
Spring Cloud	opentracing-spring-cloud-starter	Petclinic	No	Yes
JDBC (in memory database)	opentracing-jdbc	Spring Boot - Course Management	No	Yes
Servlet Filter	opentracing-web-servlet-filter	Simple servlet – Hello World Page	Yes	Yes
JAX-RS	opentracing-jaxrs2-discovery	Spring Boot - Course Management	Yes	Yes

Results – Simple Java (Sample time)



CloudNativeCon
Europe 2019

BenchmarkStringConcatenationSampleTime Sampling Time | Q | ⚖️



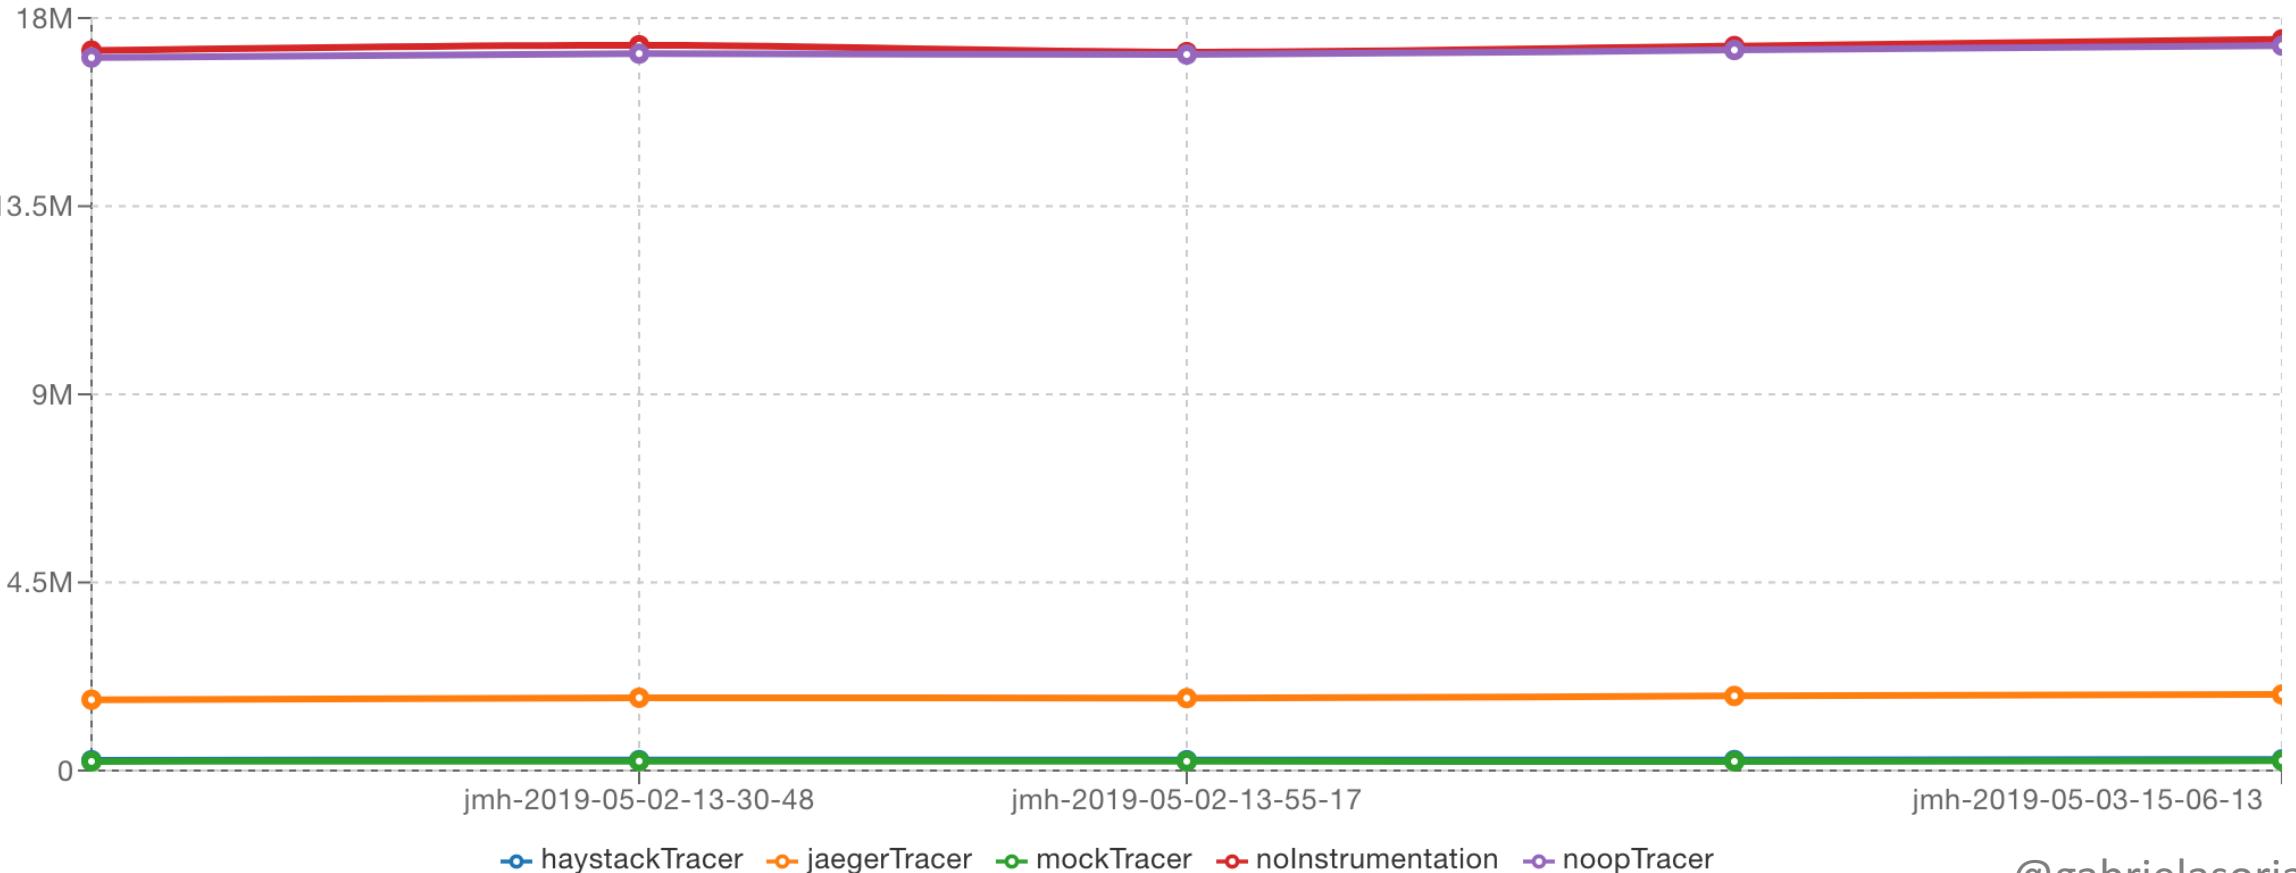
@gabrielasoriag

Results – Simple Java (Throughput)



CloudNativeCon
Europe 2019

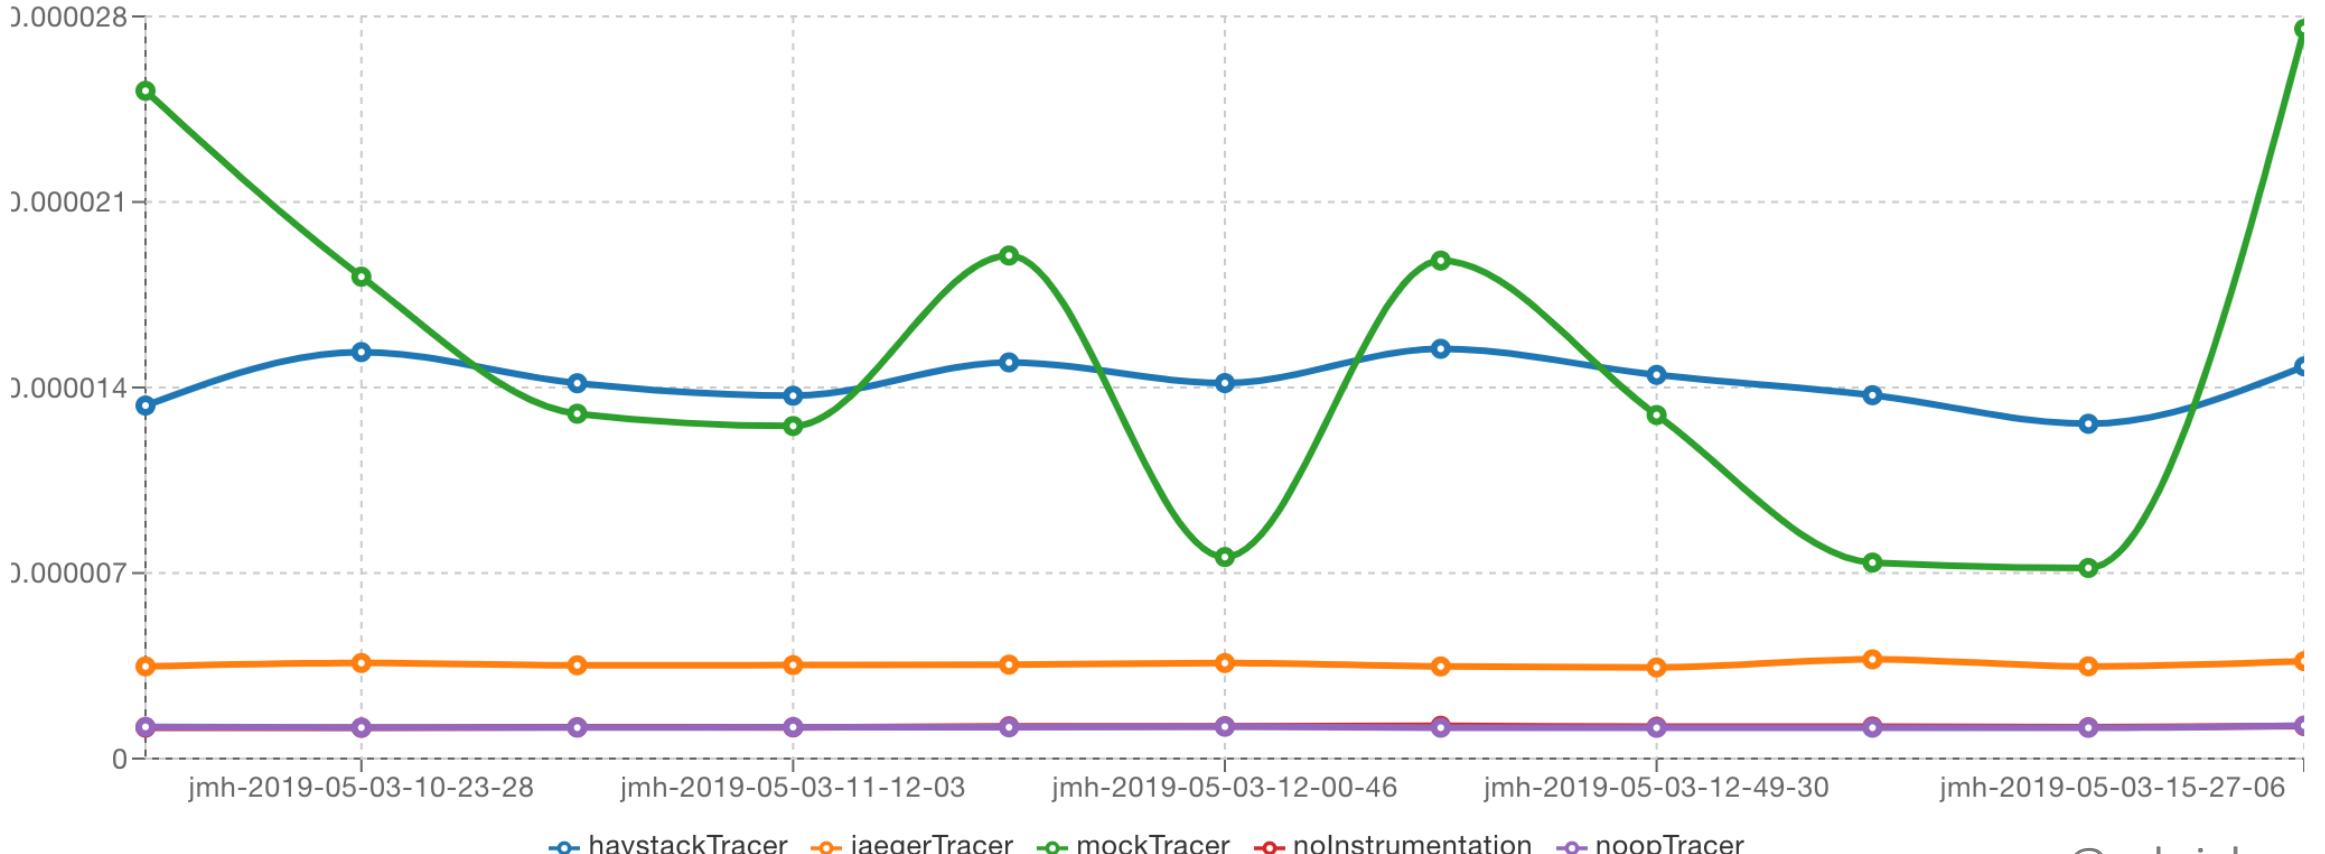
BenchmarkStringConcatenationThroughput Throughput | Q | Δ



@gabrielasoriag

Results – Spring boot (Sample time)

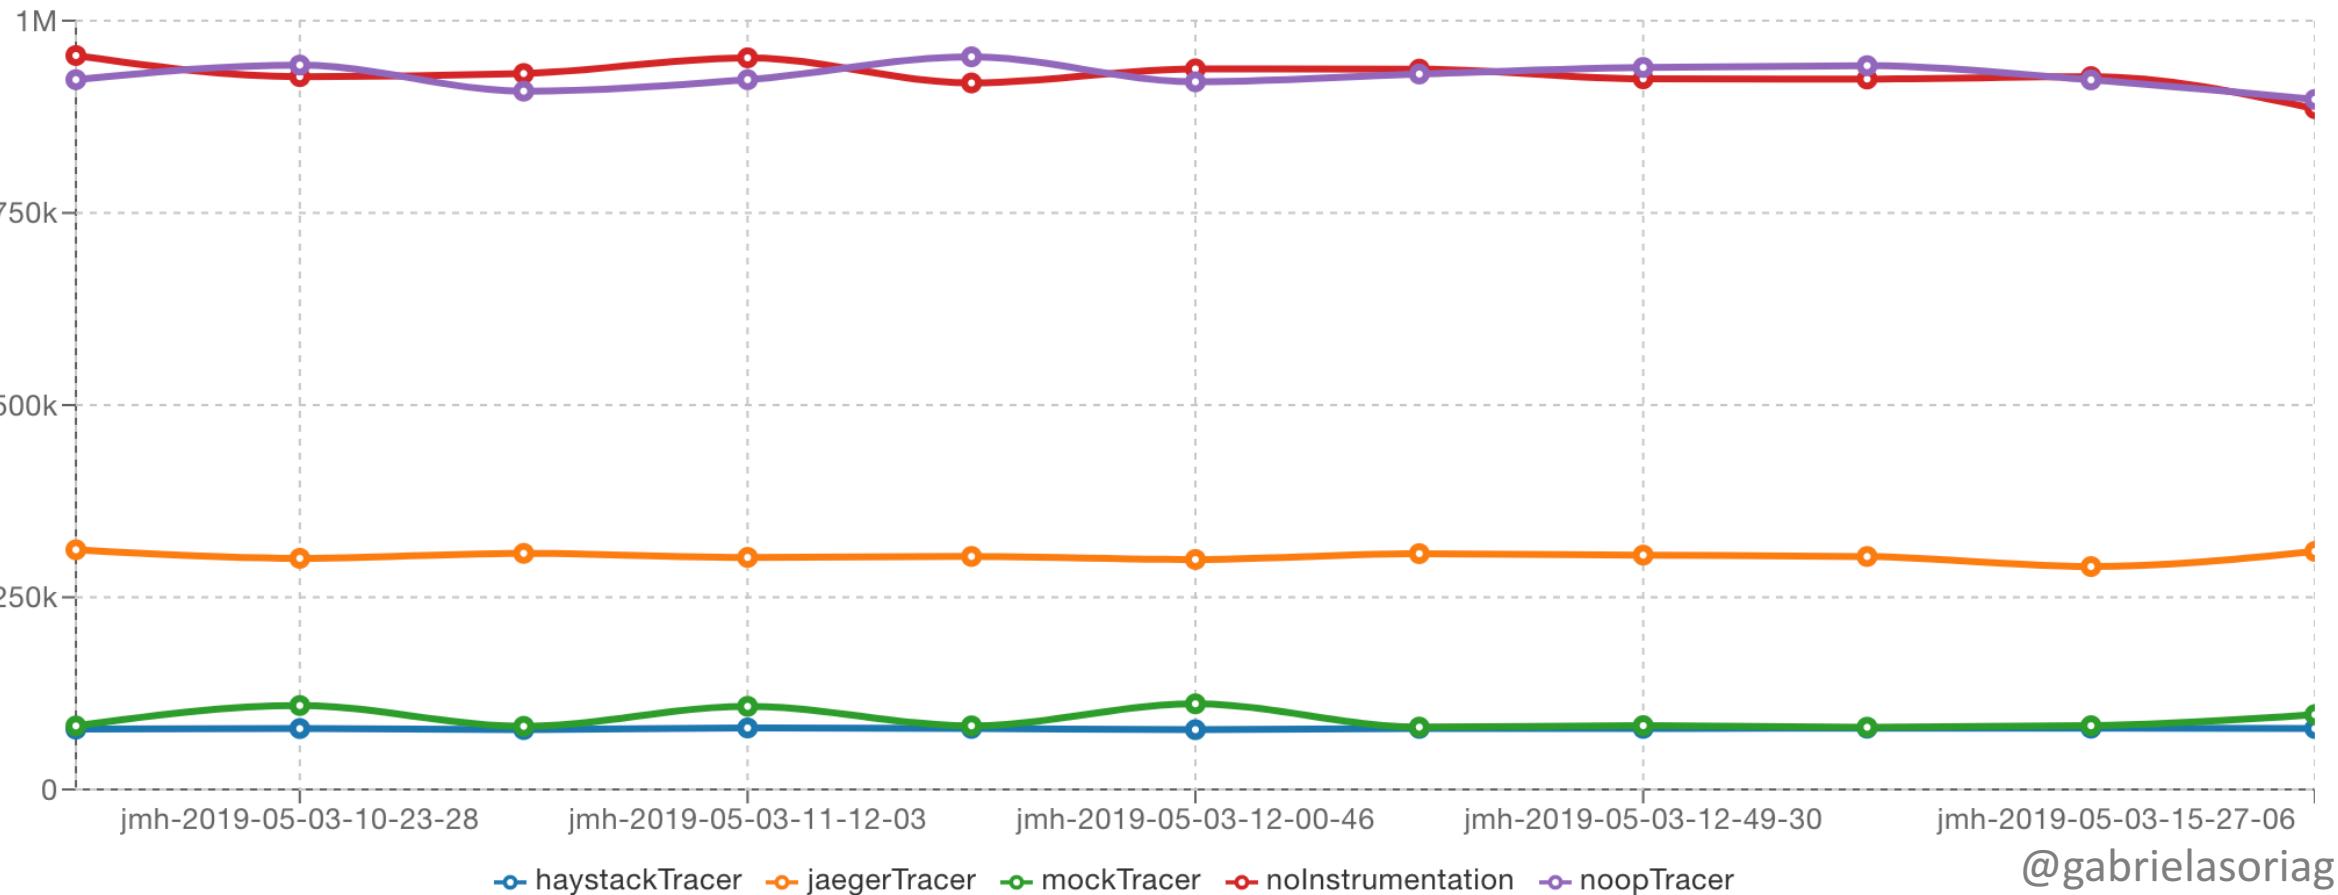
BenchmarkBillingSampleTime | Sampling Time | Q | ▾



@gabrielasoriag

Results – Spring boot (Throughput)

Benchmark Billing Throughput Throughput | + | ⚖️



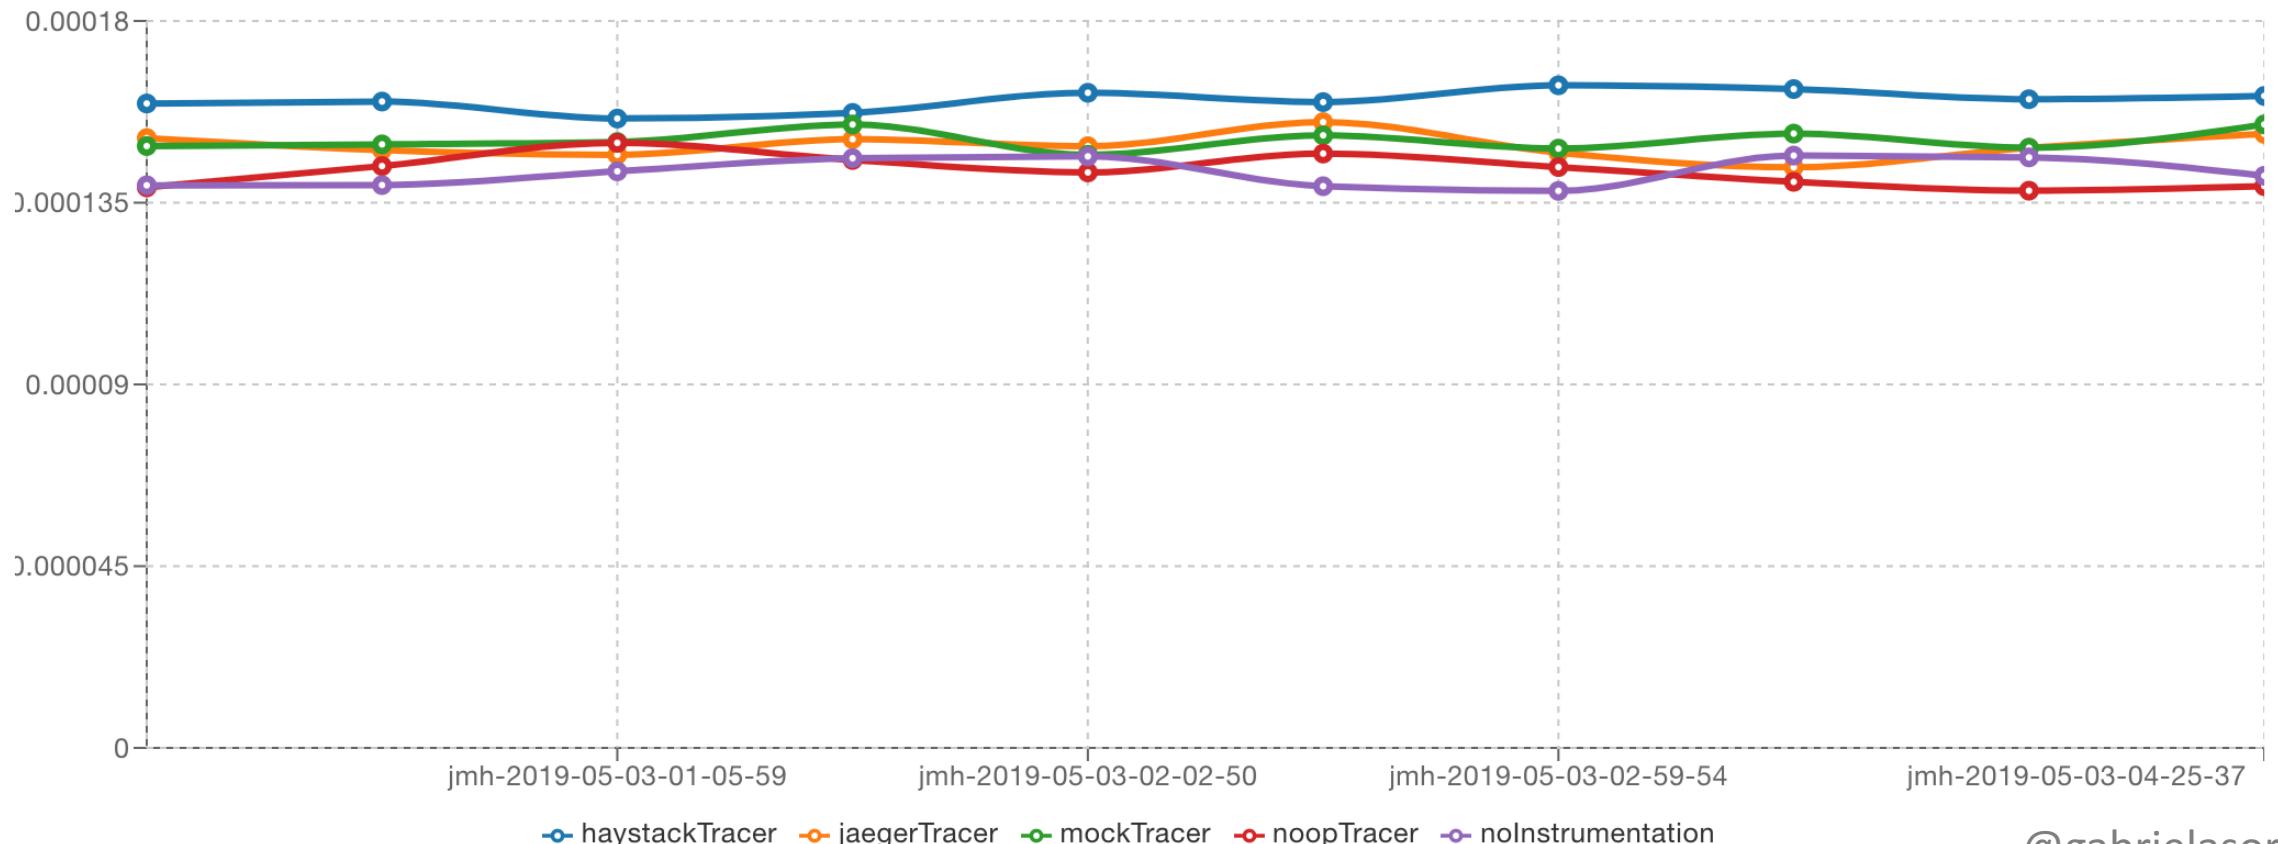
@gabrielasoriag

Results – Spring Cloud (Sample time)



CloudNativeCon
Europe 2019

BenchmarkPetclinicSampleTime | Sampling Time | + | ⚖



@gabrielasoriag

Results – Spring Cloud (Throughput)



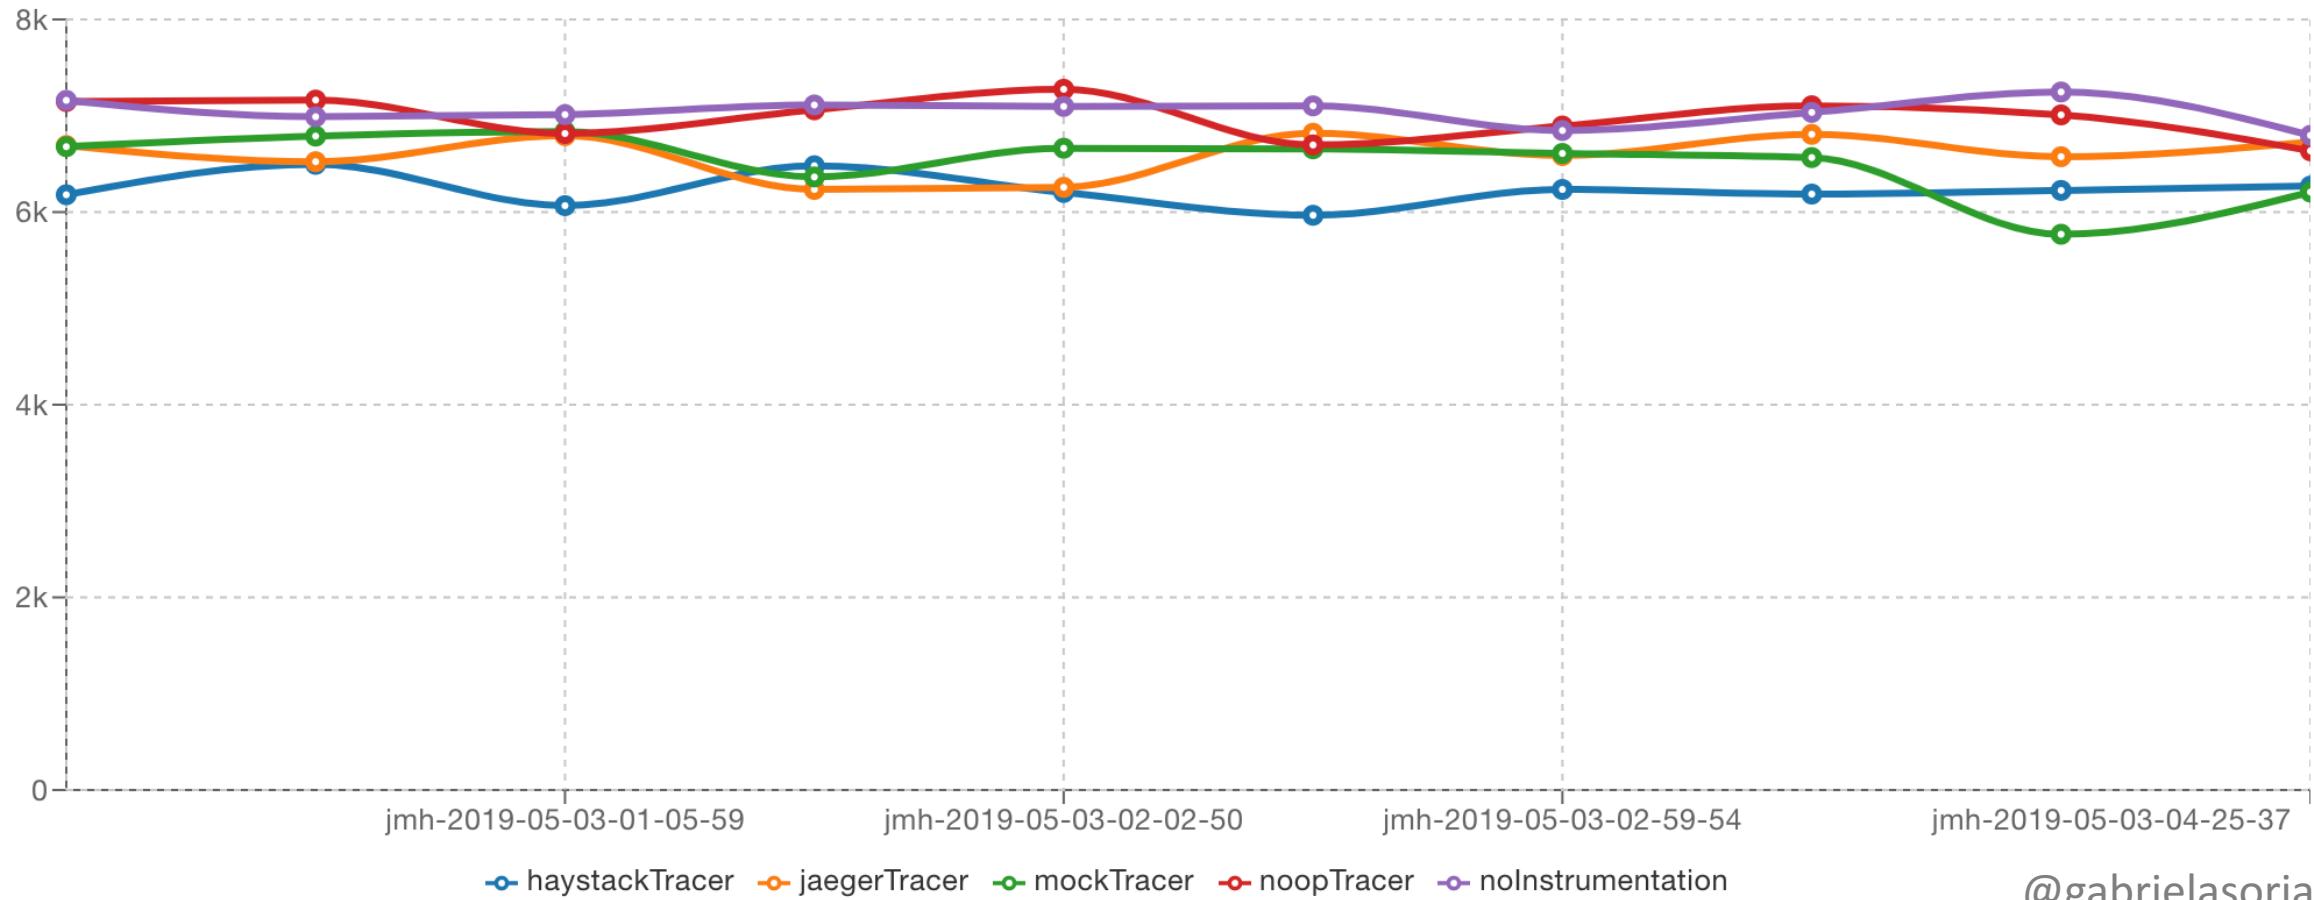
KubeCon

CloudNativeCon

Europe 2019

Benchmark Petclinic Throughput

Throughput



@gabrielasoriag

Results – JDBC (Sample time)

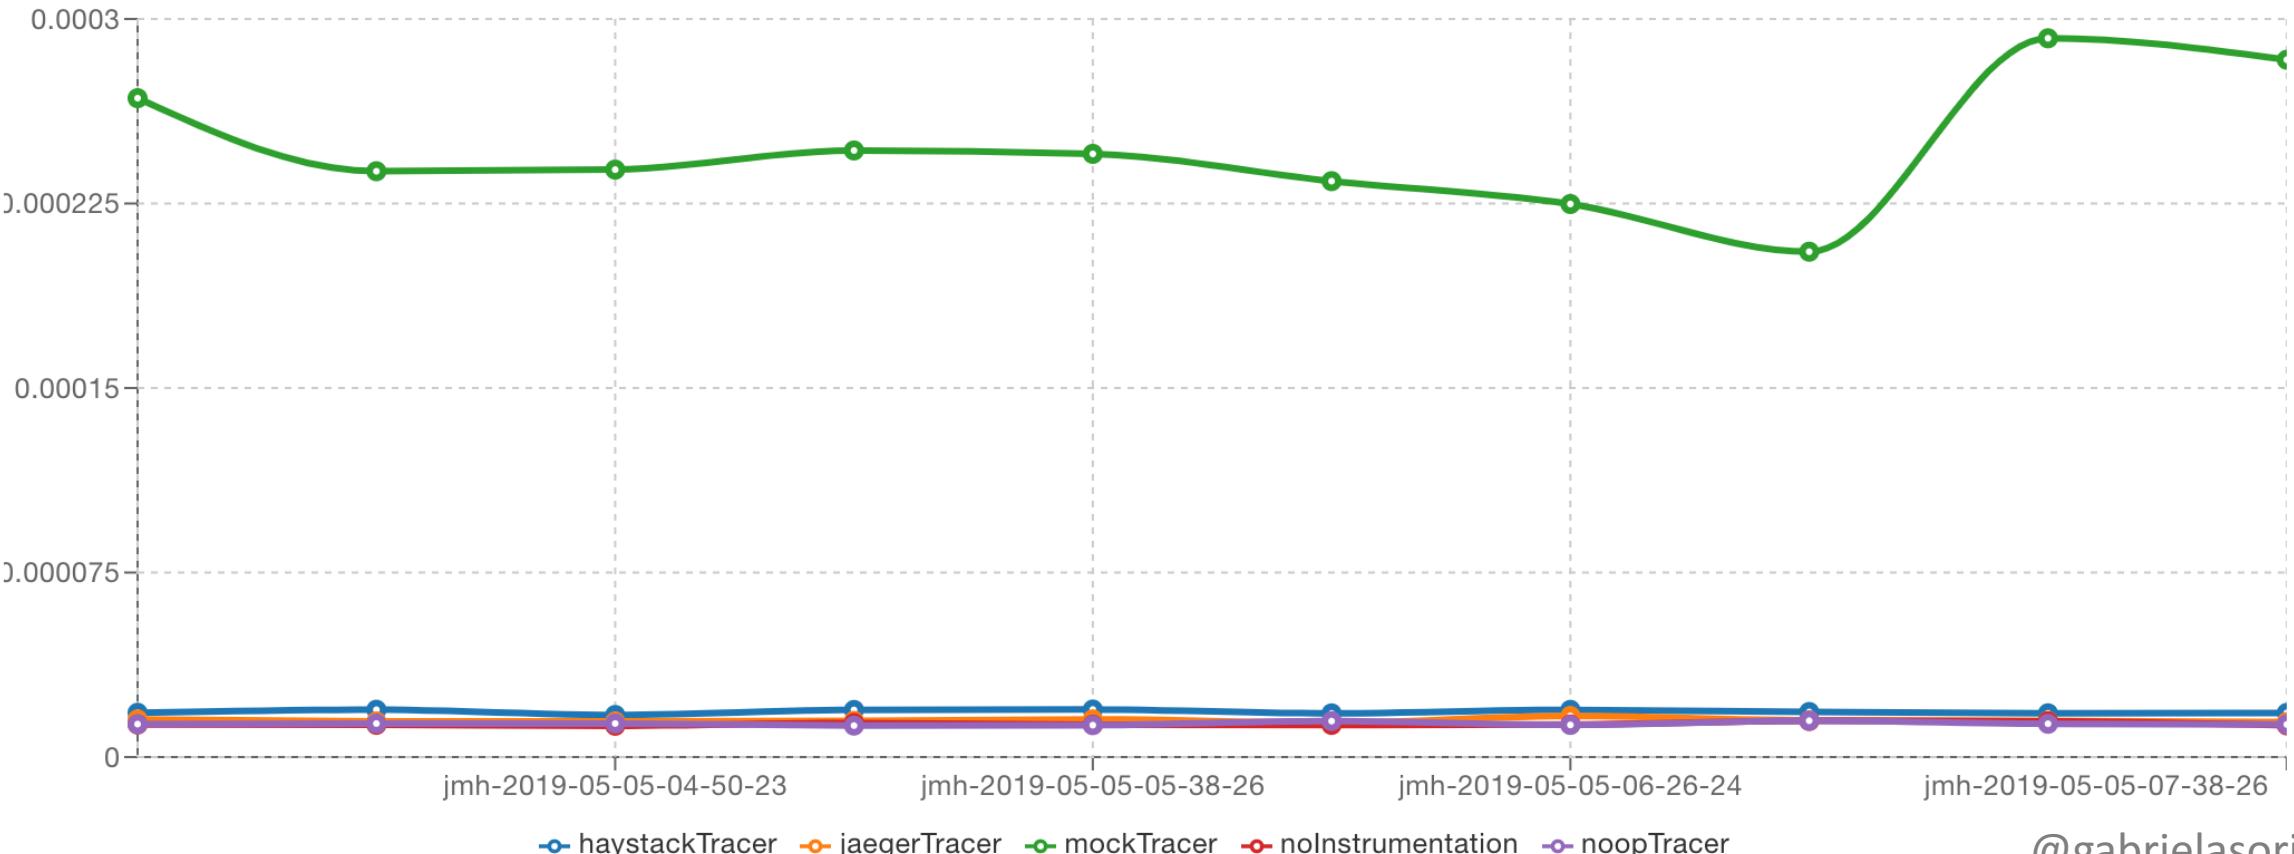


KubeCon

CloudNativeCon

Europe 2019

BenchmarkCourseManagementSampleTime | Sampling Time | |



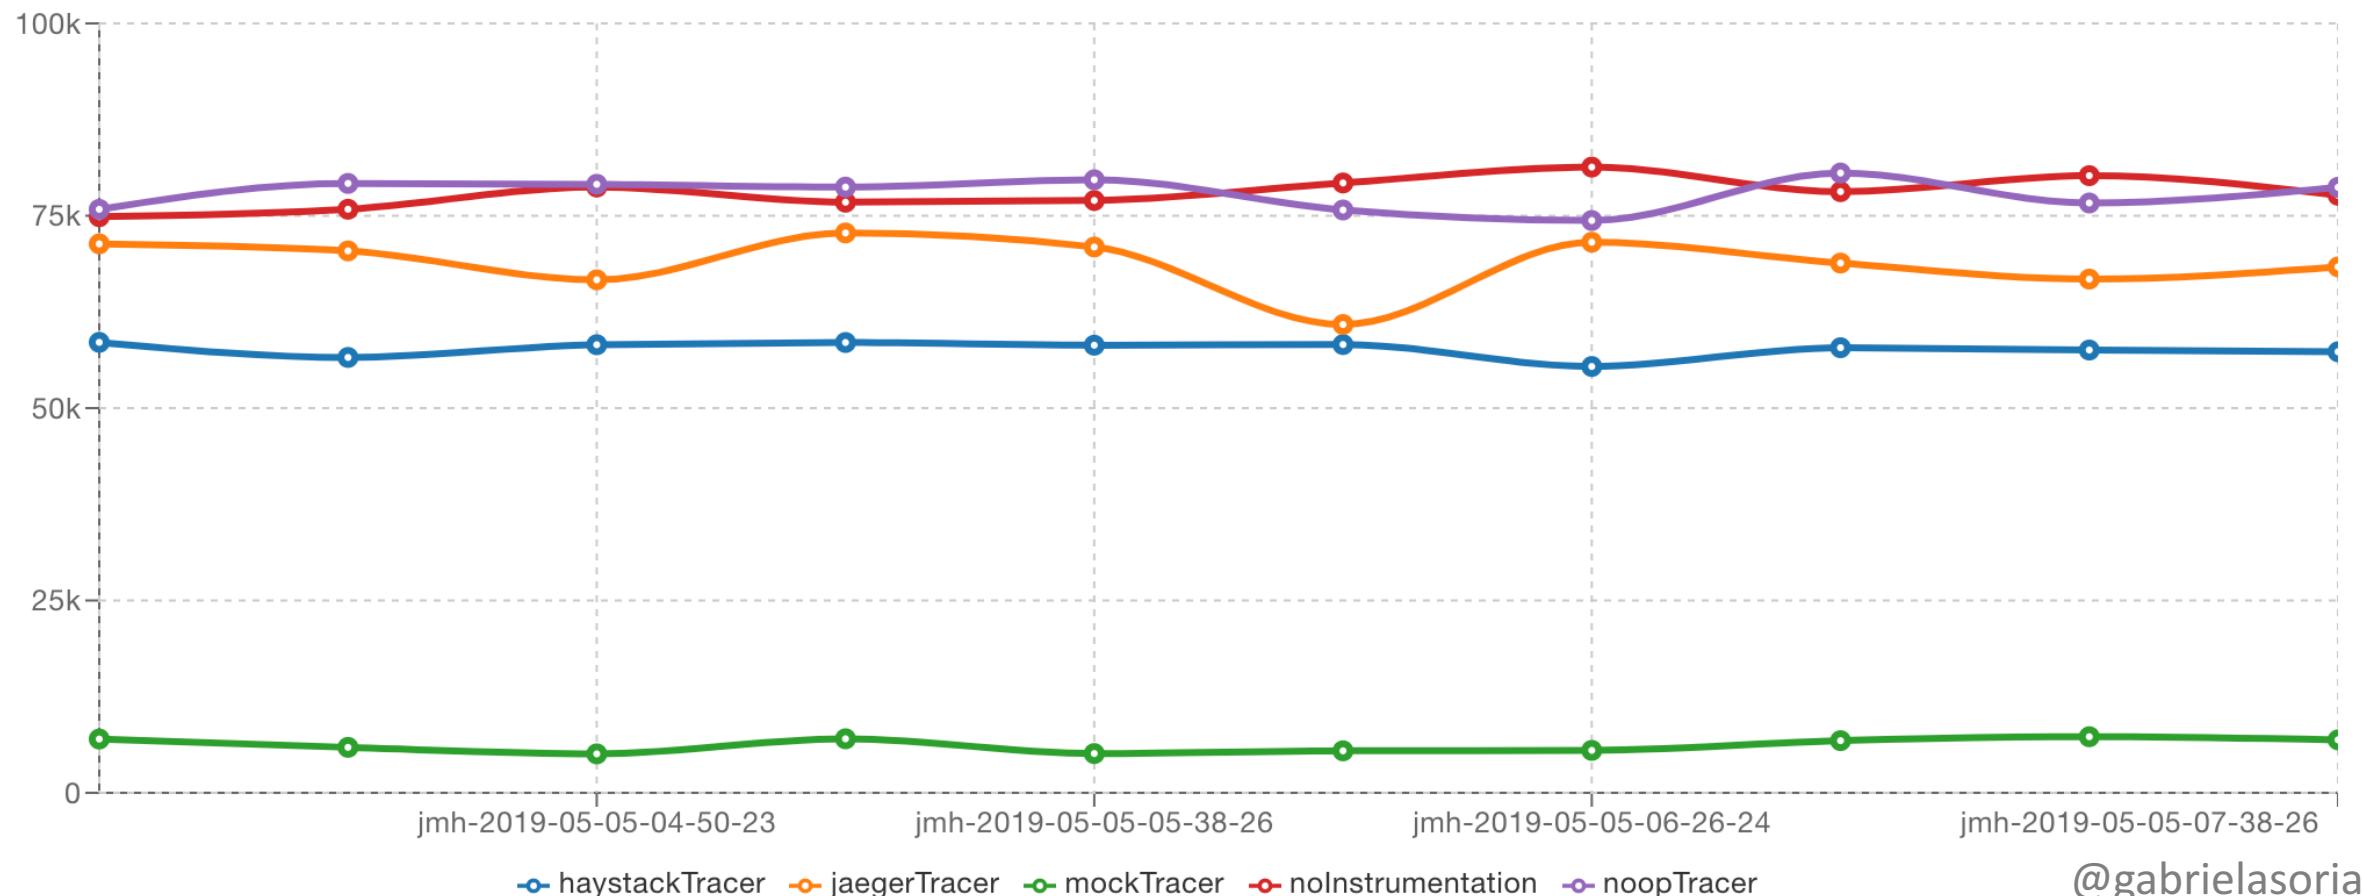
@gabrielasoriag

Results – JDBC (Throughput)



CloudNativeCon
Europe 2019

BenchmarkCourseManagementThroughput Throughput | Q | ⚖️



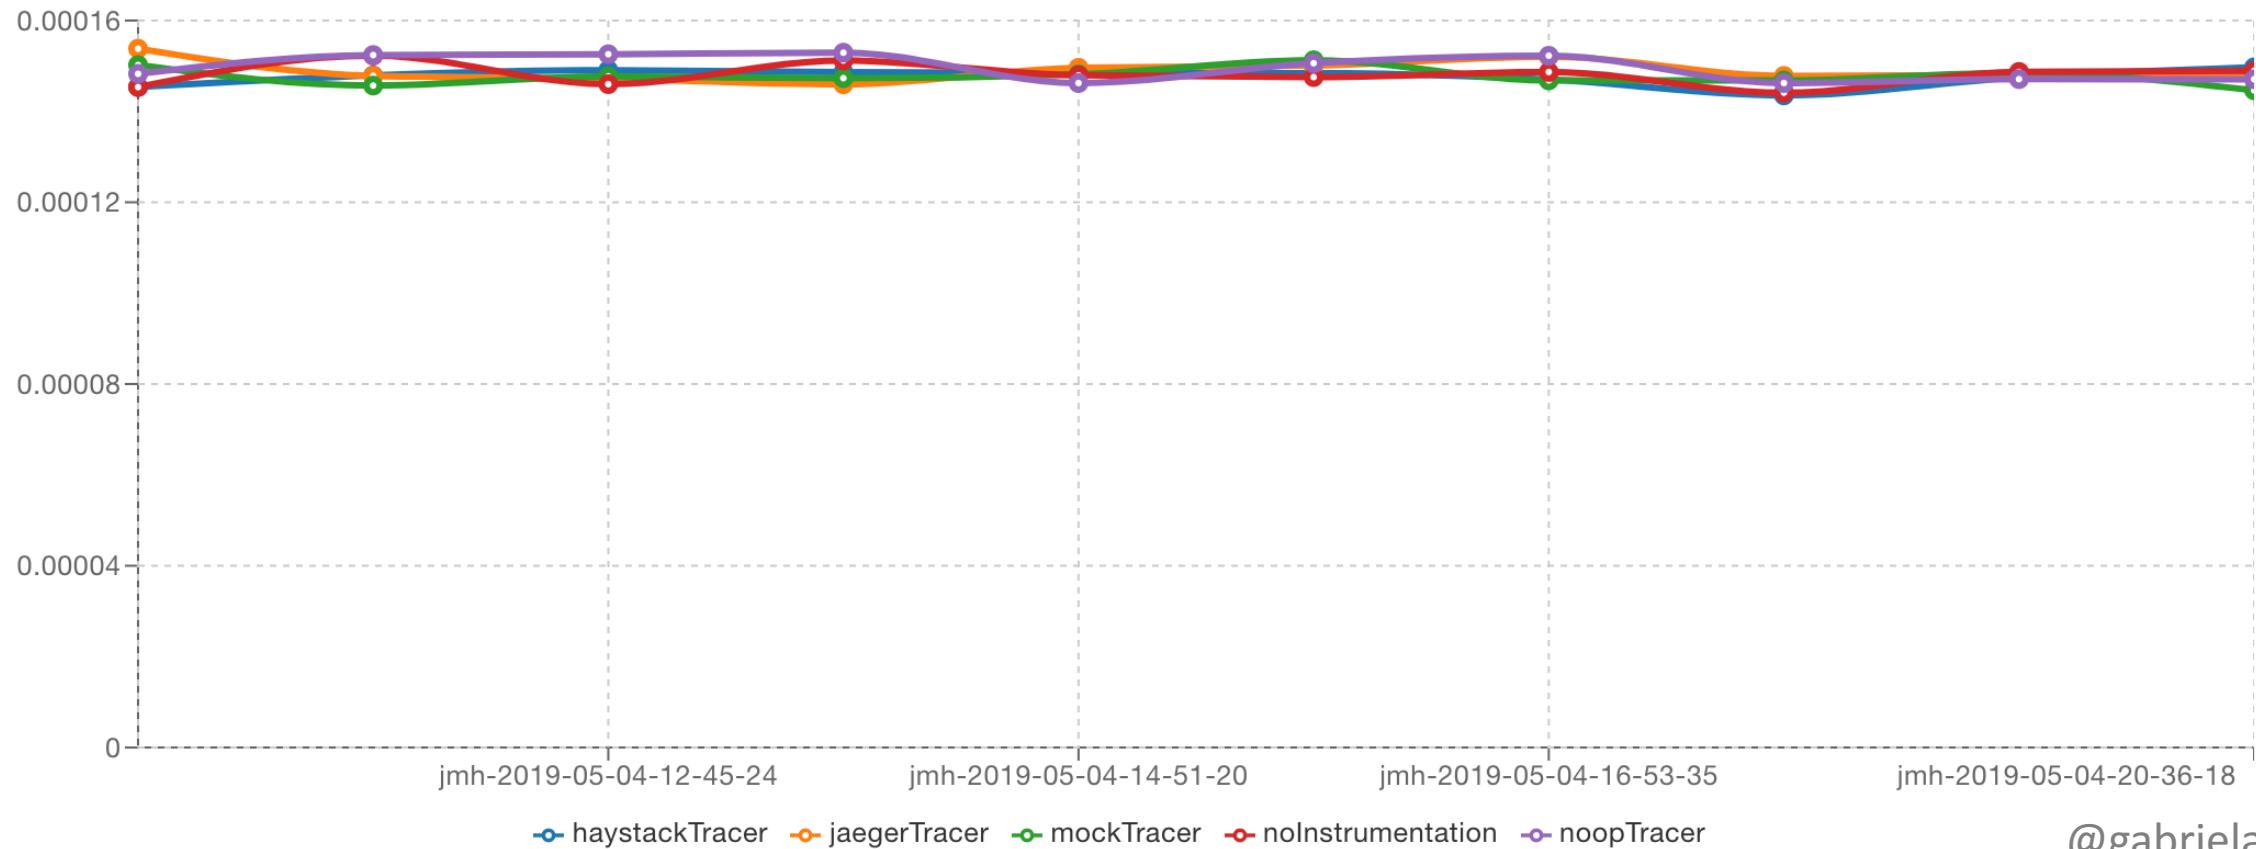
@gabrielasoriag

Results – Servlet Filter (Sample time)



CloudNativeCon
Europe 2019

BenchmarkSimpleServletSampleTime Sampling Time | Q | ⚖️



@gabrielasoriag

Results – Servlet Filter (Throughput)

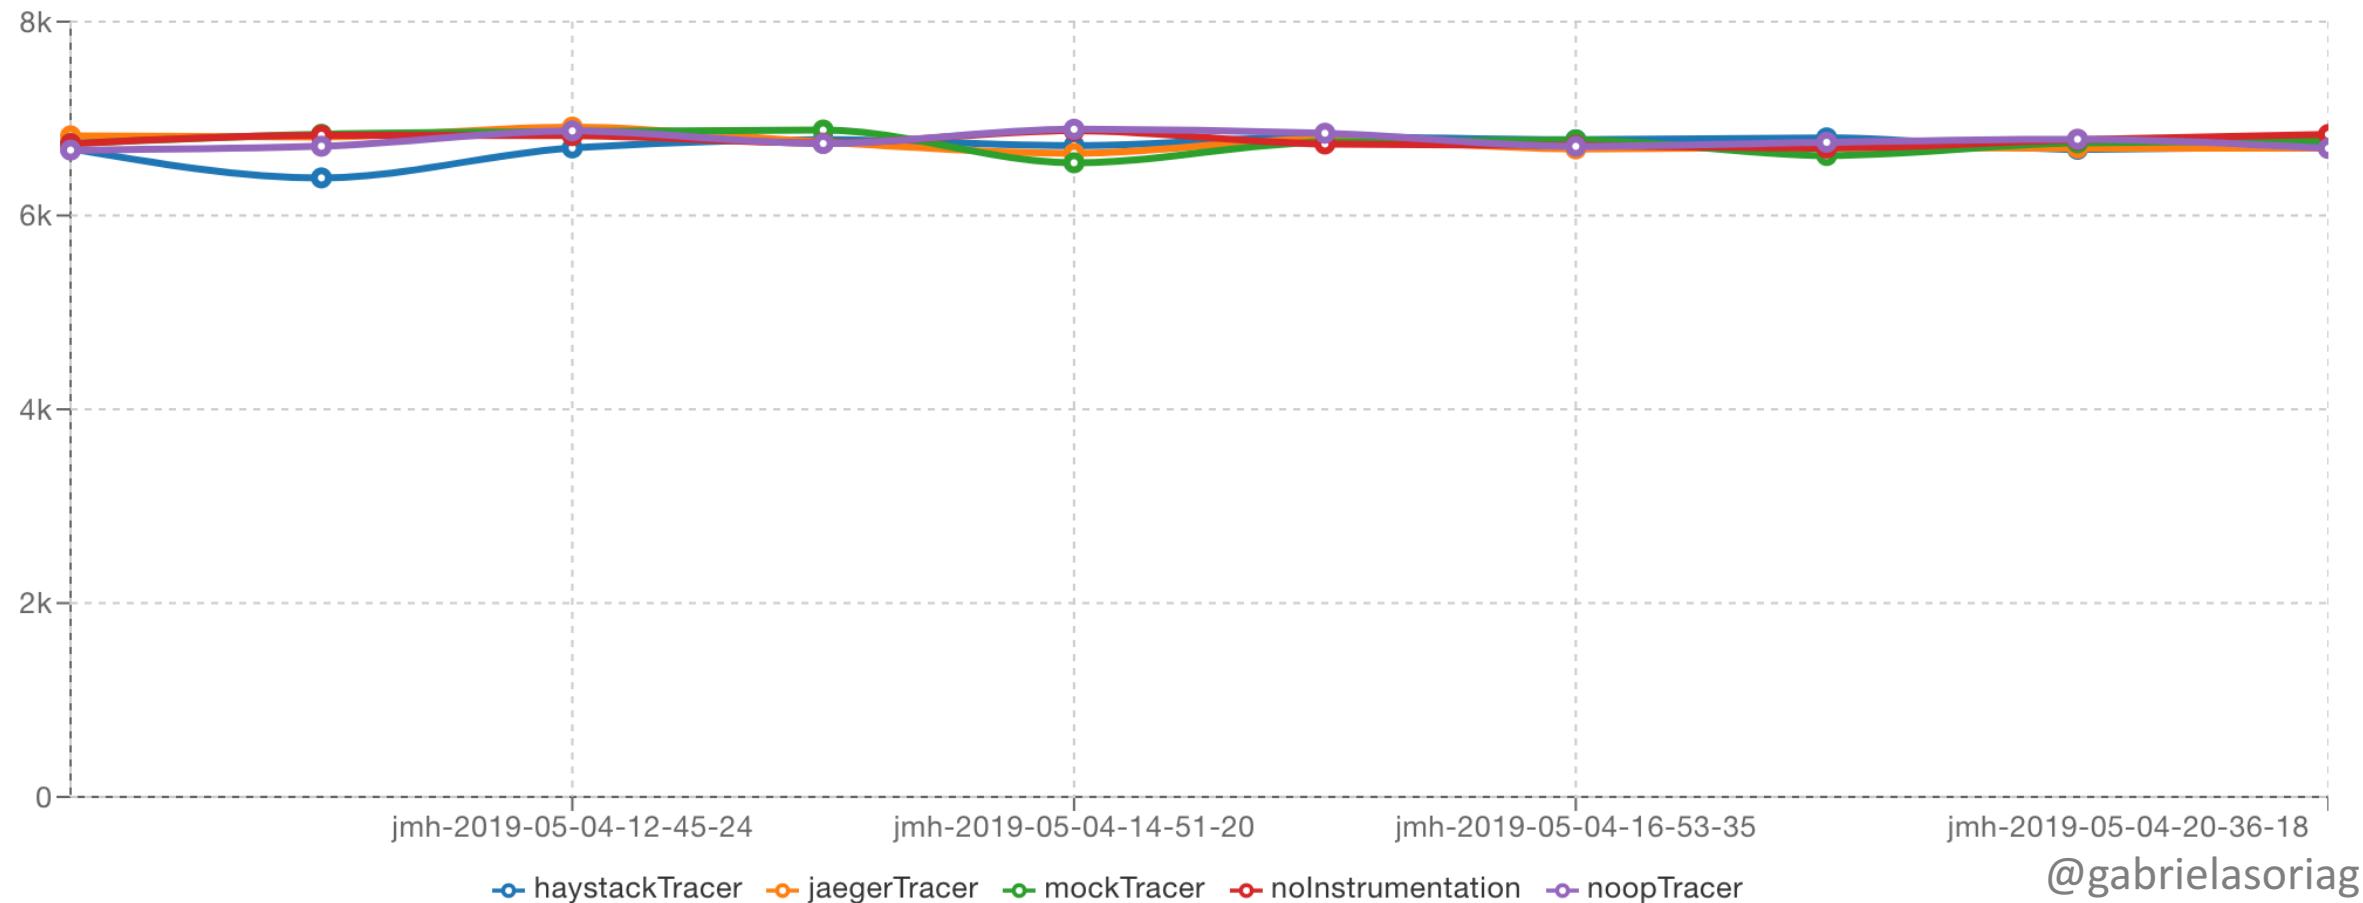


KubeCon

CloudNativeCon

Europe 2019

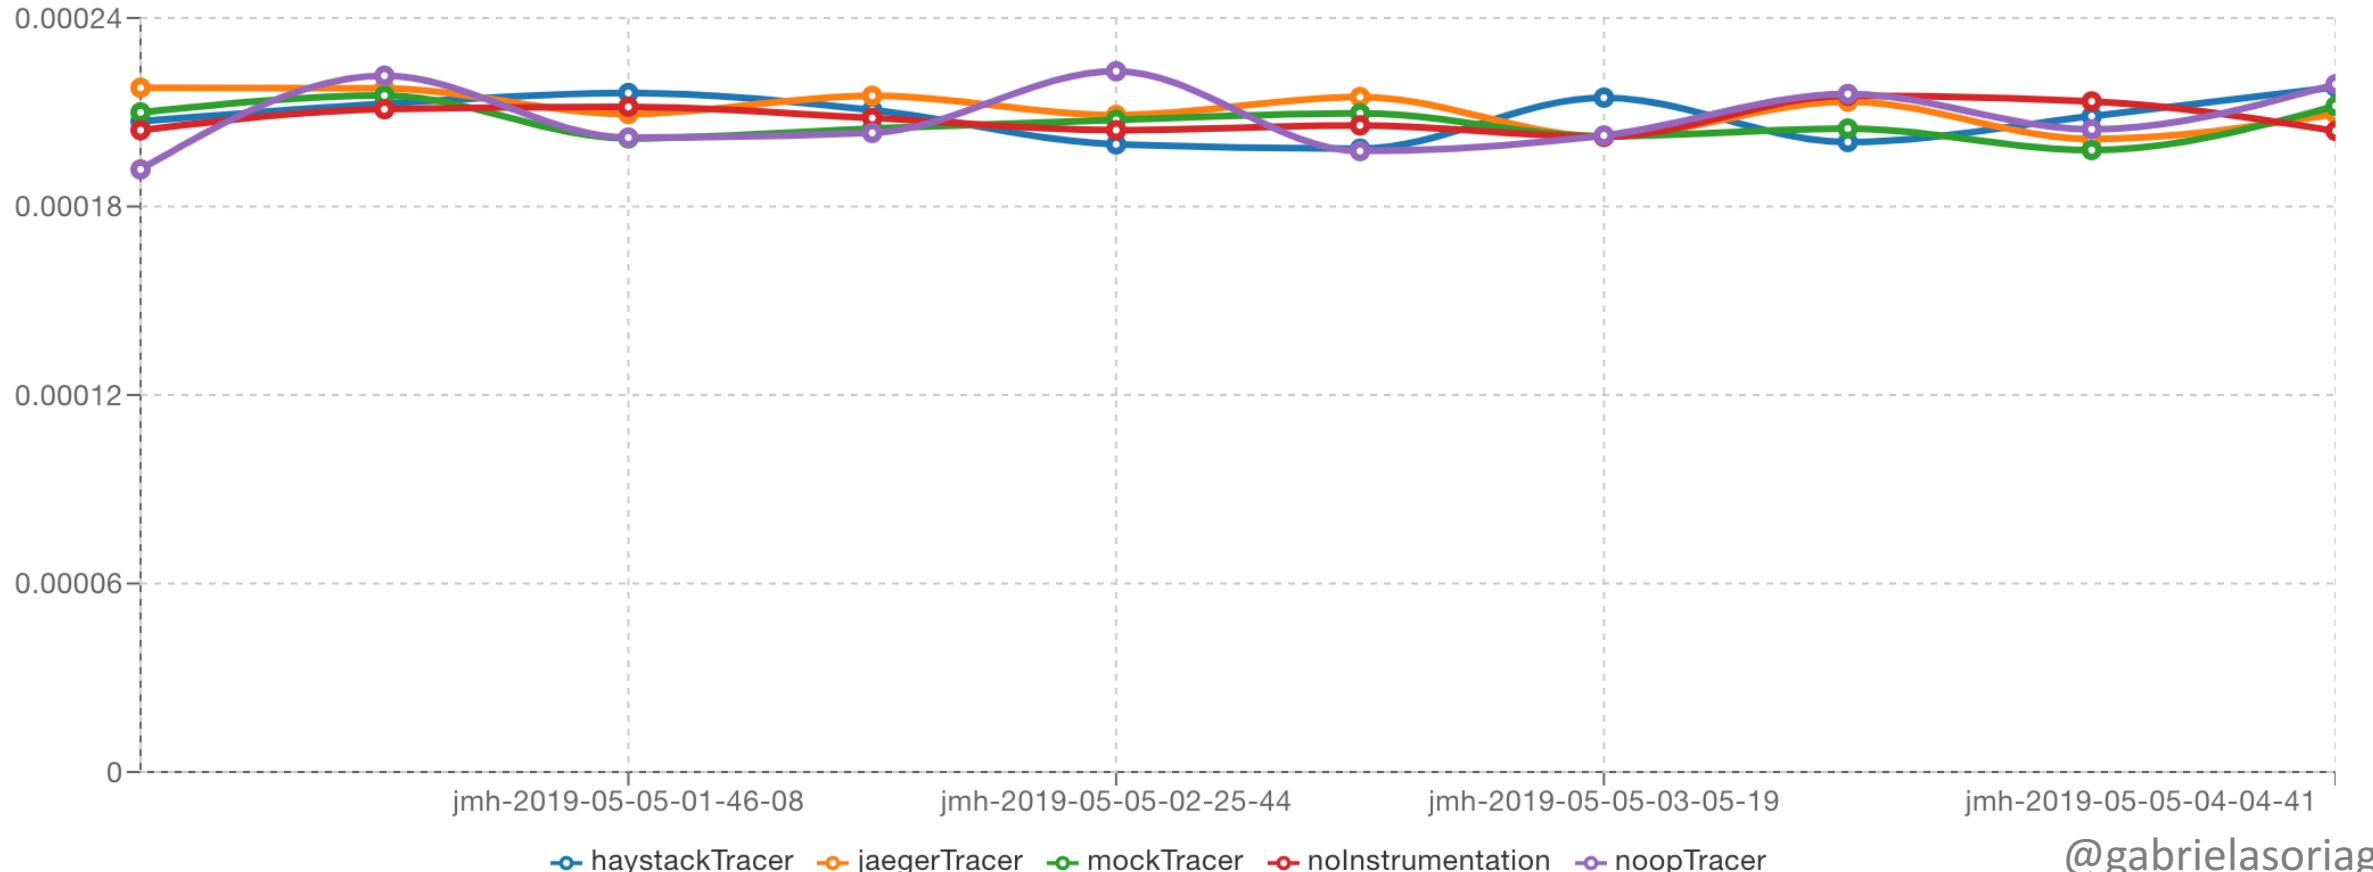
BenchmarkSimpleServletThroughput Throughput | Q | ⚖



@gabrielasoriag

Results – JAX-RS (Sample time)

BenchmarkCourseManagementSampleTime Sampling Time | 🔍 | ⚖️



@gabrielasoriag

Results – JAX-RS (Throughput)

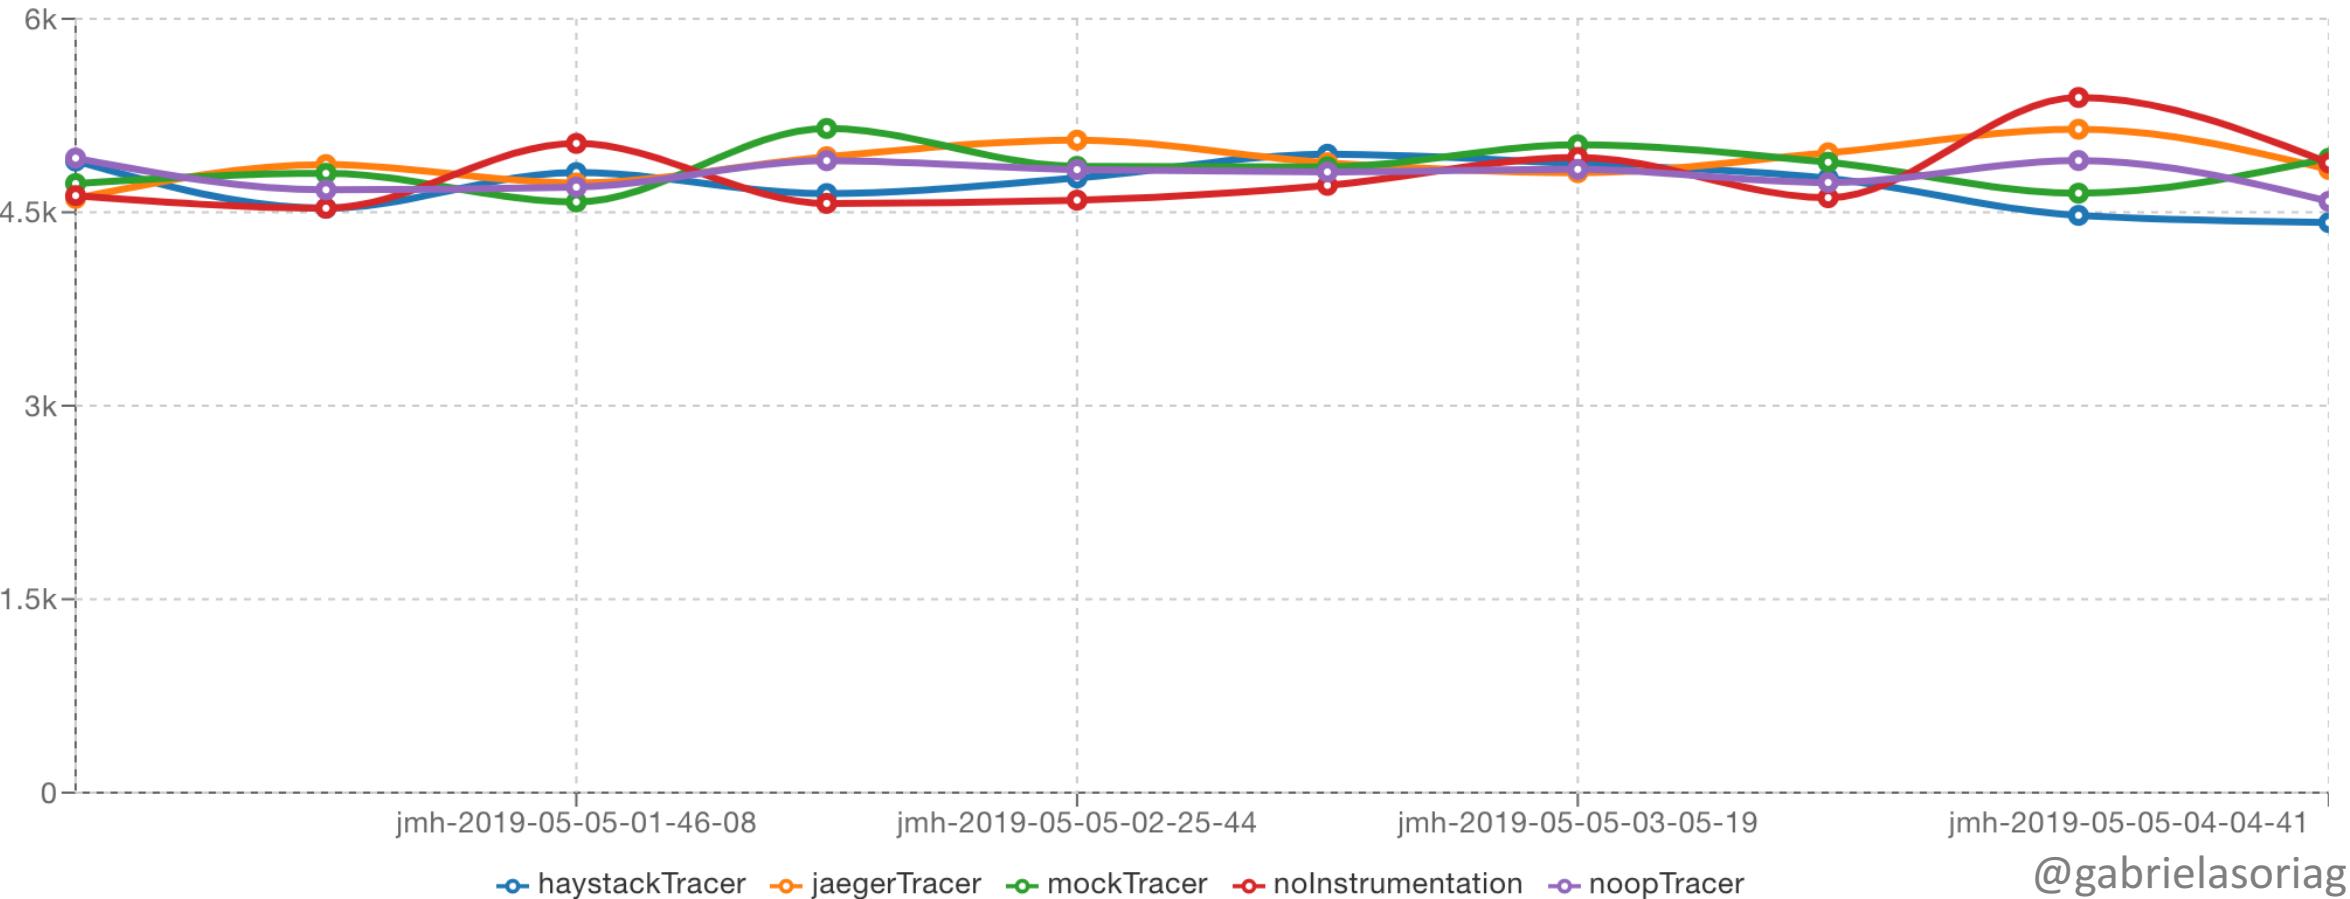


KubeCon

CloudNativeCon

Europe 2019

BenchmarkCourseManagementThroughput Throughput | Q | ⚖



@gabrielasoriag

Conclusions



KubeCon



CloudNativeCon

Europe 2019

- The overhead comes from the actual tracer and is closely related to how they work and how they are configured.
- In simple java scenarios the throughput **decreases ~90%** and the sample time **increases ~440%**
- In scenarios that include calls through the framework (spring boot, spring cloud, JDBC), on average, the throughput **decreases 12%** and the sample time **increases 14%**.

Conclusions



Europe 2019

- In the scenarios with **client calls through HTTP** (Servlet Filter, JAX-RS), the metrics show **no evidence of overhead**, as the deltas of throughput and sample time are not representative.

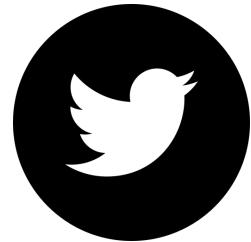
Next steps



Europe 2019

- Benchmark tests for Jaeger in different scenarios:
 - Sampling
 - gRPC vs. Thrift
 - Agent (UDP) vs. Collector (gRPC)
 - Different backend config
- Re-run the tests after merging OpenTracing and OpenCensus

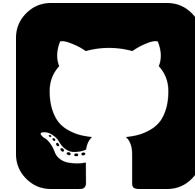
Keep in touch :)



@gabrielasoriag



OpenTracing Java Benchmarks



KubeCon



CloudNativeCon

Europe 2019