



KubeCon



CloudNativeCon

Europe 2019

Securing Kubernetes with Trusted Platform Module

Who are we, anyway?



KubeCon



CloudNativeCon

Europe 2019

Alex Tcherniakhovski

alextc@google.com

Security Engineer, Google
Kubernetes Engine

Andrew Lytvynov

awly@google.com

Software Engineer, Google
Kubernetes Engine

What's a Trusted Platform Module (TPM)?



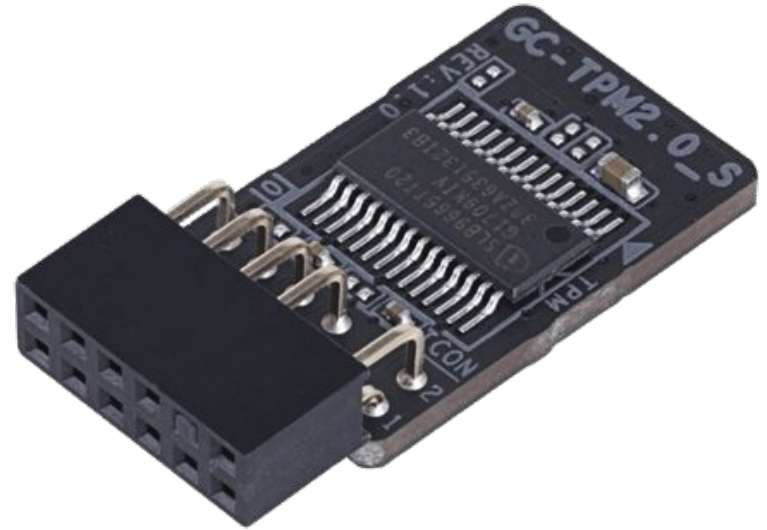
Crypto coprocessor

Hardware or software

Cheap, low-powered

Spec designed by Trusted Computing Group (TCG)

Spec versions 1.2 and 2.0



This talk

Goals:

- sample of TPM capabilities
- match k8s security challenges to TPM capabilities
- fuel exploration by users and sig-auth

Tricky security problems



KubeCon



CloudNativeCon

Europe 2019

Node trust bootstrap

- provide kubelet with credentials
- fully automated
- periodic rotation
- protect during Pod or Node compromise

Tricky security problems



KubeCon



CloudNativeCon

Europe 2019

First secret problem

- encrypt Secrets at rest
- store encryption key
- protect encryption key

Tricky security problems



KubeCon



CloudNativeCon

Europe 2019

Tamper-evident audit logging

- audit access to Secrets
- cryptographically-signed log
- verifiable log
- tamper-evident
- even with full master compromise

Agenda



KubeCon



CloudNativeCon

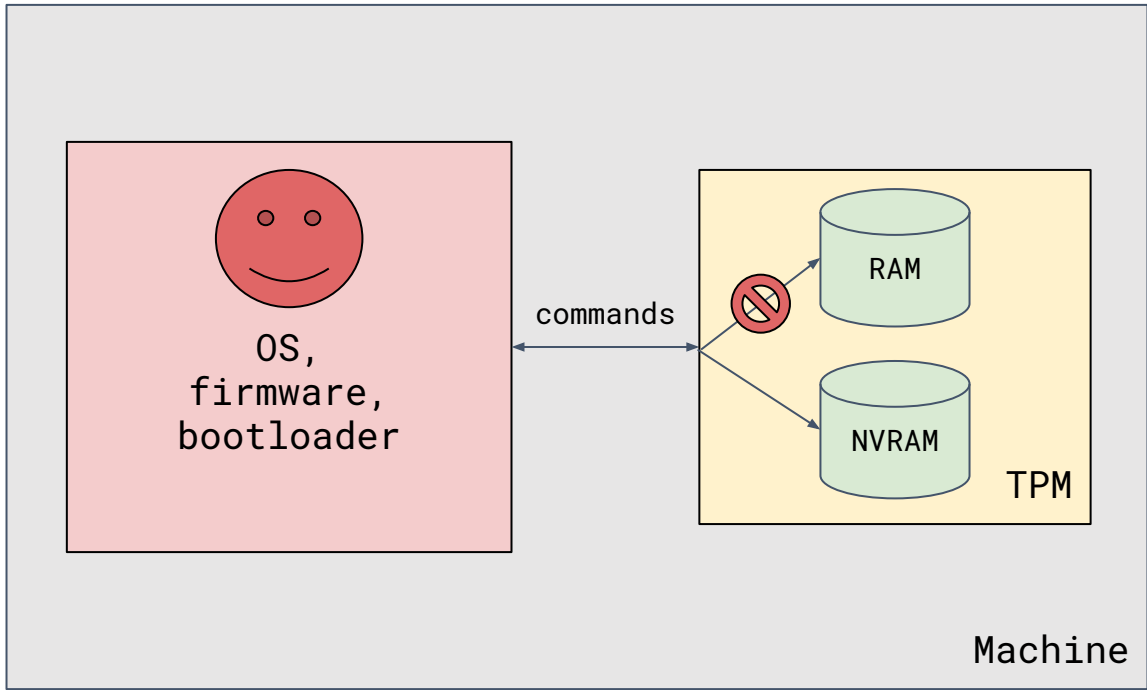
Europe 2019

1. Trusted Platform Module (TPM) crash course
2. Node trust bootstrap
3. First secret problem
4. Cryptographically protected audit log

1. Trusted Platform Module (TPM) crash course

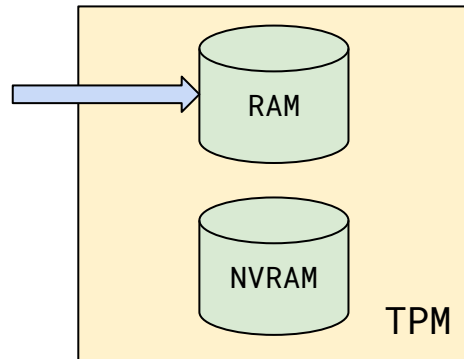
Classic use cases:

- Platform integrity
 - “is this corp machine in an expected state?”
- Disk encryption
 - BitLocker, dm-crypt, etc
 - protect encryption keys
 - verify integrity of bootloader/kernel/drivers



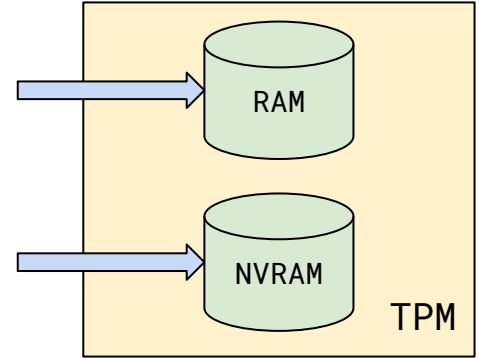
TPM keys

- RSA or ECDSA
- Encryption or signing
- Symmetric or asymmetric
- TPM-bound
 - no exfiltration
 - can export from TPM, but only encrypted
- Used via specialized commands

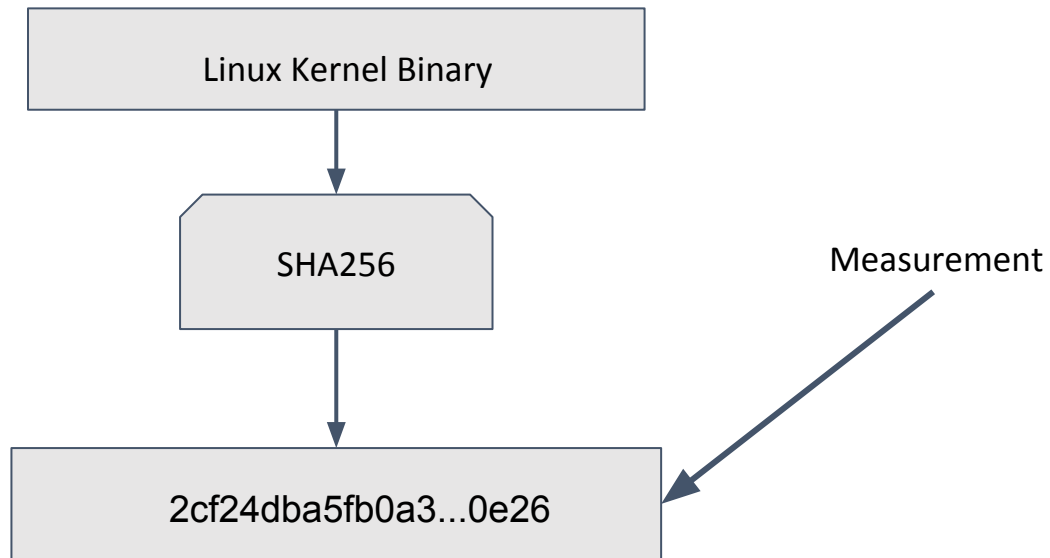


Endorsement Key (EK)

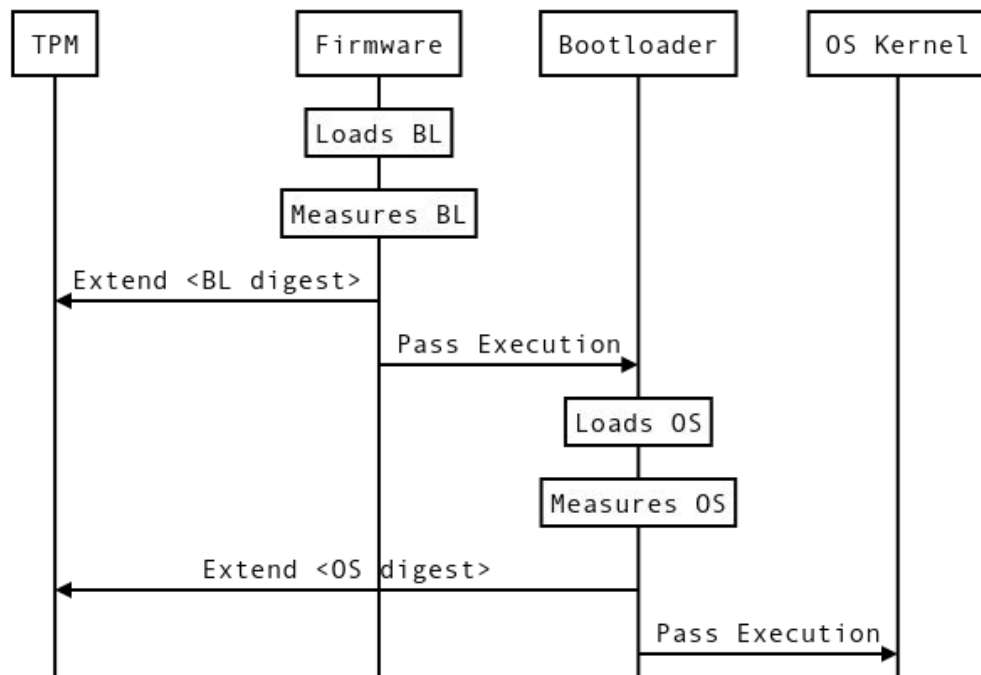
- Key baked into TPM
- Certificate signed by TPM vendor in NVRAM
- Used as machine identity



Measurements/Digests



Measured Boot



Platform Configuration Registers (PCRs)

PCR0

<-- Firmware

...

PCR4

<-- UEFI

PCR5

<-- Partition Table

PCR7

<-- Secure Boot Policy

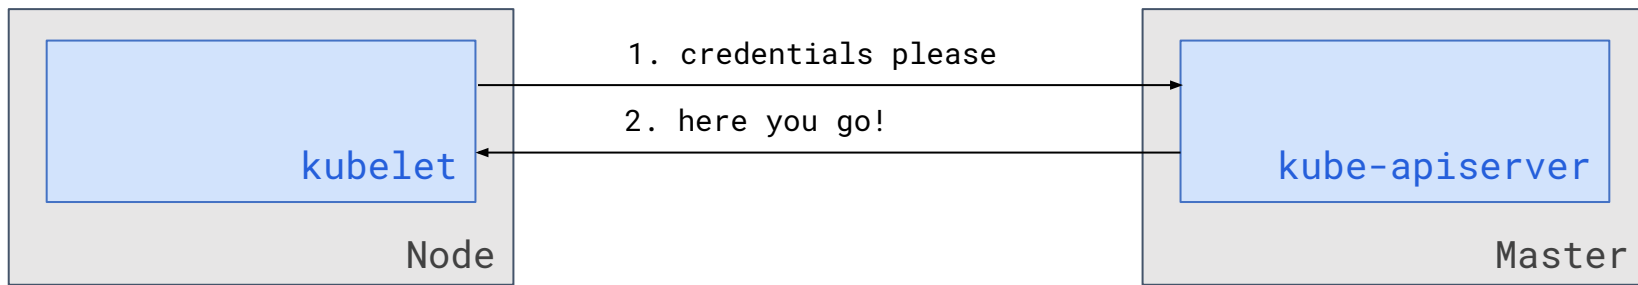
Extend Operation

$$\text{PCR} = \text{HASH}(\text{PCR} \parallel \text{datanew})$$

A whole lot more...

- RNG
- key hierarchies
- authorization policies
- certification
- dictionary attack protection
- command audit
- external/transferable keys

2. Node trust bootstrap



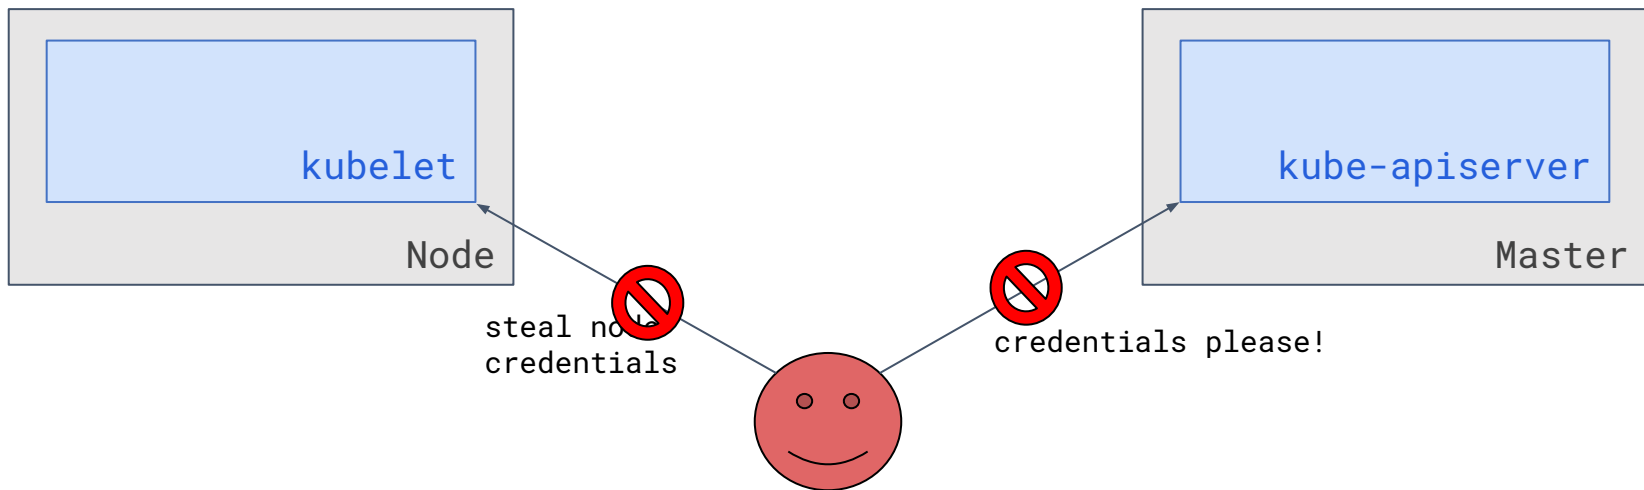
Threat model

Attacker has:

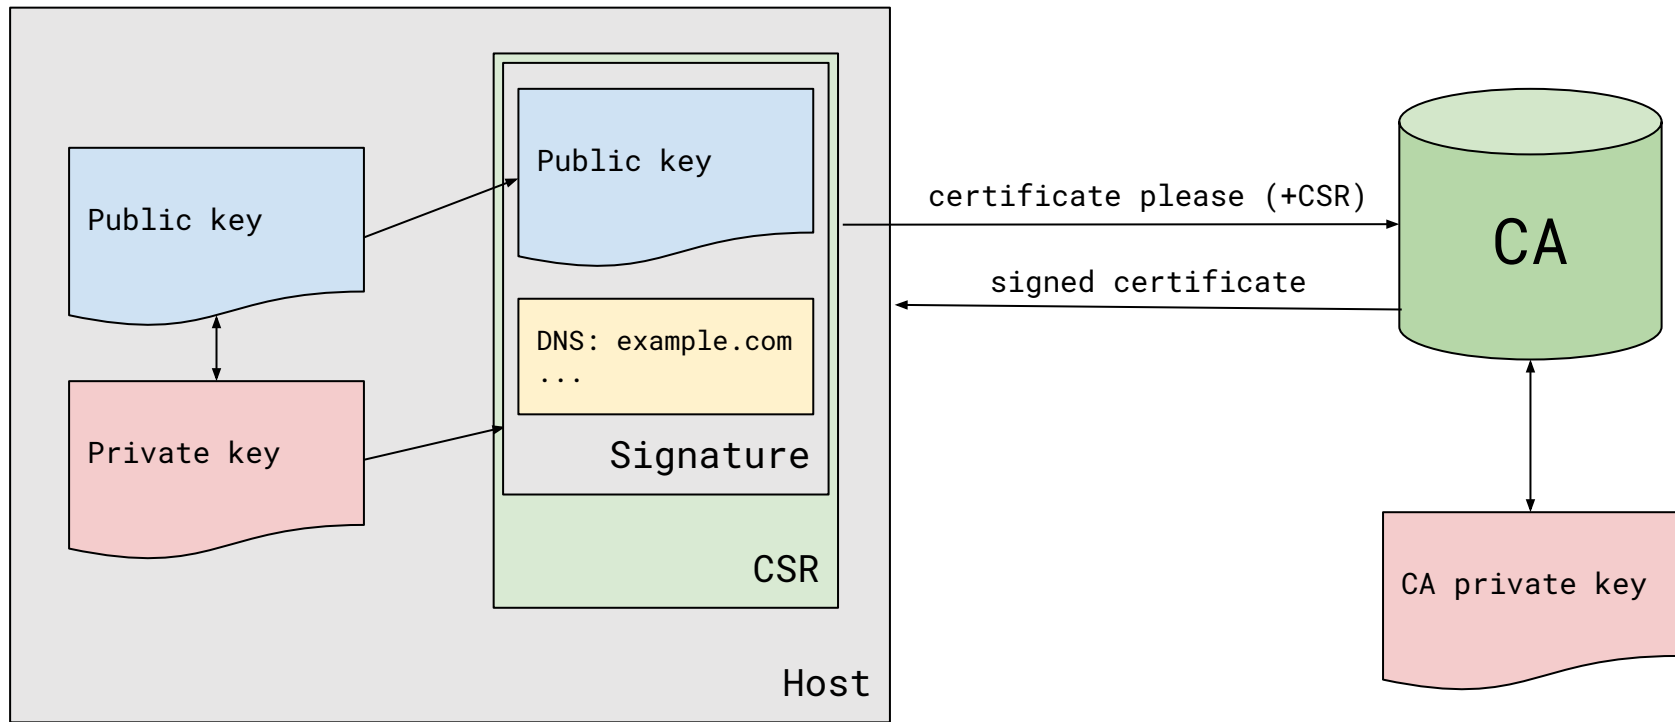
- compromised Pod
- compromised Node

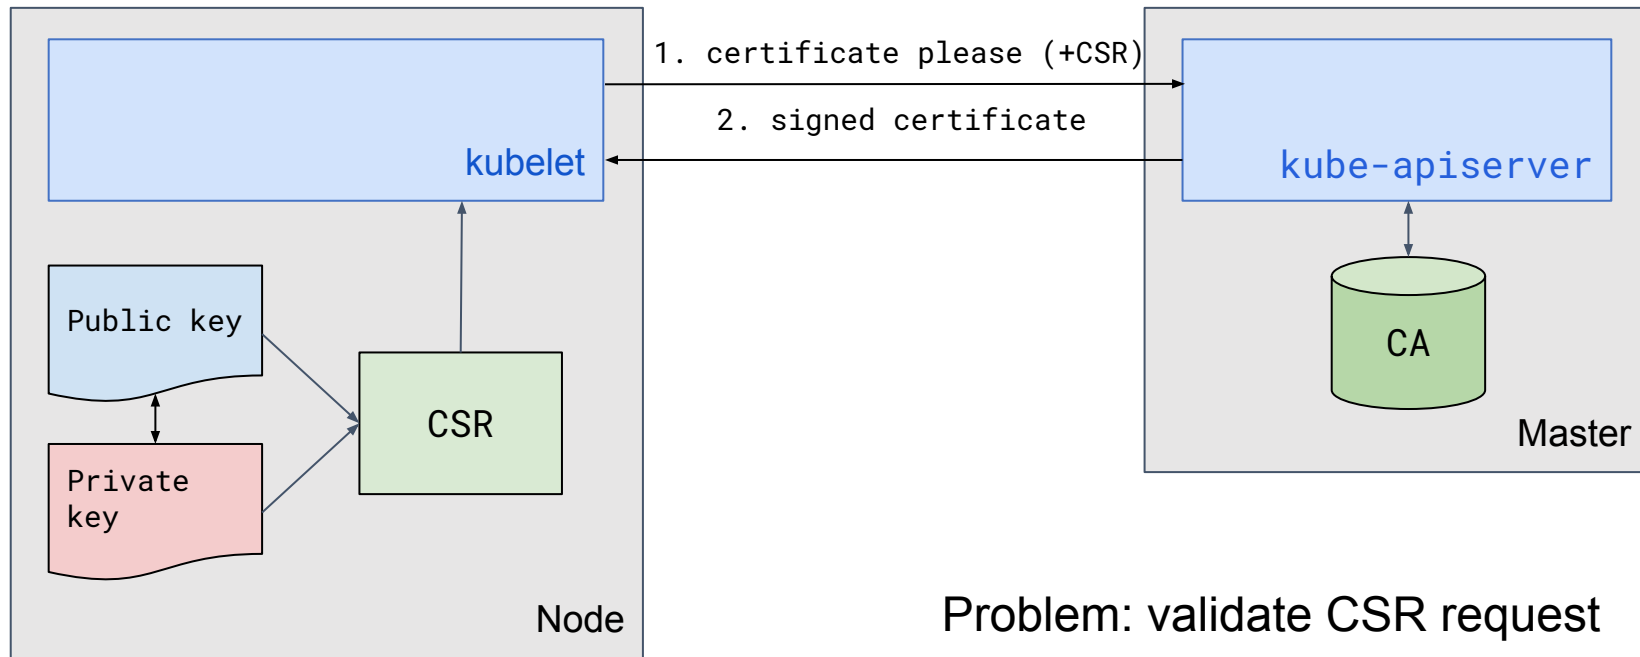
Attacker wants:

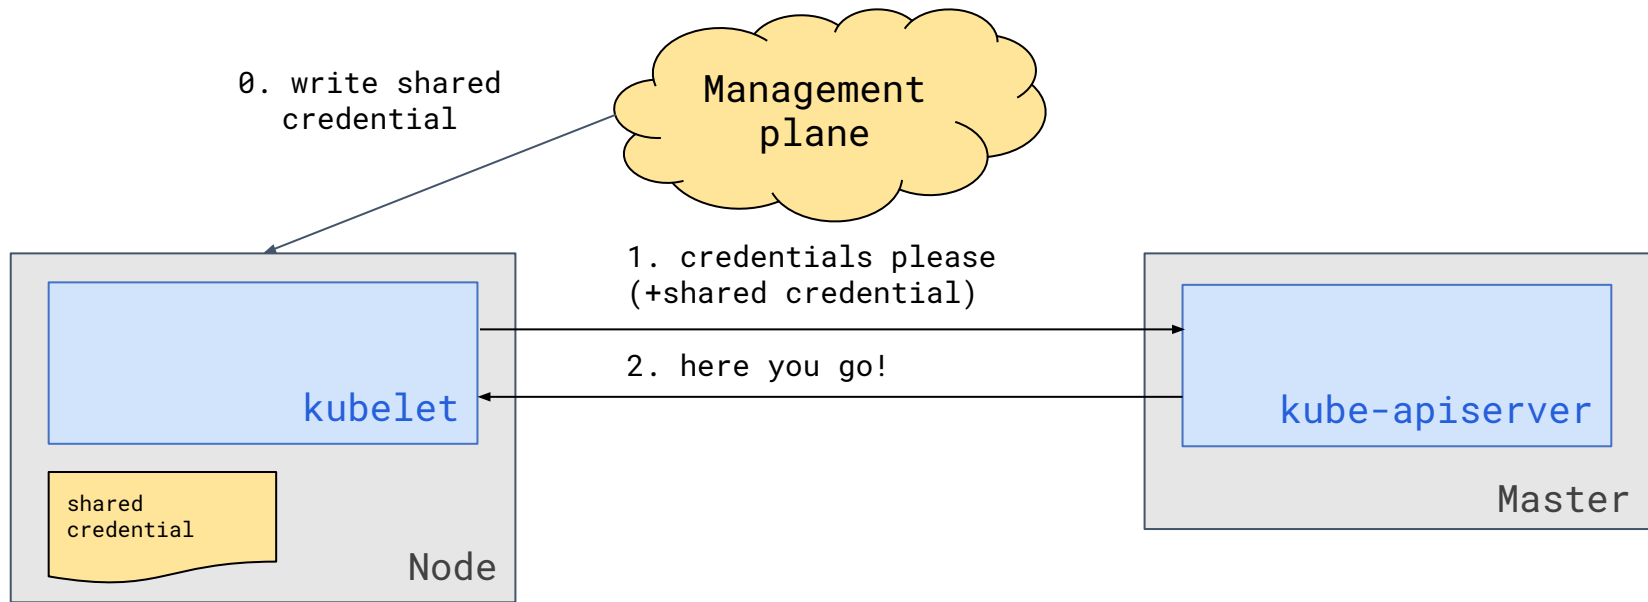
- exfiltrate application configs
- exfiltrate application Secrets
- persist access

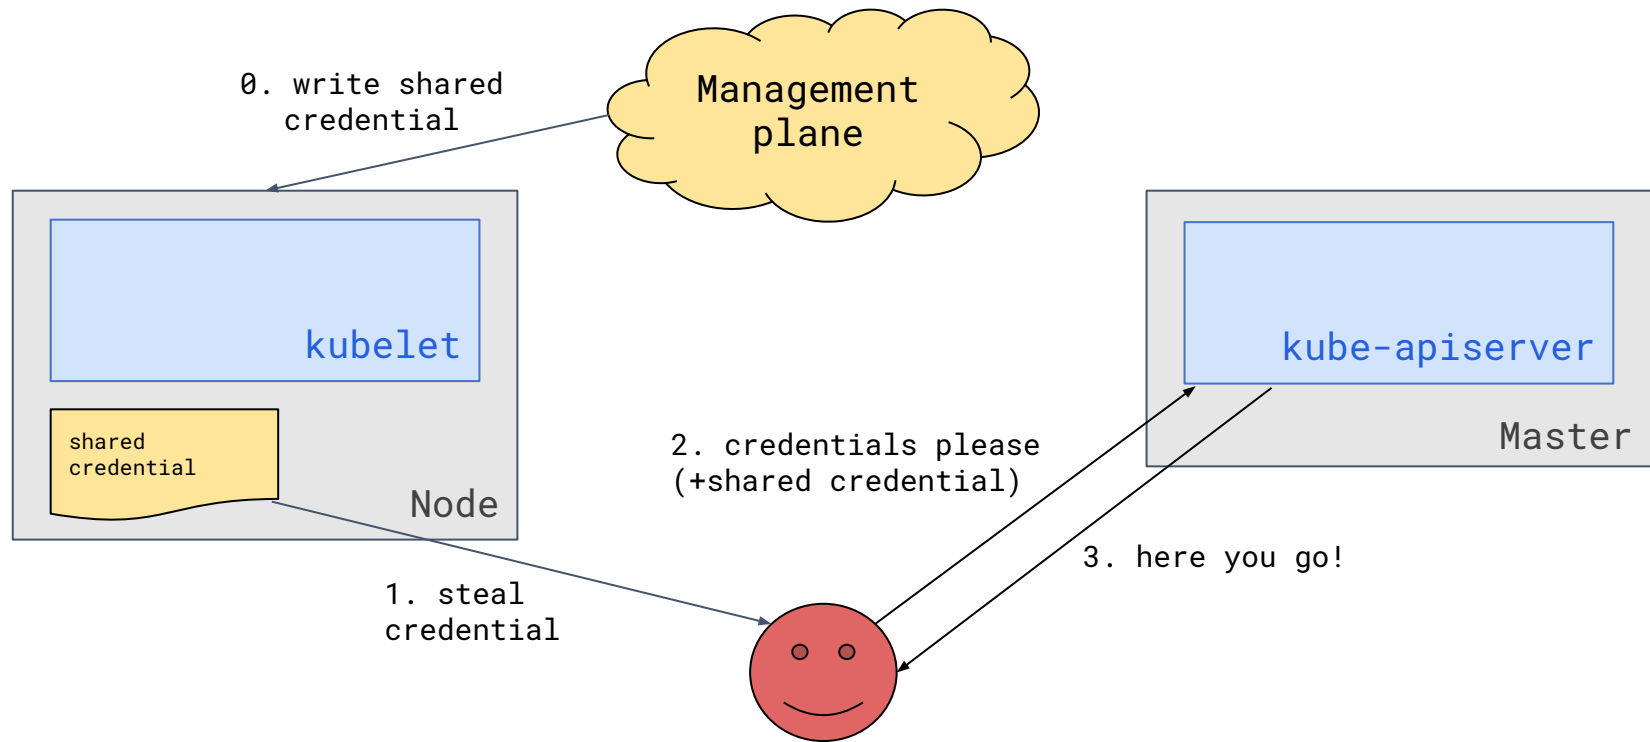


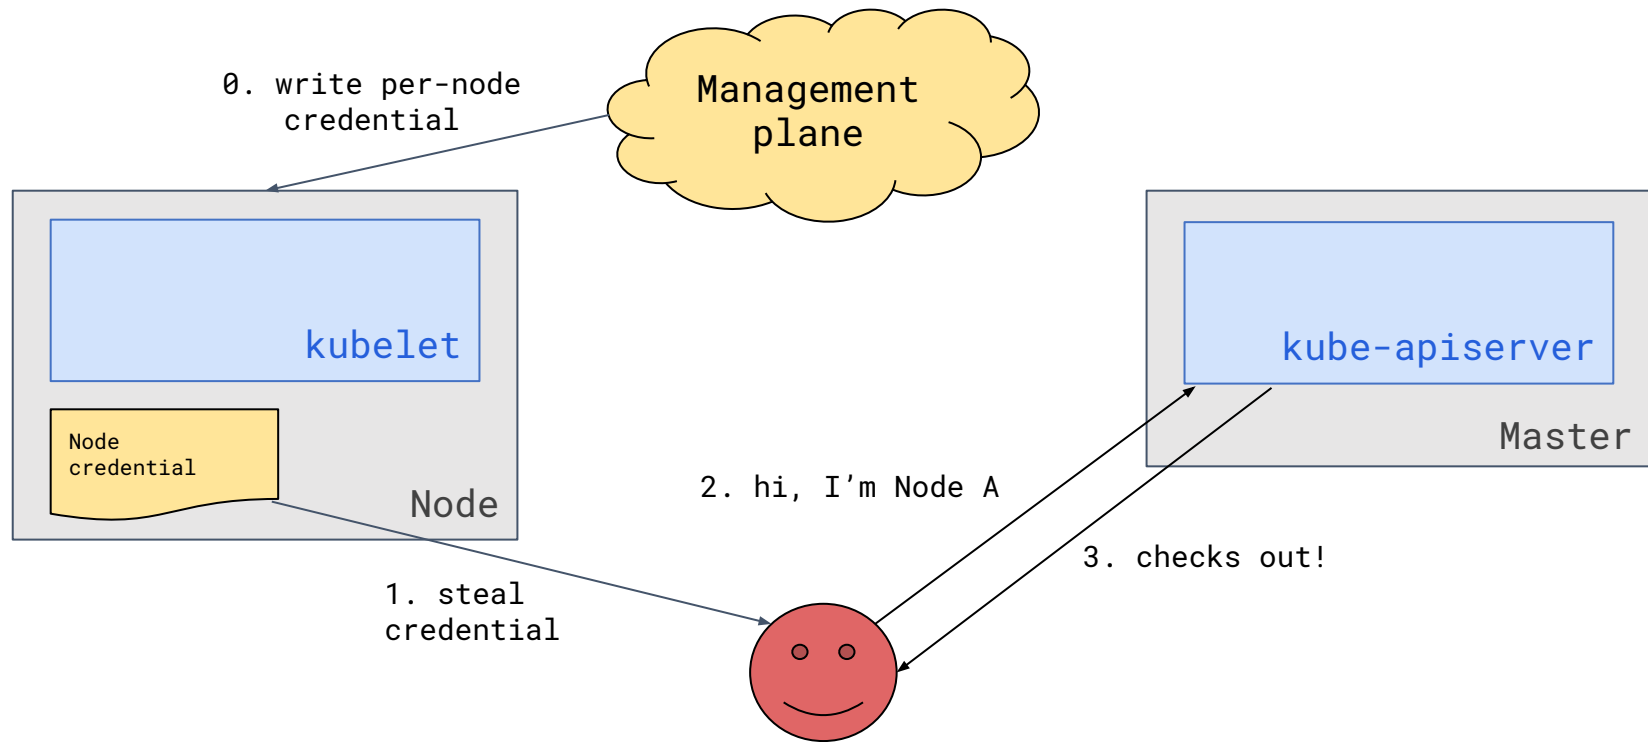
Enter X.509 CSRs and Certificates





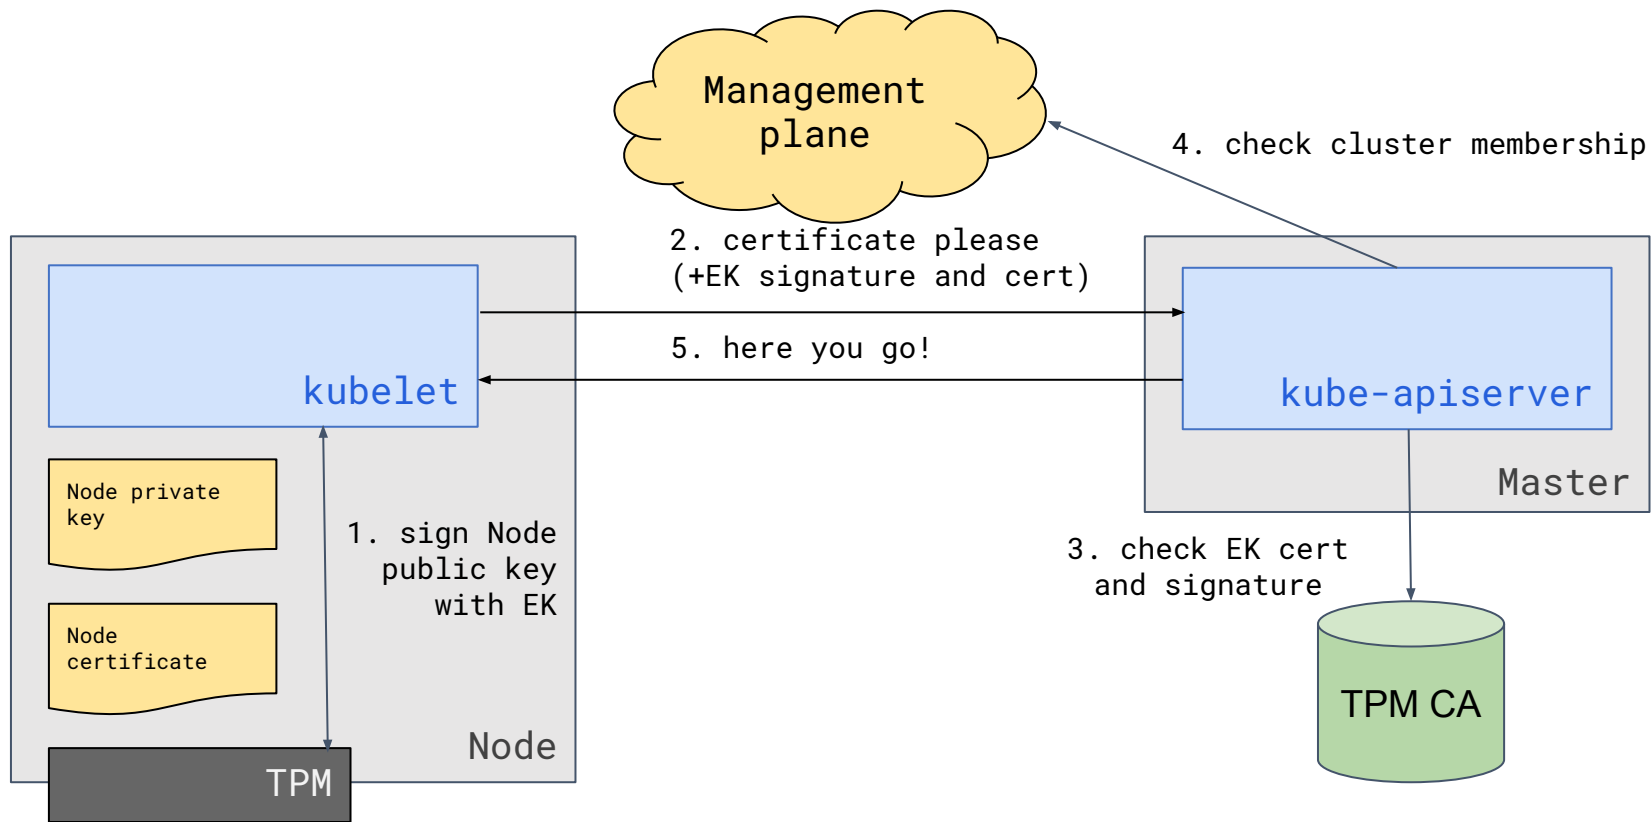






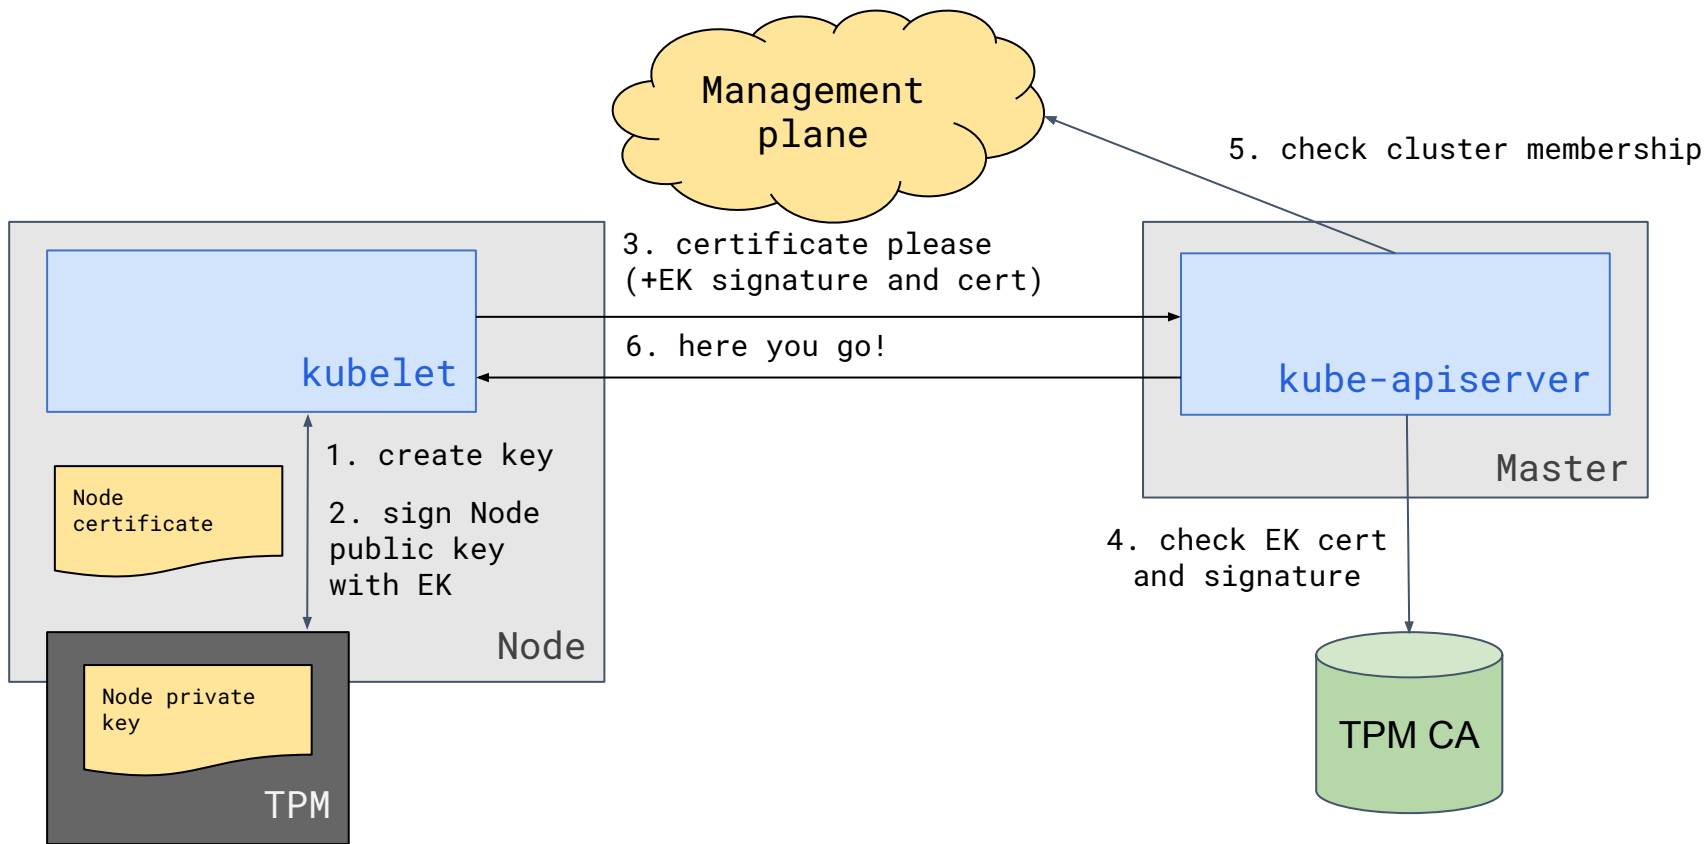
Let's use a TPM!

EK as proof of machine identity



But what about exfiltration of the Node credential after provisioning?

Put it in a TPM!



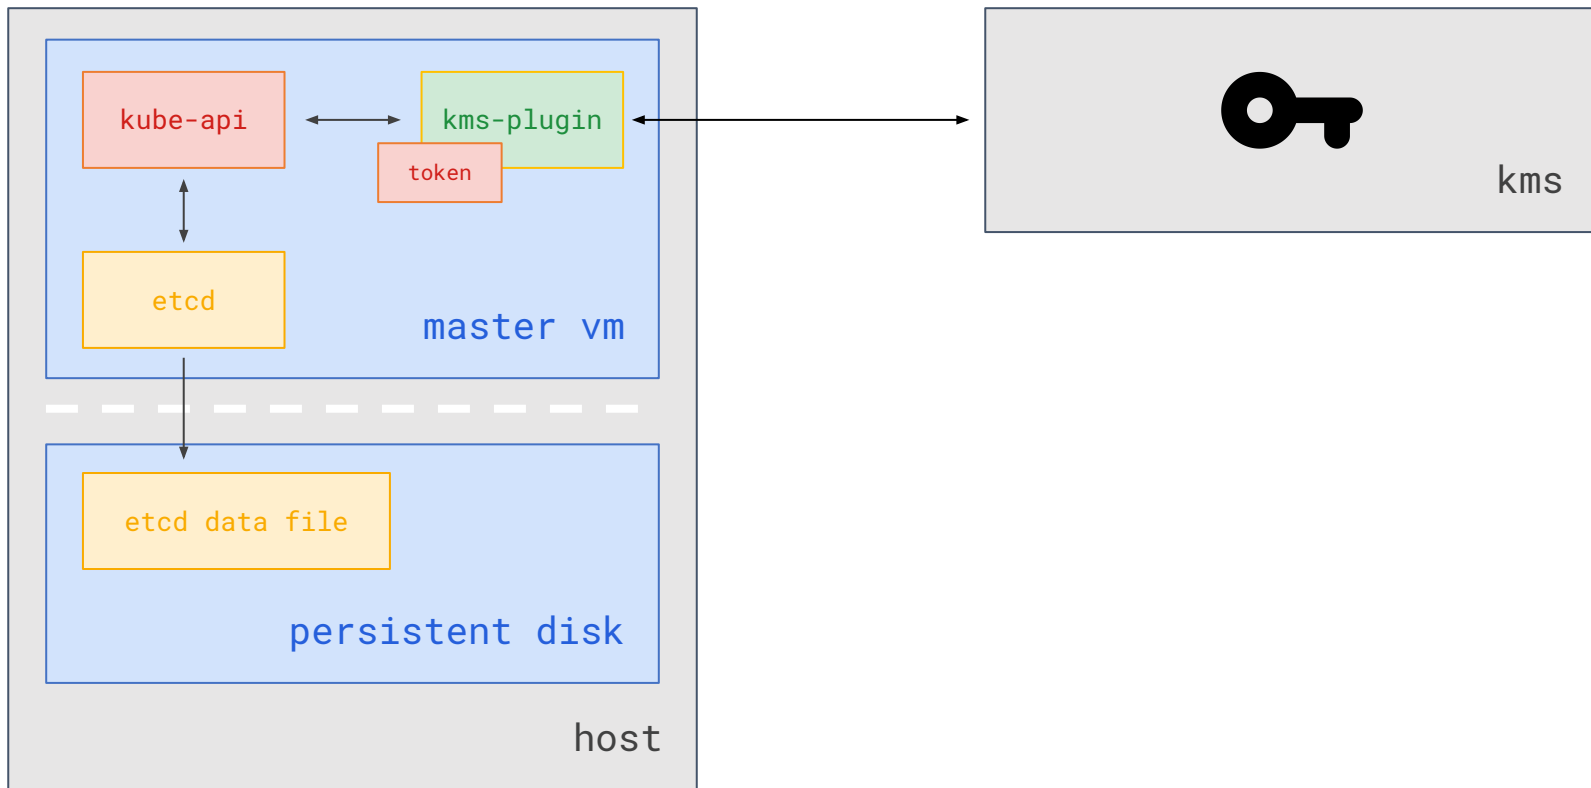
Not 100% solution

Attacker can still use Node credential via RCE on the Node.

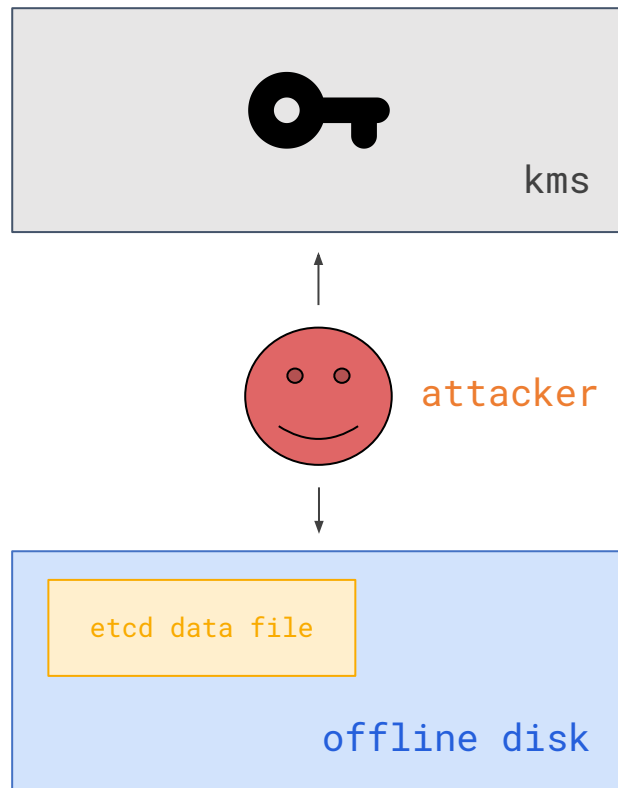
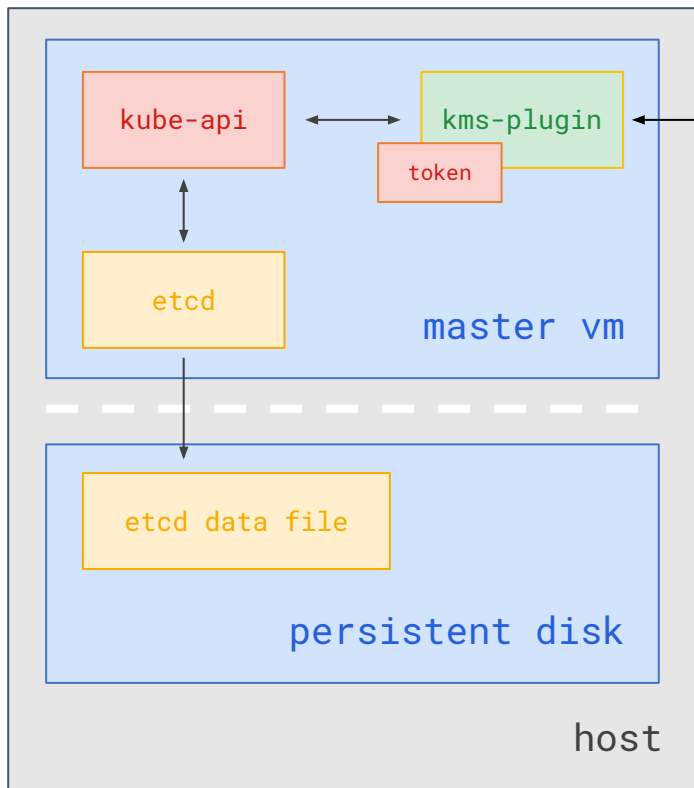
But things are in a much better state!

- requires constant Node access
- mitigated after patching vuln
- use industry standard for trust bootstrap

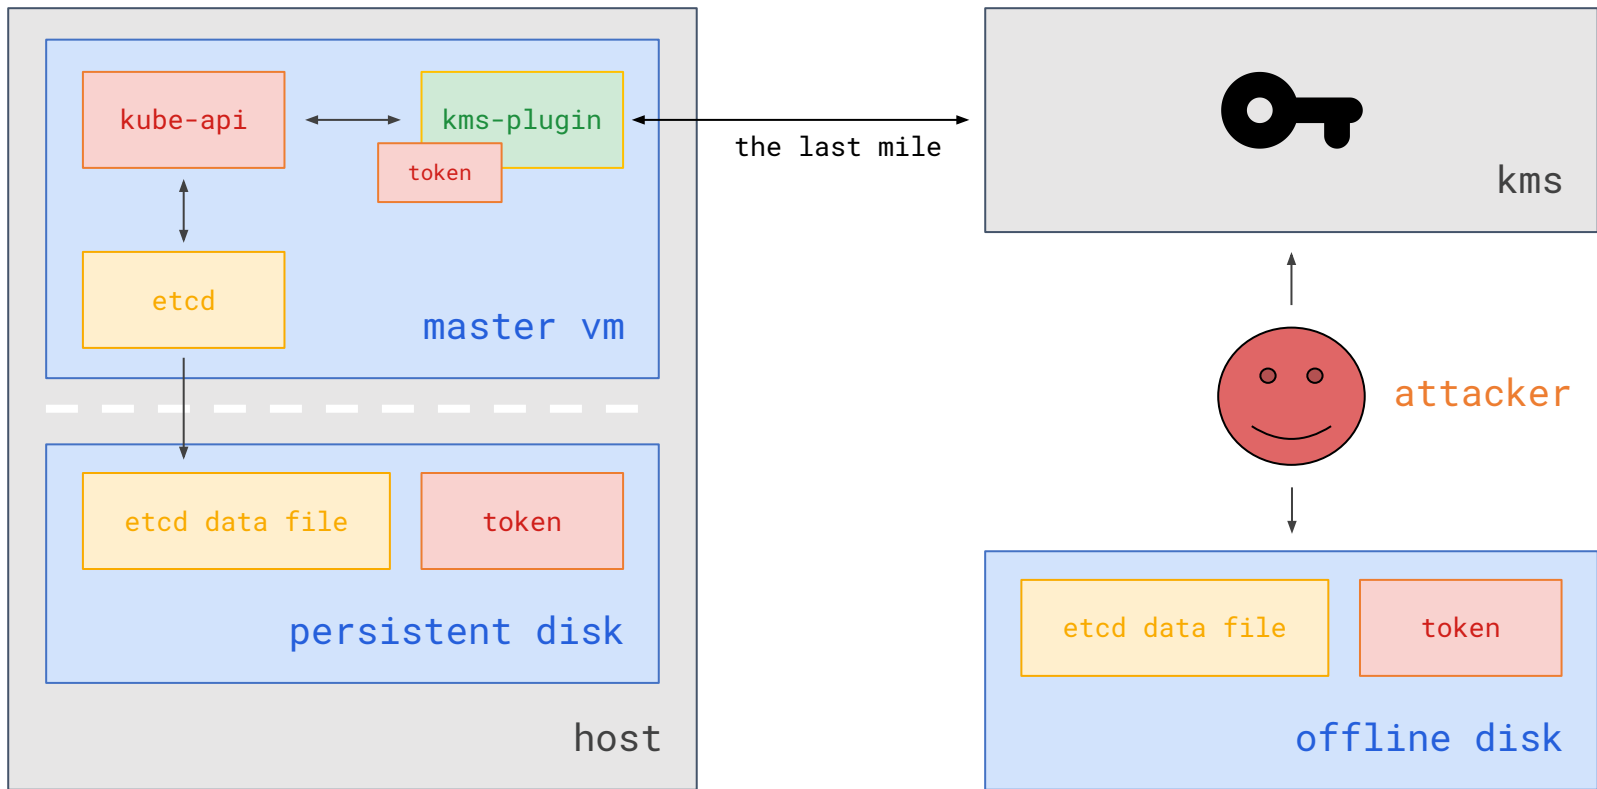
3. Solving the first secret problem



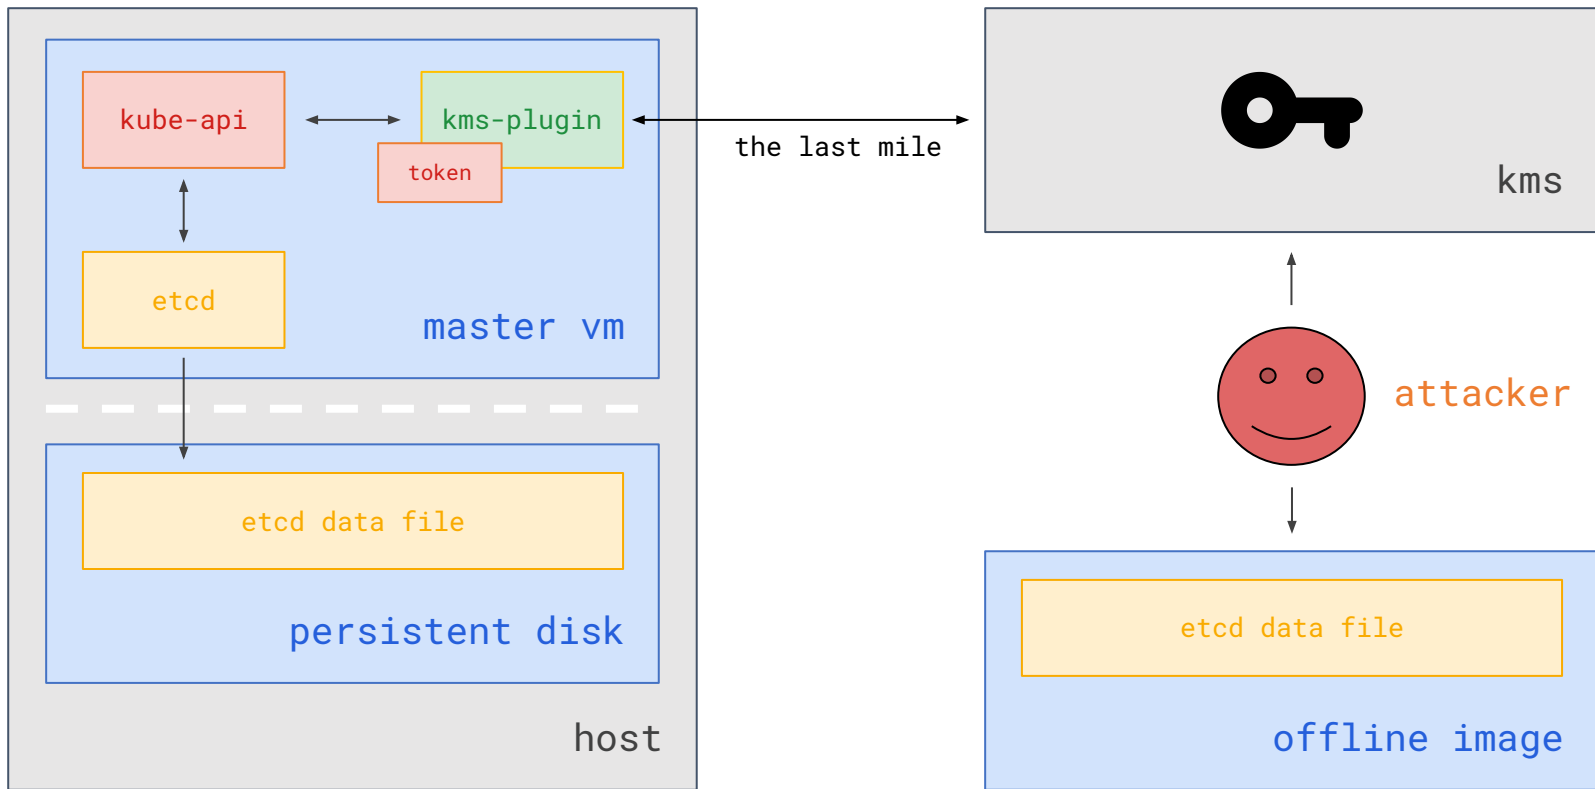
KMS Plugin



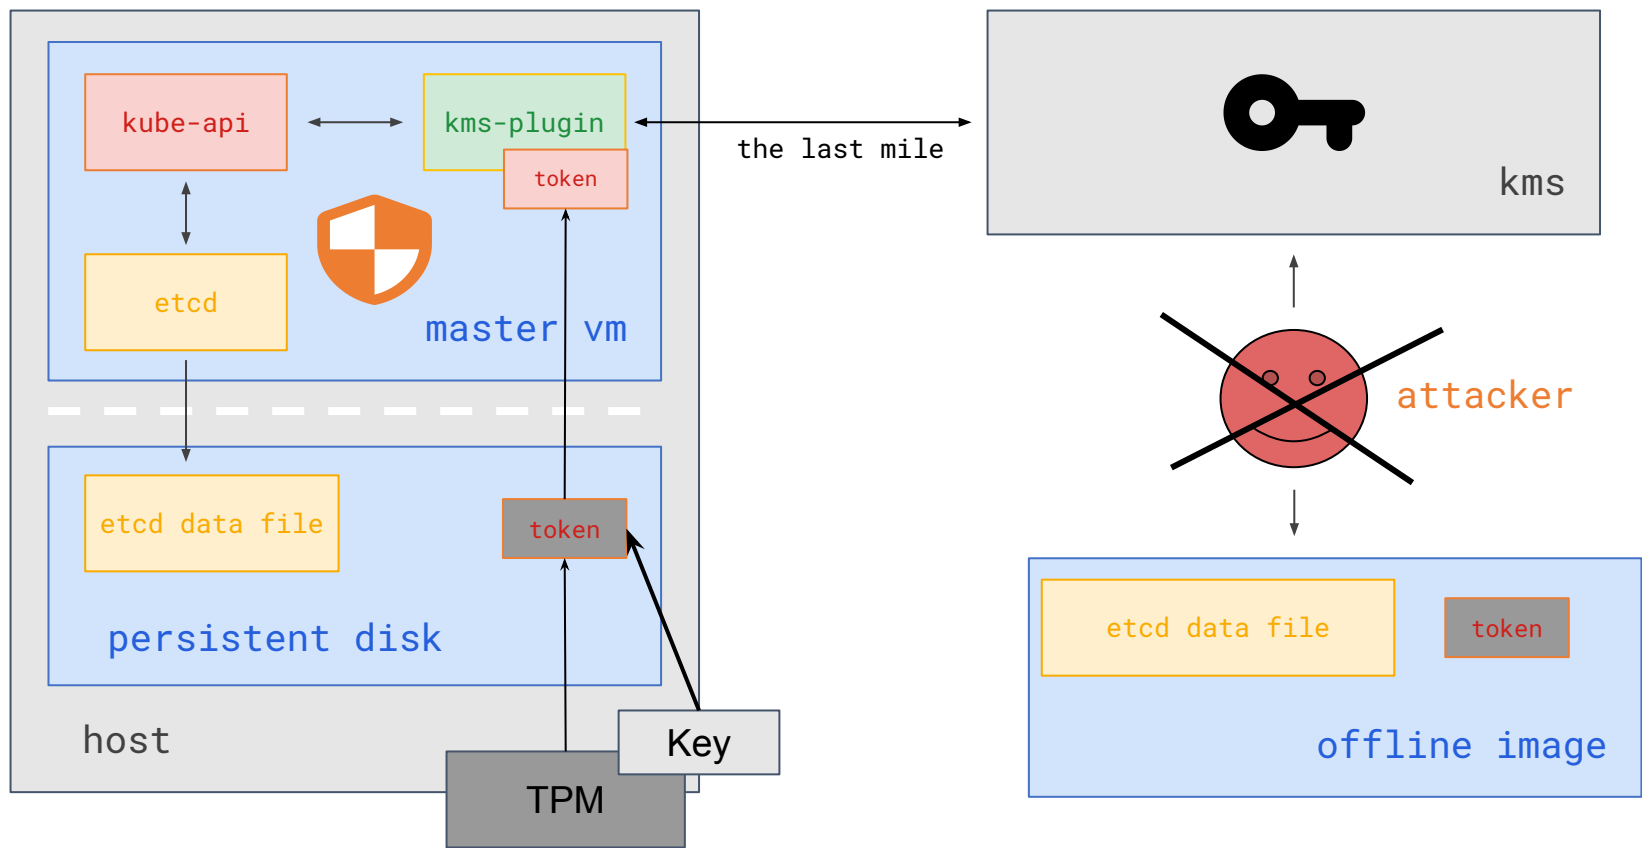
Threat Model



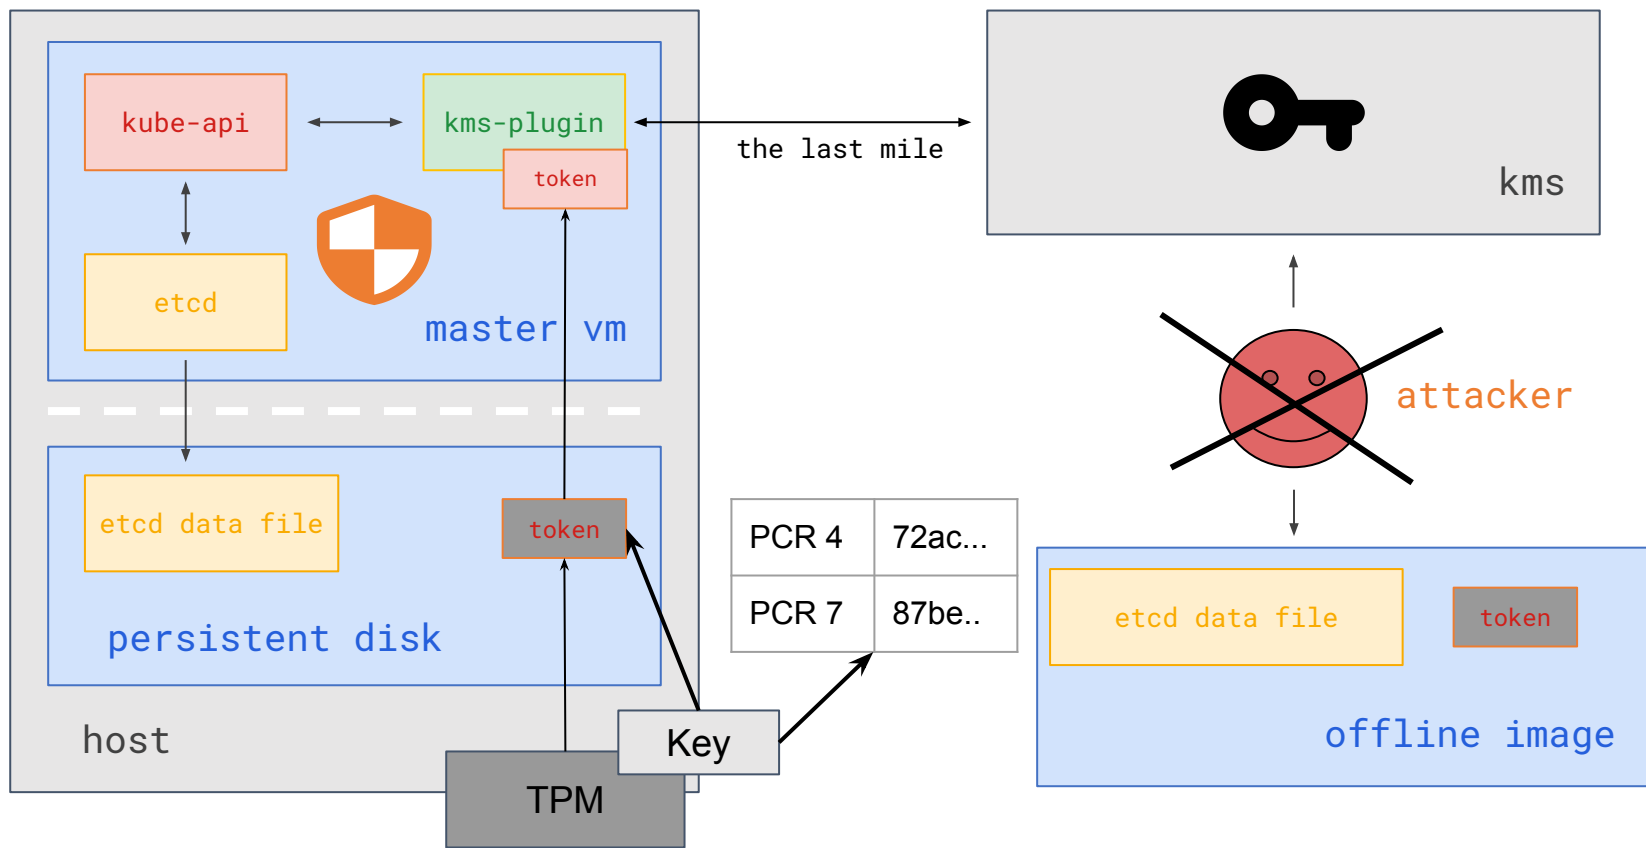
The last-mile problem



Goal: Do **NOT** get access to keys



Solution: Seal KMS Credential to TPM

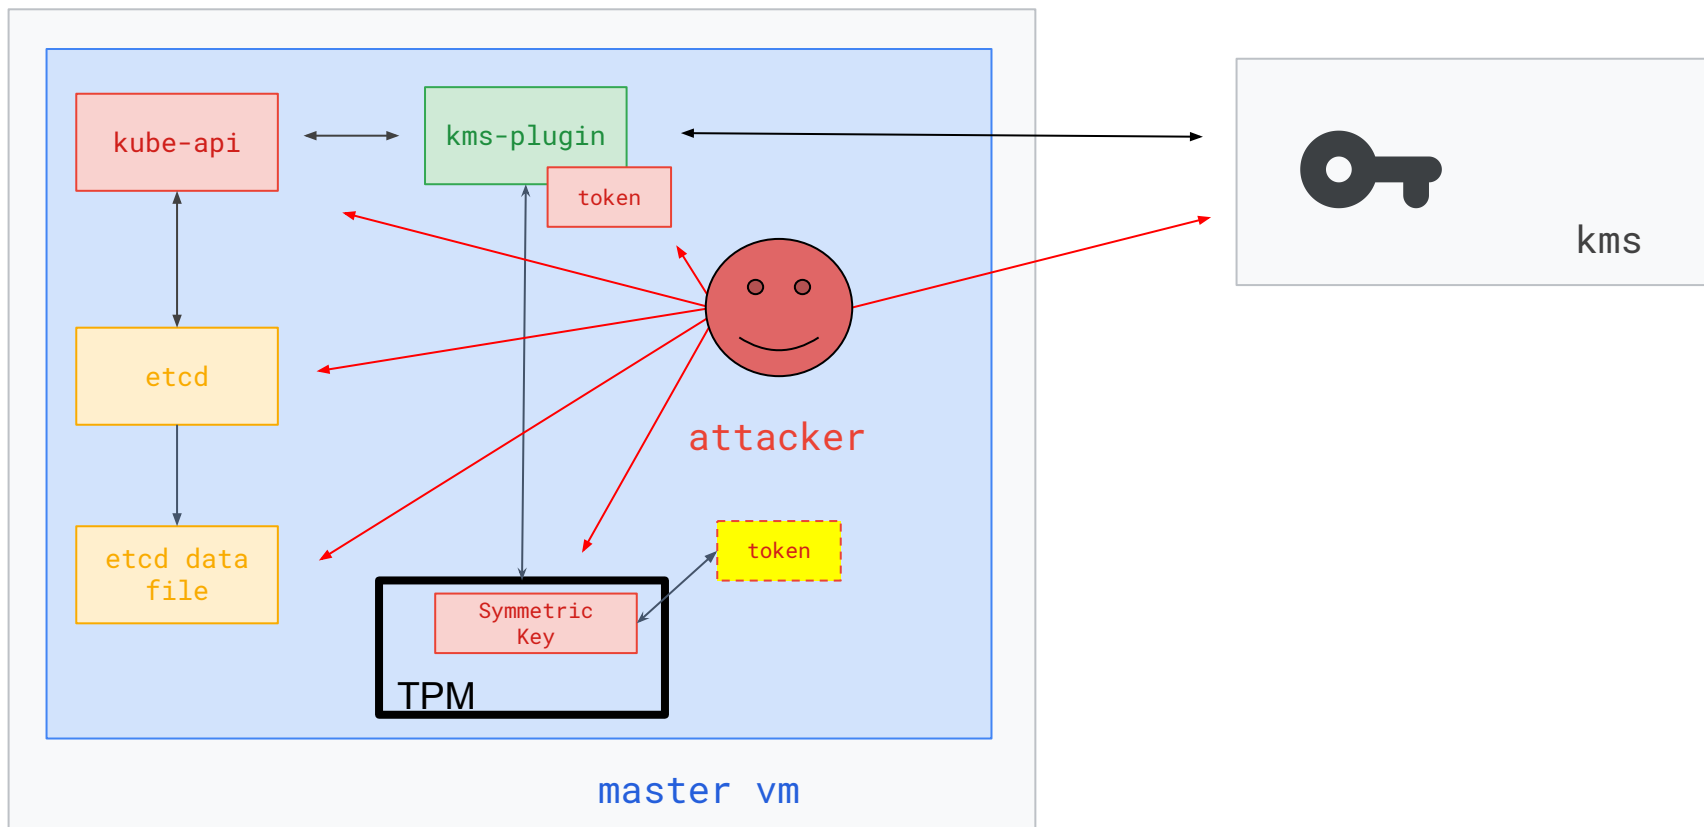


Apply: PCR Policy

Practical Use

- Embed into your applications:
<https://github.com/google/go-tpm/tree/master/examples/tpm2-seal-unseal>
- Script via an Init container:
<https://github.com/tpm2-software/tpm2-tools>

4. Tamper-evident audit logs*



Threat Model

$$\text{audit}_{\text{new}} = \mathbf{H}_{\text{auditAlg}}(\text{audit}_{\text{old}} \parallel \text{inputHash} \parallel \text{outputHash})$$

Auditing TPM Commands

Issued certificate #1:

Certificate:

Data:

Serial Number:

10:e6:fc:62:b7:41:8a:d5:00:5e:45:b6

pub:

00:c9:22:69:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:



TPM2-Sign ...



TPM Audit Register

2c503..

Logs are stored externally

Issued certificate #1:

Certificate:

Data:

Serial Number:

10:e6:fc:62:b7:41:8a:d5:00:5e:45:b6

pub:

00:c9:22:69:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

Issued certificate #2:

Certificate:

Data:

Serial Number:

19:e8:fc:62:b7:41:8a:d5:00:5e:45:b9

pub:

00:c7:22:79:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

TPM2-Sign ...

TPM Audit Register

9b863..

Logs are stored externally

Issued certificate #1:

....

Issued certificate #2

...

Issued certificate #3:

Certificate:

Data:

Serial Number:

20:e6:fc:32:b7:41:8a:d5:00:5e:45:b9

pub:

00:a7:229:79:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

TPM Audit Register

TPM2-Sign ...

8b300..

Attacker gets a certificate

Issued certificate #1:

Certificate:

Data:

Serial Number:

10:e6:fc:62:b7:41:8a:d5:00:5e:45:b6

pub:

00:c9:22:69:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

Issued certificate #2:

Certificate:

Data:

Serial Number:

19:e8:fc:62:b7:41:8a:d5:00:5e:45:b9

pub:

00:c7:22:79:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

TPM2-Sign ...

TPM Audit Register

8b300..

Logs are examined externally

TPM2_GetCommandAuditDigest

8b300

Issued certificate #1:

Certificate:

Data:

Serial Number:

10:e6:fc:62:b7:41:8a:d5:00:5e:45:b6

pub:

00:c9:22:69:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

Issued certificate #2:

Certificate:

Data:

Serial Number:

19:e8:fc:62:b7:41:8a:d5:00:5e:45:b9

pub:

00:c7:22:79:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

TPM2_GetCommandAuditDigest

9b863...

TPM Audit Register

TPM2-Sign ...

9b863...

Logs are examined externally

TPM2_GetCommandAuditDigest

Issued certificate #1:

Certificate:

Data:

Serial Number:

10:e6:fc:62:b7:41:8a:d5:00:5e:45:b6

pub:

00:c9:22:69:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

Issued certificate #2:

Certificate:

Data:

Serial Number:

19:e8:fc:62:b7:41:8a:d5:00:5e:45:b9

pub:

00:c7:22:79:31:8a:d6:6c:ea:da:c3:7f:2c:ac:a5:

TPM2-Sign ...


TPM Audit Register

Num Of Logs = 2

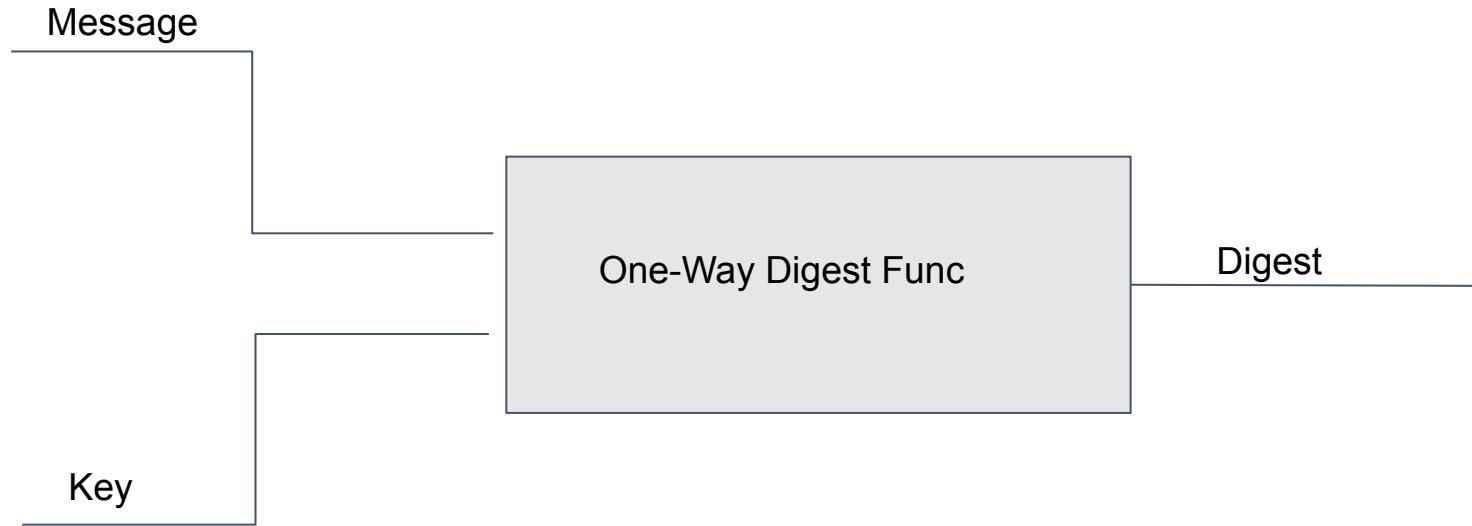
...

What if attackers reset and replay

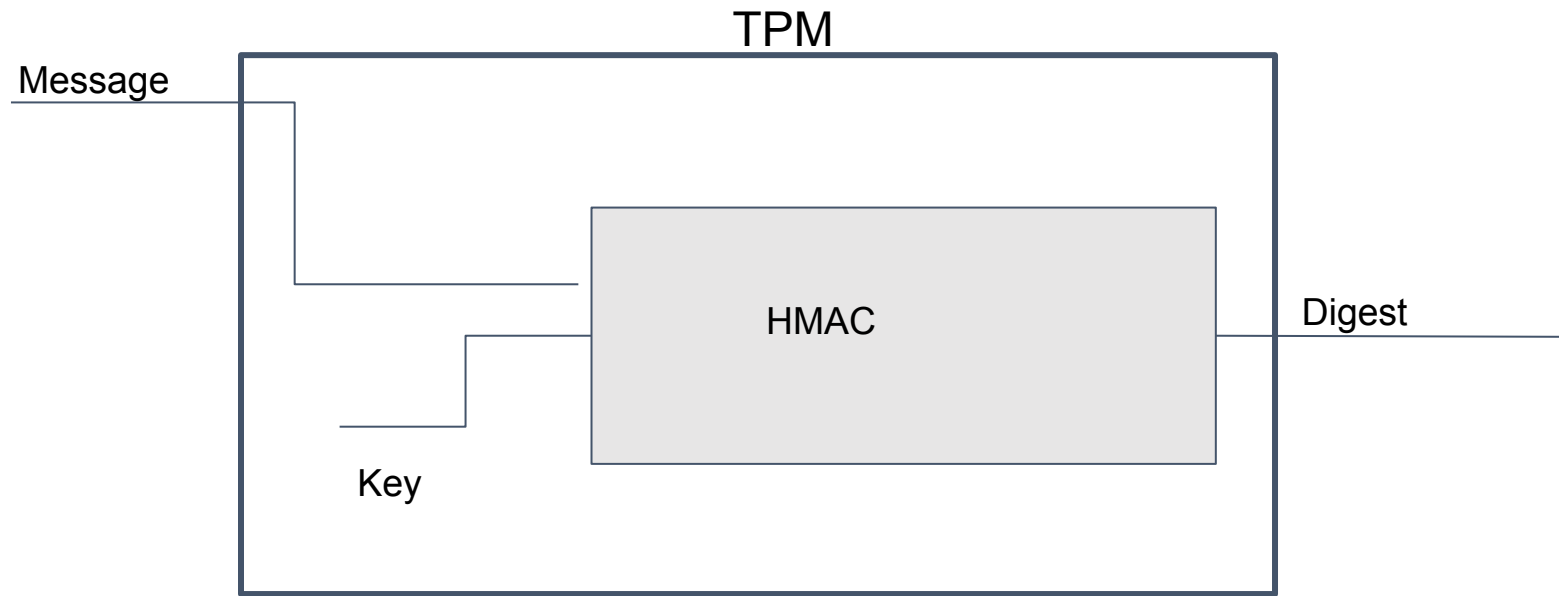
Are we done?

KMS	AEAD
TPM	HMAC
TPM	AUDIT 



Building Blocks



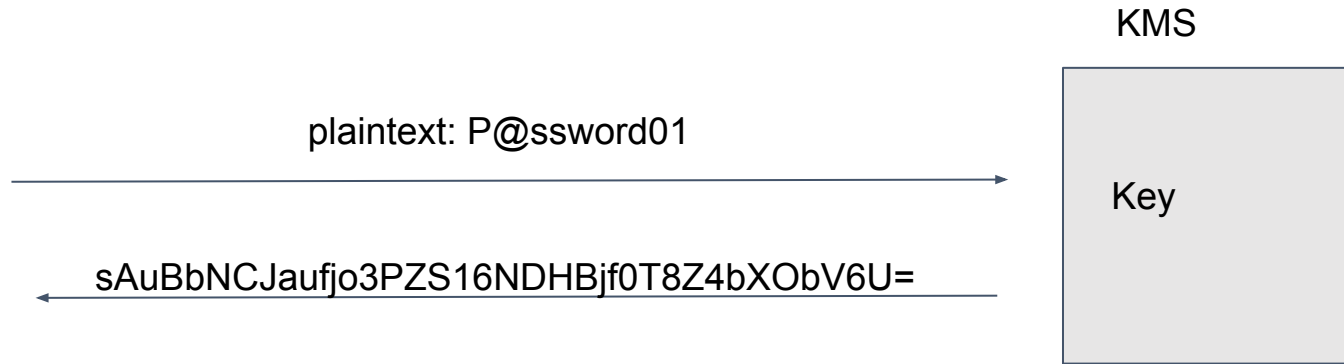
Hash-Based Message Authentication Code



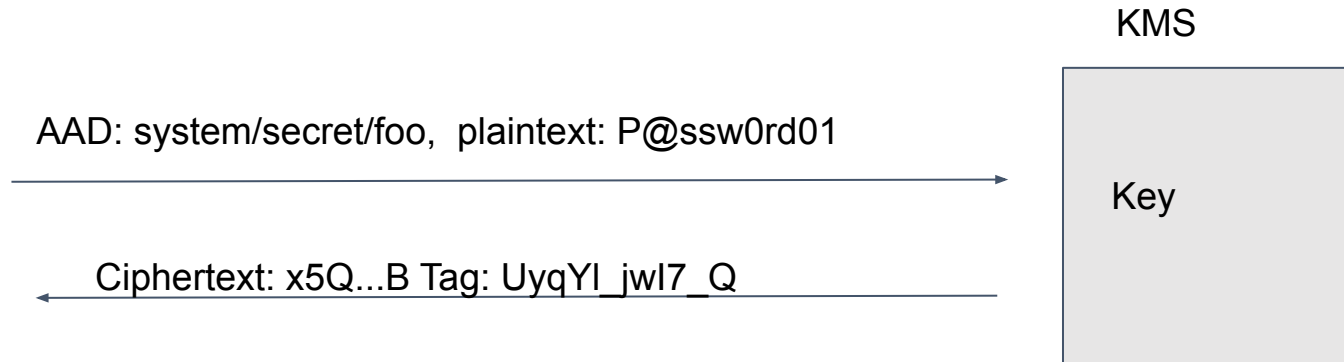
TPM2_HMAC

KMS	AEAD
TPM	HMAC 
TPM	AUDIT 

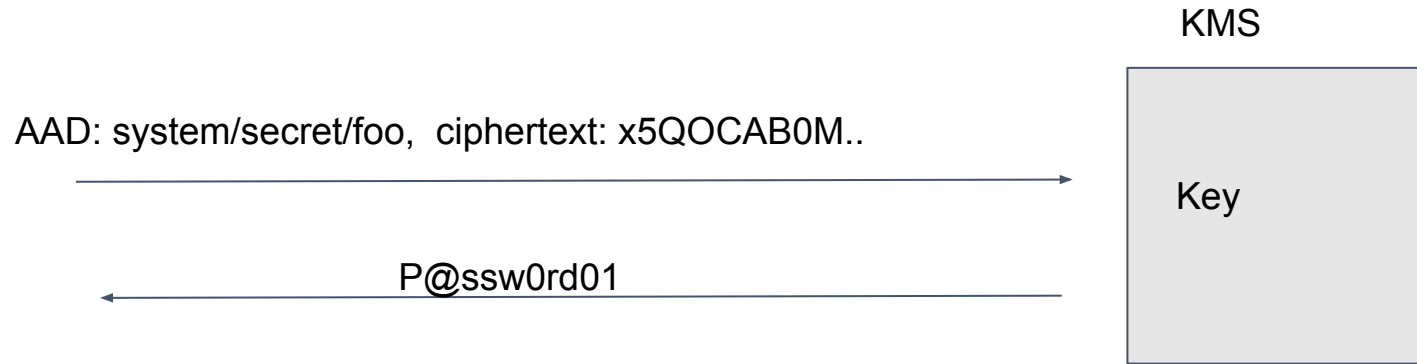
Building Blocks



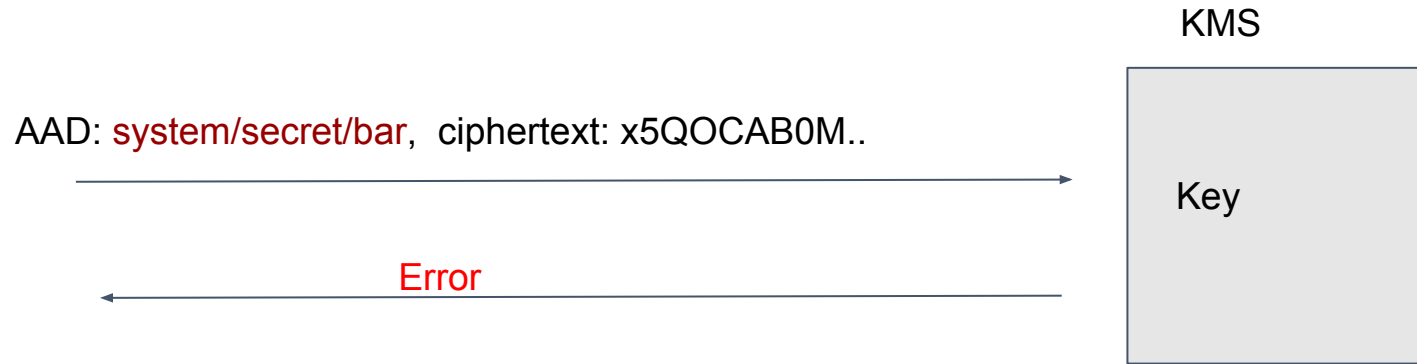
Symmetric Encryption



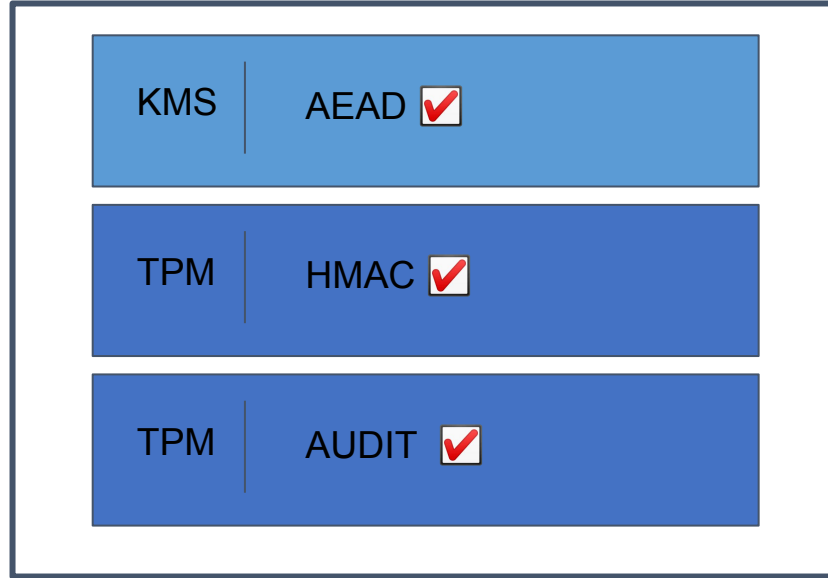
AEAD Encryption



AEAD Encryption



AEAD - AAD must match



Building Blocks

Additional Auth Data

Plaintext

TPM2_HMAC (system/my-dba-pwd)

P@ssw0rd01

Use TPM2_HMAC to generate AAD

Attacks not covered

- Reading directly from kube-apiserver cache
- Reading KEK from kms-plugin cache
- Waiting for a request from a legitimate user and intercepting the response

Summary



KubeCon



CloudNativeCon

Europe 2019

When not to use TPMs

- Performance-sensitive crypto (unless virtual)
- Bulk encryption
- As a substitute for physical security, it is tamper-resistant not tamper-proof

References



KubeCon



CloudNativeCon

Europe 2019

- [TPM 2.0 specification](#)
- [Turtles All the Way Down: Managing Kubernetes Secrets](#)
- [Securing Kubernetes Secrets](#)
- [Continuous Tamper-proof Logging using TPM2.0](#)
- [Cryptographic Support for Secure Logs on Untrusted Machines](#)
- [go-tpm library](#)
- [TPM2 Tools](#)
- [K8S KMS Plugin for Google CloudKMS](#)



KubeCon

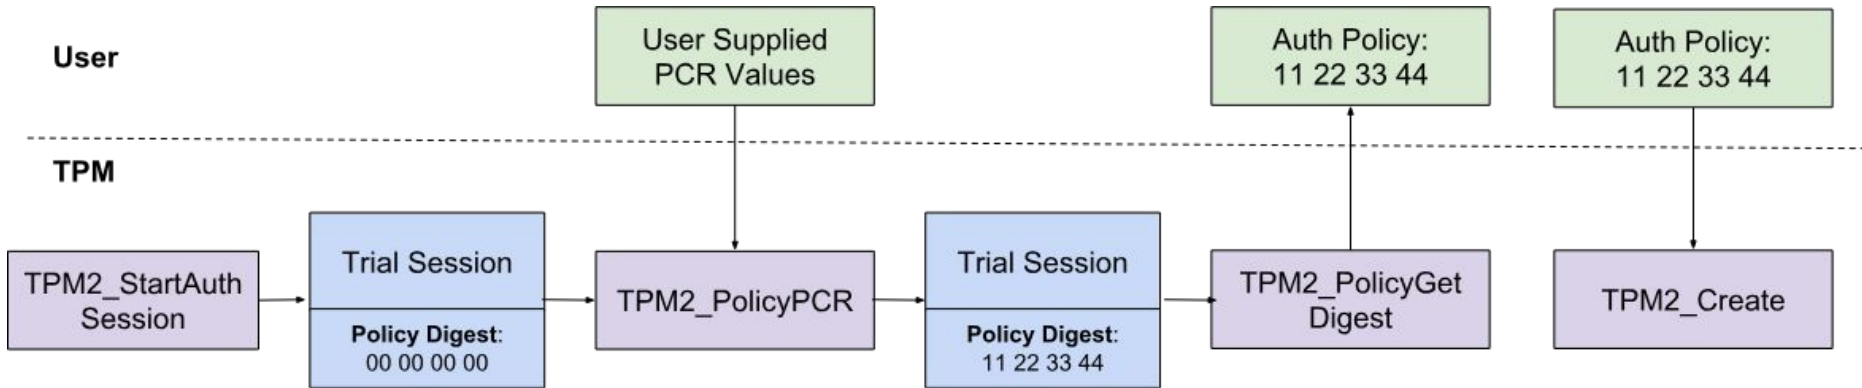


CloudNativeCon

Europe 2019

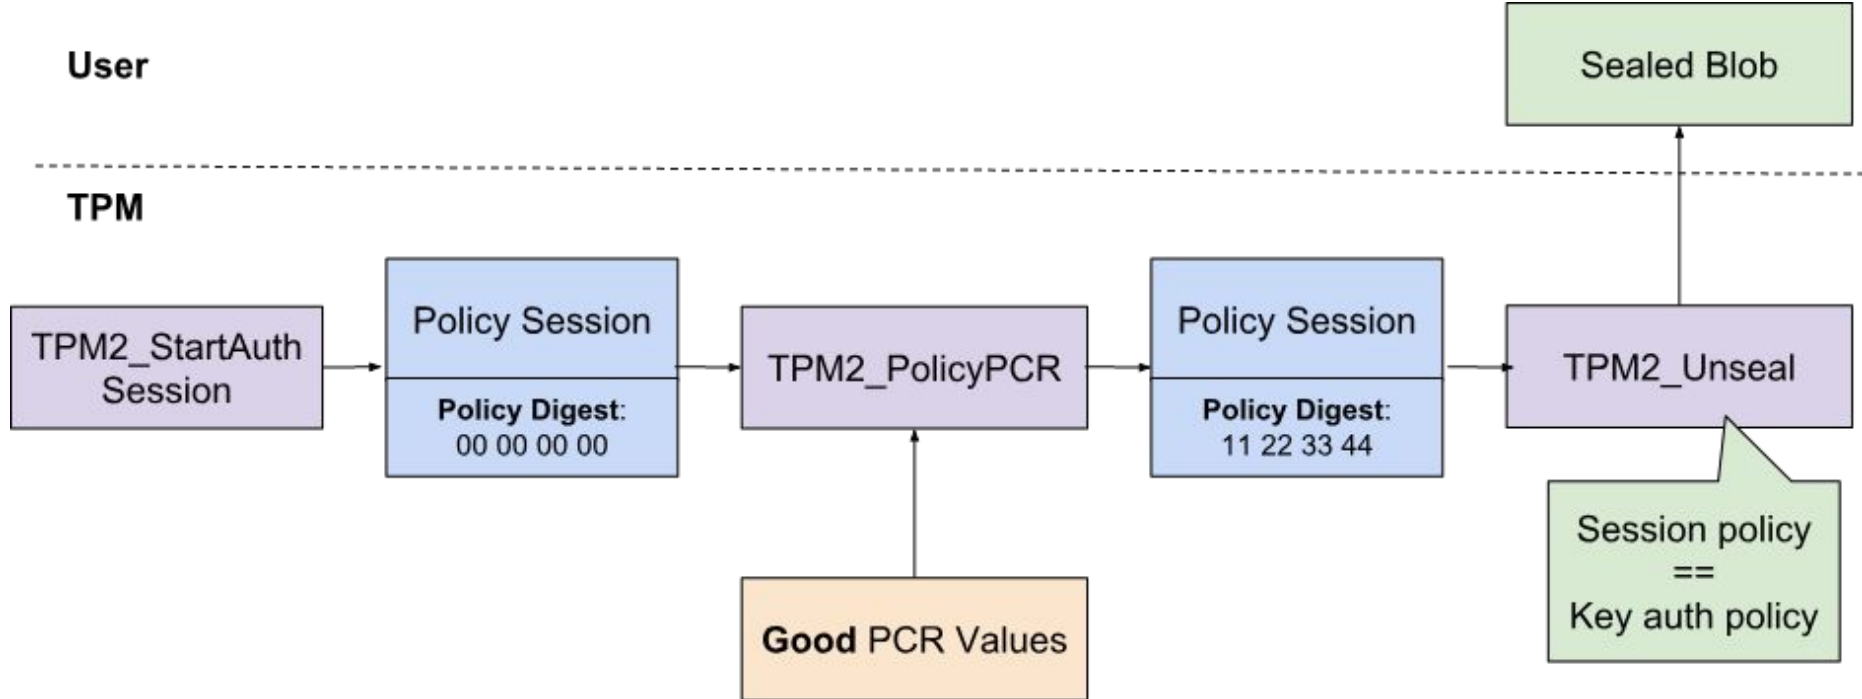
Backup Slides

Sealing to PCR Values



source <https://google.github.io/tpm-js>

Unsealing



Summarizing Tamper Evident Logging

On-prem

Cloud Provider

system/my-api-key

proxy

kube-api

kms-plugin

etcd

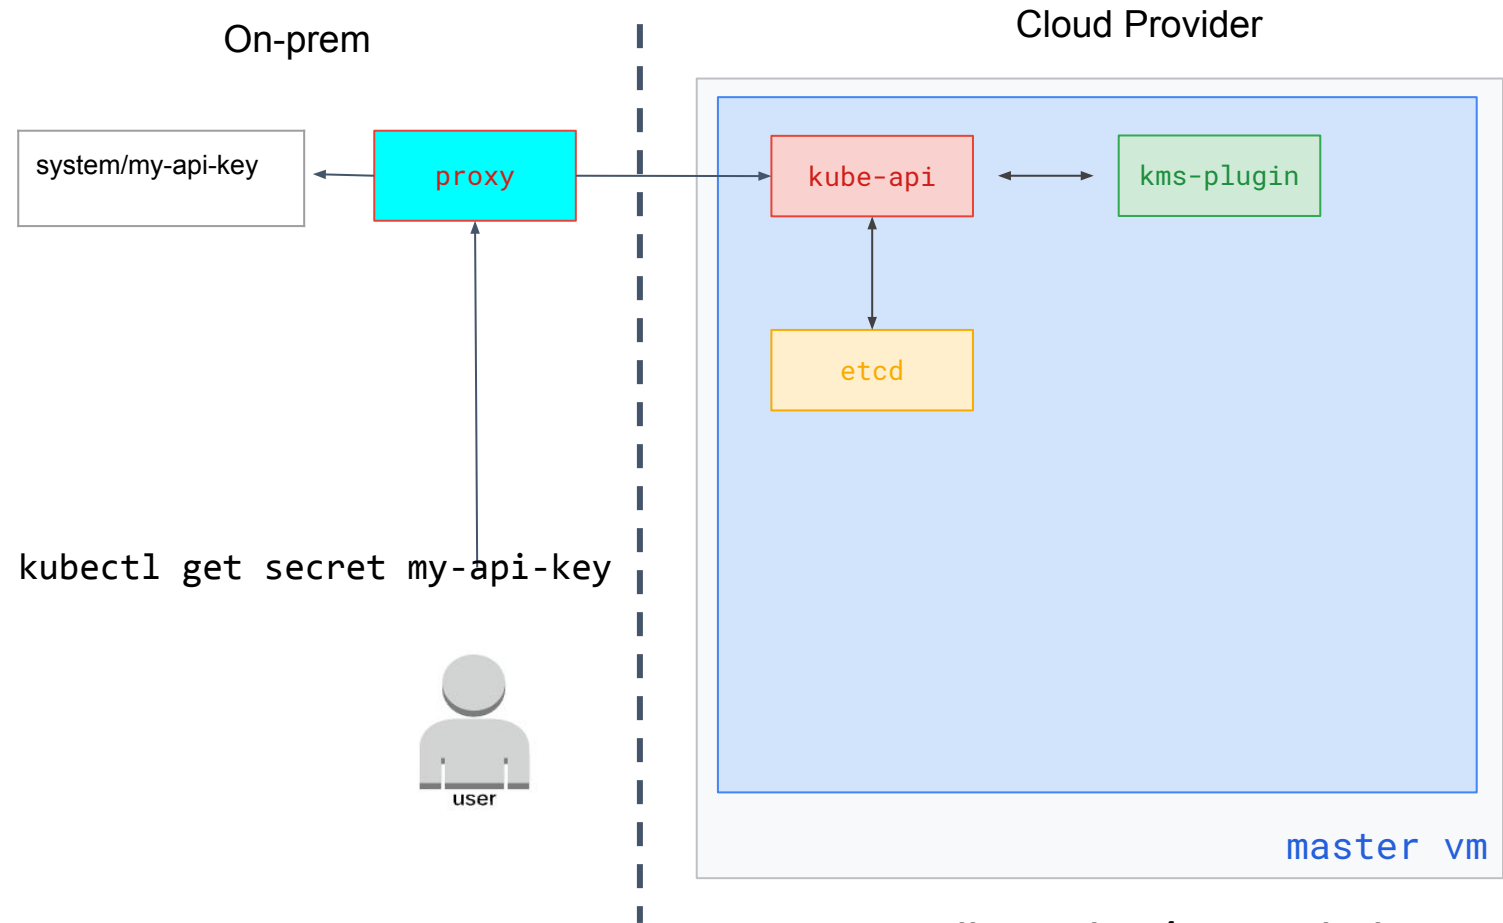
master vm

kubectl get secret my-api-key



user

Putting it all together (External Observer - Proxy)



On-prem

Cloud Provider

system/my-api-key

proxy

kube-api

kms-plugin

etcd

TPM2_HMAC system/my-api-key

Symmetric
Key

TPM

Audit: d81d4b2...

master vm

kubectl get secret my-api-key



user

Putting it all together (Creating an Audit Trail)

On-prem

Cloud Provider

system/my-api-key

proxy

kube-api

kms-plugin

etcd

Symmetric
Key

TPM

Audit: d81d4b2...

master vm

KMS

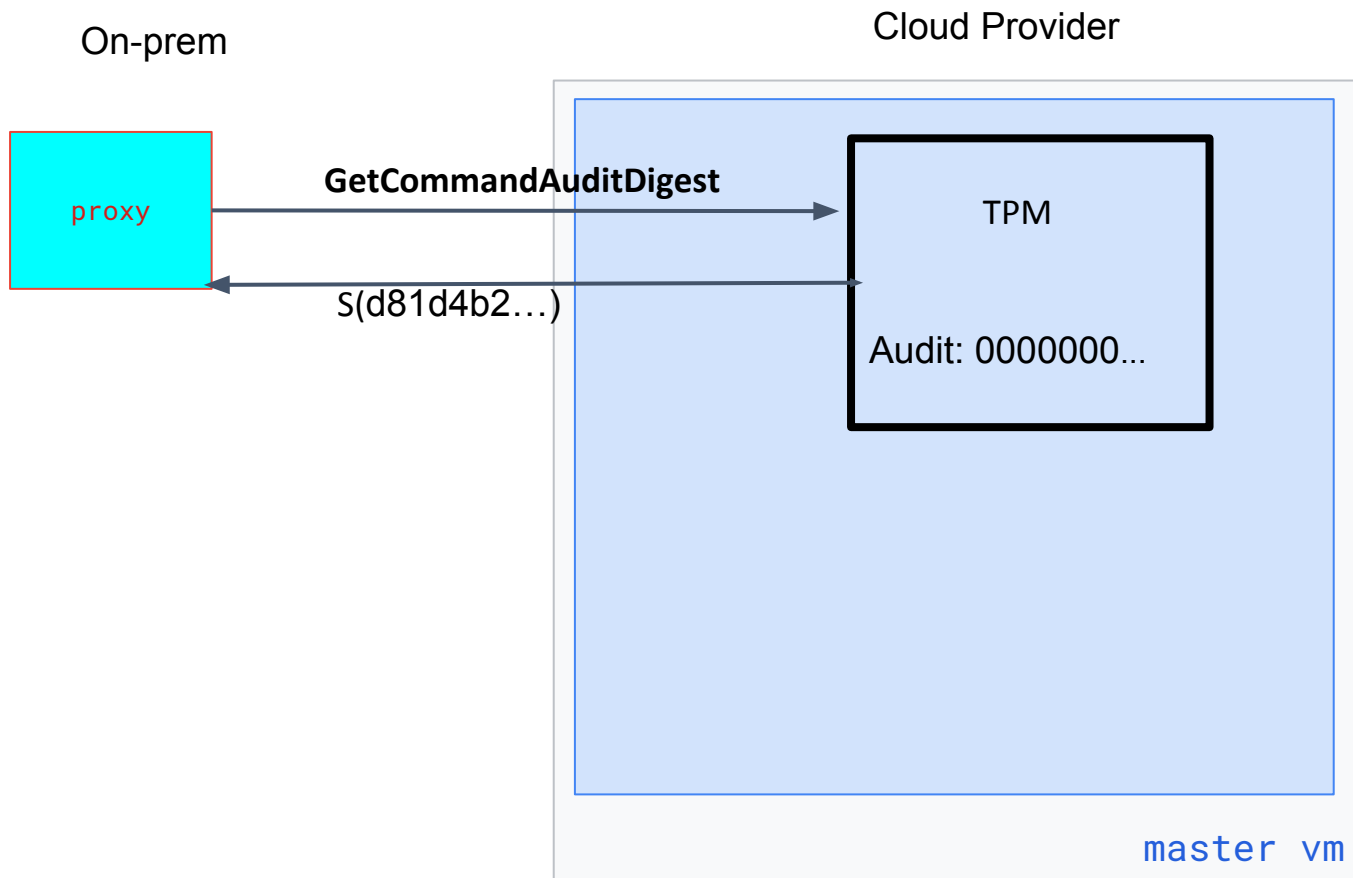
AAD: TPM2_HMAC(KEY)

kubectl get secret my-api-key

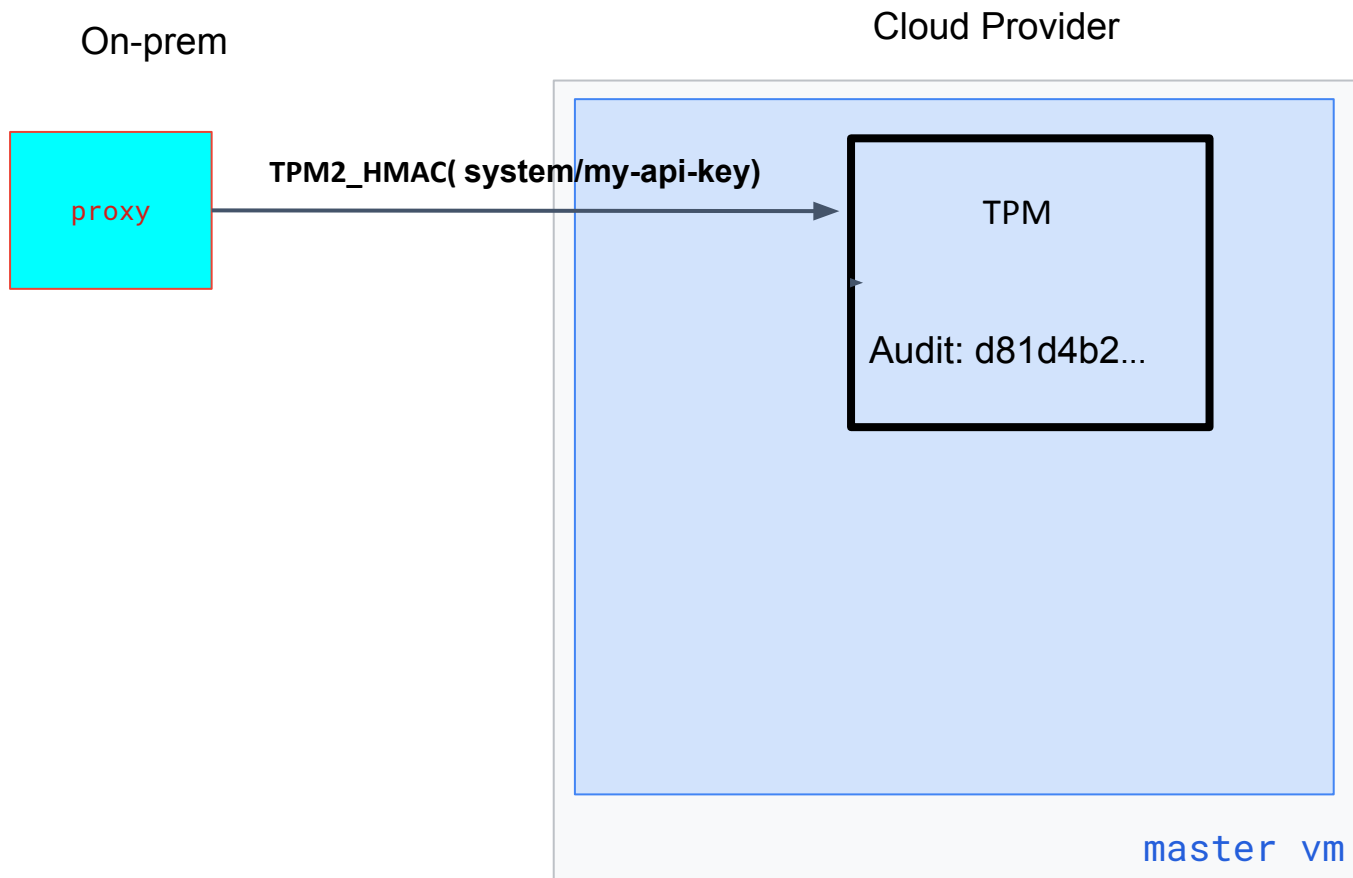


user

Putting it all together (Locking the Encrypted Secret to TPM)



Validation - Get Signed Value of the Audit Register



Validation - Replay TPM2_HMAC commands

On-prem

Cloud Provider



GetCommandAuditDigest

TPM

S(d81d4b2...)

Audit: 000000.....

D81d4b2... == D81d4b2...

master vm

Validation - Get Signed Value of the Audit Register

